

**Question 1.** Go through the points by decreasing  $x$  coordinates, maintaining as you go along the largest  $y$  coordinate encountered so far (call it  $\hat{y}$ ). A point  $p_i$  is processed by comparing its  $y$  coordinate to  $\hat{y}$ : If  $y_i < \hat{y}$  then that point is ignored (and you move on to the next point), but if  $\hat{y} < y_i$  then you update  $\hat{y}$  to be  $y_i$  and output  $p_i$  as being in  $M(S)$ . The  $O(n \log n)$  time complexity is because of sorting the points according to their  $x$  coordinates (after sorting the rest of the algorithm takes linear time).

**Question 2.** 1, 5, 7, 10, 12

**Question 3.**

1.  $g, c, a, b, d, h, e, f$
2. (a)  $g: 3; c: 3; d: 2; h: 2$   
(b) See the figure on the next page.  
(c) The total number of events ever inserted in the event list is  $2n + t$ , and similarly for the events deleted from that list. This, and the fact that a manipulation of the event list gives rise to at most two intersection discoveries, together imply that the total number of multiple discoveries is  $O(n + t)$ .

**Question 4.** The idea is to partition  $S$  into  $n/p$  contiguous chunks of size  $p$  each. Each of the windows of size  $p$  whose  $s_i$  we seek to compute either

1. coincides with one of the above-mentioned chunks,
2. overlaps with two adjacent such chunks.

Case 1 occurs for only  $n/p$  of the  $s_i$  we seek, and we can afford to spend  $O(p)$  time on each. The main difficulty is Case 2, which occurs for  $O(n)$  of the  $s_i$  we seek: We cannot afford to spend more than constant time on each. Note, however, that in Case 2 the overlap is with a suffix of the left chunk, and with a prefix of the second chunk. This observation suggests a pre-processing step in which we compute, for each of the  $n/p$  chunks, every prefix-max value (= maximum value in every prefix of that chunk) and every suffix-max value (= maximum value for every suffix of that chunk). This pre-processing takes  $O(p)$  time per chunk, hence  $O(n)$  total. It makes possible constant-time computation of every  $s_i$ : If the window for that  $s_i$  overlaps with chunks  $k$  and  $k + 1$  then  $s_i$  is the larger of (i) the suffix-max value of chunk  $k$  corresponding to the amount of overlap with that chunk; and (ii) the prefix-max value of chunk  $k + 1$  corresponding to the amount of overlap with that chunk. Because these two values are already available (from the pre-processing stage), each  $s_i$  is computed using one comparison.

If the above is not clear enough, below is a more formal description.

1. Partition  $S$  into  $n/p$  contiguous chunks of size  $p$  each (the last chunk could be smaller). We call  $S_k$  the  $k$ th such chunk, i.e.,  $S_k = x_{(k-1)p+1}x_{(k-1)p+2} \cdots x_{(k-1)p+p}$ .

2. Do the following for  $k = 1, \dots, n/p$  in turn:

(a) Compute in  $O(p)$  time the quantities  $L_{k,1}, \dots, L_{k,p}$ , where

$$L_{k,i} = \max\{x_{(k-1)p+1}, \dots, x_{(k-1)p+i}\}$$

(i.e.,  $L_{k,i}$  is the maximum of the leftmost  $i$  items of chunk  $S_k$ ). This can be done by a left to right walk along  $S_k$  that keeps track of the maximum encountered so far.

(b) Compute in  $O(p)$  time the quantities  $R_{k,1}, \dots, R_{k,p}$ , where

$$R_{k,i} = \max\{x_{(k-1)p+i}, \dots, x_{(k-1)p+p}\}$$

(i.e.,  $R_{k,i}$  is the maximum of the rightmost  $p - i + 1$  items of chunk  $S_k$ ). This can be done by a right to left walk along  $S_k$  that keeps track of the maximum encountered so far.

3. For  $i = 1, 2, \dots, n - p + 1$  compute  $s_i$  in constant time as follows:

- Let  $k = \lfloor i/p \rfloor$  and let  $j = i \bmod p$  (i.e.,  $j = i - kp$ ).
- If  $j = 1$  then  $s_i = R_{k,1}$ .
- If  $1 < j \leq p - 1$  then  $s_i = \max\{R_{k,j}, L_{k+1,j-1}\}$ .
- If  $j = 0$  then  $s_i = \max\{R_{k,p}, L_{k+1,p-1}\}$ .




Figure 1: The answer to question 3.2.b.