

ECE437: Introduction to Digital Computer Design

Chapter 2 - Instructions

Fall 2016

Announcements

- Lab1 is on basic 270/337 material and is not covered in 437 lectures
- Homeworks are due in class, at the beginning of the lecture
- Slides and homeworks posted on the course website.
 - Slides will be posted every couple of lectures

ECE437, Fall 2016 © Vijaykumar

(2)

8/26/2016

Instructions

- Instructions are the "words" of a computer
- Instruction set architecture (ISA) is its "vocabulary"
- With a few other things, this defines the **interface** of computers
 - But **implementations** vary
- We will study **MIPS™ ISA**
 - but the next ISA is not too hard after the first
- We won't write programs - EE362


ECE437, Fall 2016 © Vijaykumar

(3)

8/26/2016

Computer

- What does it do?
 - Fetch instruction from address in Program Counter (PC)
 - Execute instruction
 - Update PC to point to next instruction



ECE437, Fall 2016 © Vijaykumar

(4)

8/26/2016

Outline

- Instructions
 - Basics
 - Registers and ALU ops
 - Memory and load/store
 - Branches and jumps
 - And more . . .
- Read Chapter 2 and skim through Appendix E (on CD) for MIPS ISA

ECE437, Fall 2016 © Vijaykumar

(5)

8/26/2016

Basics

- C statement
- $f = (g + h) - (i + j)$
- MIPS instructions*
 - add $t0, g, h$ // $t0 = g + h$
 - add $t1, i, j$ // $t1 = i + j$
 - sub $f, t0, t1$ // $f = t0 - t1$
- Opcodes/mnemonic, operands, source/destination

ECE437, Fall 2016 © Vijaykumar

(6)

8/26/2016

Basics

- Opcode: Specifies the kind of operation (mnemonic)
- Operands: input and output data (source/destination)
- Operands $t0, t1$ are temps
- One operation, two inputs, one output
- Multiple instructions for one C statement

ECE437, Fall 2016 © Vijaykumar

(7)

8/26/2016

Why not have bigger instructions?

- Why not have " $f = (g + h) - (i + j)$ " as one instruction?
- Church-Turing thesis: A very primitive computer can compute anything that a fancy computer can compute
 - need only logical functions, read and write memory and data-dependent decisions
- Therefore, ISA selected for practical reasons: performance and cost, not computability
- Simplicity/regularity tend to improve both
 - E.g., H/W to handle arbitrary number of operands
 - complex, slow and rarely usable and NOT NECESSARY

ECE437, Fall 2016 © Vijaykumar

(8)

8/26/2016

Registers and ALU ops


- Ok, I lied!
- Operands must be registers, not variables
 - add \$8, \$17, \$18 // $\$8 \leftarrow \$17 + \$18$ - destination LEFT MOST
 - if you confuse this, you won't get anything right in 437
 - add \$9, \$19, \$20
 - sub \$16, \$8, \$9
- MIPS has 32 registers \$0-\$31 (figure next slide)
- Registers \$8, \$9 are temps, \$16 - f, \$17 - g, \$18 - h, \$19 - i, \$20 - j
- MIPS also allows one constant called "immediate"
 - later we will see immediate is 16 bits [-32768, 32767]

ECE437, Fall 2016 © Vijaykumar

(9)

8/26/2016

Registers and ALU



ECE437, Fall 2016 © Vijaykumar

(10)

8/26/2016

ALU ops

- Some ALU ops:
 - add, addi, addu, addiu (immediate, unsigned)
 - Caution: casting signed to unsigned in System Verilog may cause weirdness
 - sub . . .
 - mul, div - weird!
 - and, andi
 - or, ori
 - sll, srl, . . .
- Why registers? fit in instructions, smaller is faster
- Are registers enough?

ECE437, Fall 2016 © Vijaykumar

(11)

8/26/2016

Memory and load/store


- But need more than 32 words of storage
- Memory - an array of locations $M[addr]$, indexed by addr (figure next slide)
- Data movement (on words or integers)
 - load word for reg \leftarrow memory
 - lw \$17, 1002 // get input g
 - store word for reg \rightarrow memory
 - sw \$16, 1001 // save output f
 - Note: destination LAST for stores!

ECE437, Fall 2016 © Vijaykumar

(12)

8/26/2016

Memory and load/store



ECE437, Fall 2016 © Vijaykumar

(13)

8/26/2016

Byte vs. Word addresses


- I lied again!
 - we need to address bytes for characters
 - words take 4 bytes
 - therefore, word addresses must be multiples of 4 (i.e., end in 00)
 - Addresses are ALWAYS byte addresses (next slide)
 - NOT word addresses (previous slide)
- figure next slide

ECE437, Fall 2016 © Vijaykumar

(14)

8/26/2016

Memory and load/store



ECE437, Fall 2016 © Vijaykumar

(15)

8/26/2016

Memory and load/store

- Important for arrays
 - $A[i] = A[i] + h$ (figure next slide)
 - # \$8 - temp, \$18 - h, \$21 - $(i \times 4)$
 - Astart is 0x8000, A is an array of words NOT bytes
- ```
lw $8, Astart($21) // == 0x8000($21)
 // Astart or 0x8000 is offset
 // $8 <- M[Astart+$21]
 // or $8 <- M[0x8000+$21]


add $8, $18, $8
sw $8, Astart($21) //offset is signed, can be negative
```
- MIPS has other load/store for byte/halfword

ECE437, Fall 2016 © Vijaykumar

(16)

8/26/2016

## Memory and load/store



ECE437, Fall 2016 © Vijaykumar

(17)

8/26/2016

## Endian

- Word storage
- Word = bytes (0x400,0x401,0x402,0x403)?
- Word = bytes (0x403,0x402,0x401,0x400)?
  - It depends...
- **Big endian:** MSB at address xxxxxx00
  - e.g., IBM, SPARC
- **Little endian:** MSB at address xxxxxx11
  - e.g., Intel x86
- Mode selectable
  - e.g., PowerPC, MIPS

ECE437, Fall 2016 © Vijaykumar

(18)

8/26/2016

## Branches and Jumps

```
While (i != j) {
 j = j + i;
 i = i + 1;
}
```

• HLL → Assembly, \$8 is i, \$9 is j

```
Loop: beq $8, $9, Exit
 // note that the condition is reversed
 add $9, $9, $8
 addi $8, $8, 1
 j Loop
Exit:
```



ECE437, Fall 2016 © Vijaykumar

(19)

8/26/2016

## Branches and Jumps

- Better yet:
 

```
beq $8, $9, Exit // not !=
Loop: add $9, $9, $8
 addi $8, $8, 1
 bne $8, $9, Loop
Exit:
```
- Let compilers worry about such optimizations

ECE437, Fall 2016 © Vijaykumar

(20)

8/26/2016

## Branches and Jumps

- What does `bne` do really?
  - read `$8`; read `$9`, compare
  - set `PC = PC + 4` or `PC = Target`
- To do compares other than "equal" or "not equal"
  - e.g., `blt $8, $9, Target` // assembler PSEUDO-instruction
- expands to
  - `slt $1, $8, $9` // set less than `$1 = ($8 < $9)?1:0`
  - `bne $1, $0, Target` // `$0` is always 0

ECE437, Fall 2016 © Vijaykumar

(21)

8/26/2016

## Branches and Jumps

- Other MIPS branches/jumps
  - `beq $8, $9, imm`
  - // if (`$8 == $9`) `PC = PC + imm<<2` else `PC += 4`
- `bne` ...
- `slt, sle, sgt, sge`
  - with immediate, unsigned
- `j addr` // `PC = addr`
- `jr $12` // `PC = $12`
- `jal addr` // `$31 = PC+4; PC = addr; used for procedure calls`

ECE437, Fall 2016 © Vijaykumar

(22)

8/26/2016

## Assembly Exercise

|                                         |       |                               |
|-----------------------------------------|-------|-------------------------------|
| for(i=0;i<n;i++) {<br>blah<br>blah<br>} | Init: | <code>mov \$8,0</code>        |
|                                         | Test: | <code>slt \$5,\$8,\$9</code>  |
|                                         |       | <code>beq \$5,\$0,Exit</code> |
|                                         | Loop: | <code>blah</code>             |
|                                         |       | <code>blah</code>             |
|                                         | Ind:  | <code>addi \$8,\$8,1</code>   |
|                                         |       | <code>j Test</code>           |
|                                         | Exit: |                               |

ECE437, Fall 2016 © Vijaykumar

(23)

8/26/2016

## Assembly Exercise

- HLL-> assembly
- if (`a=b`)
  - `a = a+b`
- else
  - `b = a+b`

ECE437, Fall 2016 © Vijaykumar

(24)

8/26/2016

## Assembly Exercise

- HLL → assembly
- ```

if (a=b)
    a = a+b
else
    b = a+b
    
```
- ```

lw $8, a // a → $8
lw $9, b // b → $9
beq $8,$9,L_IF
j L_ELSE
L_IF:
add $8,$8,$9 // a = a+b
sw $8,a // write a to mem
j EXIT
L_ELSE:
add $9,$8,$9 // b = a+b
sw $9,b // write b to mem
EXIT:

```

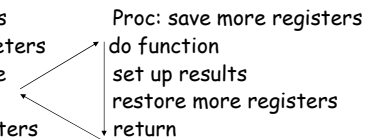
ECE437, Fall 2016 © Vijaykumar

(25)

8/26/2016

## Procedure Calls

- See section 2.7 for more details
  - save registers
  - set up parameters
  - call procedure
  - get results
  - restore registers
- jal is the ONLY operation in hardware: the rest is in software
- jal addr // \$31 = PC+4 (return addr); PC = addr (UNLIKE 362 where return addr is saved on stack - too complicated/slow to do in fast clock)



ECE437, Fall 2016 © Vijaykumar

(26)

8/26/2016

## Procedure Calls

- An important data structure is the stack
- Stack grows from larger to smaller addresses
- \$29 is the stack pointer, it points just beyond valid data
- push \$2:
  - addi \$29, \$29, -4
- pop \$2:
  - lw \$2, 4(\$29)
  - addi \$29, \$29, 4
- push, pop NOT real instructions
- Addi-sw/lw-addi order cannot change. Why?

ECE437, Fall 2016 © Vijaykumar

(27)

8/26/2016

## Procedure Example

- HLL code for swap
 

```

{ temp = v[k]; // code for swap */
 v[k] = v[k+1];
 v[k+1] = temp;
 return (temp); }

```
- Corresponding assembly code
 

```

swap: addi $29, $29, -8 // save registers
 sw $15, 4($29)
 sw $16, 8($29)

```

ECE437, Fall 2016 © Vijaykumar

(28)

8/26/2016

## Procedure Example

```

multi* $2, $5, 4 // procedure body: $2 <- k*4
add $2, $4, $2 // $2 <- v + k*4 == address of v[k]
lw $15, 0($2) // $15 <- v[k]
lw $16, 4($2) // $16 <- v[k+1]
sw $16, 0($2) // v[k] <- $16
sw $15, 4($2) // v[k+1] <- $15
mov* $2, $15 // return value
lw $15, 4($29) // restore registers
lw $16, 8($29)
addi $29, $29, 8
jr $31 // return

```

- All push/pop in s/w

ECE437, Fall 2016 © Vijaykumar

(29)

8/26/2016

## Procedure Call Example

- Now, calling the procedure
    - HLL code
- ```
return_value = swap(arr_A, i); /* call swap */
```

```

mov     $4, $18      // first parameter
mov     $5, $17      // second parameter
jal     swap
mov     $20, $2       // copy return value

```

ECE437, Fall 2016 © Vijaykumar

(30)

8/26/2016

Layers of Software

- Notation - program: input data → output data
 - executable: input data → output data
 - loader: executable file → executable in memory
 - linker: object files → executable file
 - assembler: assembly file → object file
 - compiler: HLL file → assembly file
 - editor: editor commands → HLL file
- Only possible because programs can be manipulated as data
 - "stored program computer", "Von Neumann architecture".

ECE437, Fall 2016 © Vijaykumar

(31)

8/26/2016

362 vs. 437

- Shifts in 362 vs. 437
 - Freescale HCS (362) shifts include the Carry
 - MIPS (437) shifts do not
- In MIPS, the Carry is not part of the ISA
 - Carry is an internal signal to detect overflows
 - Carry is not visible to the ISA
 - branches use general-purpose registers and not condition codes

ECE437, Fall 2016 © Vijaykumar

(32)

8/26/2016

MIPS Machine Language

- All 32-bit instructions
 - Assembly: add \$1, \$2, \$3
 - 14 ASCII (or Unicode) codes
 - 61 64 64 20 24 31 2C 20 24 32 2C 20 24 33 (hex)
 - Machine code: 1 word
- 000000 00010 00011 00001 00000 010000
- | | | | | | |
|--------|---|---|---|------|------------|
| alu-rr | 2 | 3 | 1 | zero | add/signed |
|--------|---|---|---|------|------------|


ECE437, Fall 2016 © Vijaykumar

(33)

8/26/2016

Instruction Format

- R-format:



- Digression:
 - How do you store the number 4,392,976?
 - Same as add \$1, \$2, \$3
- Stored program: instructions are represented as numbers
 - programs can be read/written in memory like numbers
- Other R-format: addu, sub*, and, or etc

ECE437, Fall 2016 © Vijaykumar

(34)

8/26/2016

Instruction Format

- Assembly: lw \$1, 100(\$2)
- Machine:

100011 00010 00001 00000000001100100

lw 2 1 100 (in signed binary)
- I-format:

opcode	rs	rt	address/immediate
6	5	5	16 (signed)

 - Addr/immediate is 16 bits signed 2's complement

ECE437, Fall 2016 © Vijaykumar

(35)

8/26/2016

Instruction Format

- I-format also used for ALU ops with immediates
 - addi \$1, \$2, -4
 - 001000 00010 00001 111111111111100
- What about constants larger than 16 bits = [-32768, 32767]?
 - e.g., 0000 0000 0000 1100 0000 0000 0000 1111 or 0x000C000F?
- lui \$4, 12 // \$4 <- 0x000C0000
- ori \$4, \$4, 15 // \$4 <- 0x000C000F
- All loads and stores use I-format

ECE437, Fall 2016 © Vijaykumar

(36)

8/26/2016

Instruction Format

- I-format for branches: beq \$1, \$2, 7
- 000100 00001 00010 0000 0000 0000 0111
- $PC = PC + (0000\ 0111 \ll 2)$ // word offset

Finally, **J-format**: opcode addr
6 26

- addr is weird in J-format:
 - Target address = 4 MSB of PC:addr:00 ($4+26+2 = 32$)
 - Why MSB of PC? We have 26, we need 30
 - Can we Pad with 0s? Restricts jumps to top 16^{th} of address space


ECE437, Fall 2016 © Vijaykumar

(37)

8/26/2016

MIPS: branch vs jump

- Branch: 16 bit displacement
 - Current word +/- 2^{15}
 - Range is symmetric
 - $\text{Range}_{(\text{forward})} = \text{Range}_{(\text{backward})} = 2^{15}$
- Jump: 26 lower-bit concatenation in word address
 - Jump anywhere within a 2^{26} segment
 - Range may be asymmetric
 - $\text{Range}_{(\text{forward})} + \text{Range}_{(\text{backward})} = 2^{26}$



ECE437, Fall 2016 © Vijaykumar

(38)

8/26/2016

Summary: Instruction Formats

- R-format: opcode rs rt rd shamt funct
6 5 5 5 5 6
- I-format: opcode rs rt address/immediate
6 5 5 16
- J-format: opcode addr
6 26
- Instr decode
 - Theory: Inst bits \rightarrow identify instrs \rightarrow control signals
 - Practice: Instruction bits \rightarrow control signals


ECE437, Fall 2016 © Vijaykumar

(39)

8/26/2016

Addressing modes

- There are many ways of accessing data
- 1. Register addressing
 - add \$1, \$2, \$3




ECE437, Fall 2016 © Vijaykumar

(40)

8/26/2016

Addressing Modes

- 2. Base addressing (aka displacement)
 - `lw $1, 100($2)` // $\$2 == 400, M[500] == 42$




ECE437, Fall 2016 © Vijaykumar

(41)

8/26/2016

Addressing Modes

- 3. Immediate addressing
 - `addi $1, $2, 100`




ECE437, Fall 2016 © Vijaykumar

(42)

8/26/2016

Addressing Modes

- 4. PC relative addressing
 - `beq $1, $2, 100` // if $(\$1 == \$2) PC = PC + 100$




ECE437, Fall 2016 © Vijaykumar

(43)

8/26/2016

Addressing Modes

- Pseudodirect addressing
 - `j Loop` // instruction contains address*



ECE437, Fall 2016 © Vijaykumar

(44)

8/26/2016

Addressing Modes


NOT found in MIPS

- 5. Indexed: add two registers - base + index
- 6. Indirect: $M[M[addr]]$ - two memory references
- 7. Autoincrement/decrement: add data size
- 8. Autoupdate - found in IBM PowerPC, HP PA-RISC
 - like displacement but update register

Addressing Modes

• Autoupdate

lwupdate \$1, 4(\$2) // $\$1 == M[4+\$2]$; $\$2 == \$2 + 4$



Addressing Modes

for ($i = 0, i < N, i += 1$)
 sum += A[i];

- \$7 - sum, \$8 - address of a[i], \$2 - tmp

inner: new inner:
 lw \$2, 0(\$8) lwupdate \$2, 4(\$8)
 addi \$8, \$8, 4 }
 add \$7, \$7, \$2 add \$7, \$7, \$2

- any problems with new inner ?

How to Choose ISA

- Minimize what?
 - Two way constraint satisfaction
 - What can be implemented in hardware? From below
 - What ISA is better for optimized compilation? From above
- In 1985-1990 technology, simple modes like MIPS good
 - Easy to implement, not a real reason in the business world
 - Very compelling reason in the classroom
- As technology changes, computer design options change
- For small memory, dense instructions important and dense usually complex/irregular
- For high speed, pipelining important and for pipelining simple/regular important

Intel 80386 ISA

- 8 32-bit registers
- Two register machine: src1/dst src2
 - reg - reg, reg - imm, reg - mem, mem - reg, mem - imm
- seven addressing modes
- 16-bit and 32-bit ops on memory and registers
 - 8-bit prefix to override default data size
 - condition codes
 - part of normal operation, extends operation time, often not looked at

ECE437, Fall 2016 © Vijaykumar

(49)

8/26/2016

Intel 80386 ISA

- Decoding nightmare
 - instructions 1 to 17 bytes
 - prefixes, postfixes
 - crazy "formats" - e.g., register specifiers move around
 - but key instructions not terrible
 - yet have to make ALL work correctly

ECE437, Fall 2016 © Vijaykumar

(50)

8/26/2016

Current Approach

- Current Intel chips
 - Instruction decode logic translates into "RISCy Ops"
 - RISC - Reduced Instruction Set Computer
 - Execution unit runs RISCy ops
 - Backward compatibility
 - Complex decoding
 - Execution unit as fast as RISC
- We work with MIPS to keep it simple
- Learn x86 on the job!

ECE437, Fall 2016 © Vijaykumar

(51)

8/26/2016

Complex Instructions

- More powerful instructions not necessarily faster execution
- E.g. - string copy
 - Option 1: move with repeat prefix for memory to memory move
 - special-purpose
 - Option 2: use loads into register and then stores
 - generic instructions
- Option 2 faster on the same machine!

ECE437, Fall 2016 © Vijaykumar

(52)

8/26/2016

Outline

- So far, too MIPS-centric
 - Rich design space:
 - Other possibilities: Stack and Accumulator architectures :Ch 2 (on CD)

ECE437, Fall 2016 © Vijaykumar

(53)

8/26/2016

General Purpose Architecture

- General Purpose Register architecture
 - 3 operands: 2 source + 1 dest
 - Depends on addressing modes
 - Register-register only
 - MIPS : all register operands except load/store
 - Memory-memory (all memory operands)

- Eg

- $f = (g+h) * (i-j)$

```
Add t1, g, h
Sub t2, i, j
Mul f, t1, t2
```

ECE437, Fall 2016 © Vijaykumar

(54)

8/26/2016

Stack Machine

- Push, Pop with one address
 - Last in, first out
- Binary/Unary Op operates on
 - Implicit operands
 - Top one/two element(s) of stack
- Eg
 - $f = (g+h) * (i-j)$

```
Push g
Push h
Add
Push i
Push j
Subtract
Mul
```

ECE437, Fall 2016 © Vijaykumar

(55)

8/26/2016

Accumulator Machine

- One address
 - 2nd implicit operand: accumulator
 - accumulator is 2nd source operand as well as destination
- Eg
 - $f = (g+h) * (i-j)$

```
Load g // Acc = g
Add h // Acc = Acc+h
Store t // t = g+h
Load i // Acc = i
Subt j // Acc = Acc-j
Mul t1 // Acc = Acc*t
Store f
```

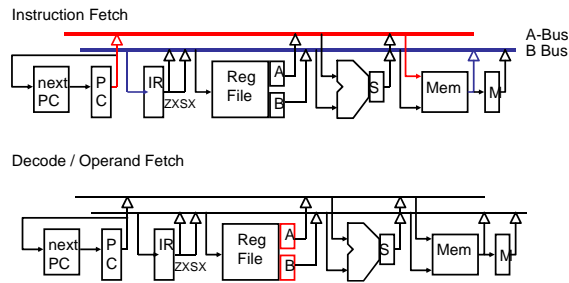
ECE437, Fall 2016 © Vijaykumar

(56)

8/26/2016

2-Bus Datapath

- Alternate datapath: Shared bus
- E.g. Load instruction




ECE437, Fall 2016 © Vijaykumar

(57)

8/26/2016

2-bus Datapath



ECE437, Fall 2016 © Vijaykumar

(58)

8/26/2016

Concluding Remarks

- Simple and regular
 - same length instructions, fields in same place
- Small and fast
 - Small number of registers
- Compromises inevitable
 - There is **NO PERFECT ISA!**
 - Pipelining (more later) should not be hindered
- Common case fast
- Read Chapter 2
- Read Chapter 4a 4.1-4.4 (coming up next)
 - Ch 4a+b is 1/3rds of this course, lab

ECE437, Fall 2016 © Vijaykumar

(59)

8/26/2016