

ECE437: Introduction to Digital Computer Design

Chapter I/O

Fall 2016

Motivation

- I/O needed for
 - To/from users
 - To/from computers
 - To/from non-volatile storage media
- I/O Performance matters
- Total time = CPU + I/O - overlap
 - $10 + 4 - 4 = 10$; 1x
 - $5 + 4 - [0,4] = [9,5]$; [1.1x, 2x]
 - $1 + 4 - [0,1] = [5,4]$; [2x, 2.5x]

ECE437, Fall 2016

(2)

I/O Performance

- What is performance?
 - With CPU : Iron Law
 - With I/O
 - Applications have different I/O needs
- Supercomputers read and write 1G of data
 - High data throughput
- Transactions processing does many independent, small I/O ops.
 - Fast I/O throughput (# of I/Os per sec)
- File systems
 - Fast response time, locality

ECE437, Fall 2016

(3)

I/O Device Characteristics

See Figure 6.2

Device	I or O?	Partner	Data Rate Mb/s
Mouse	I	Human	0.0038
Graphics Display	O	Human	800-8000
Modem	I/O	Machine	0.016-0.064
LAN	I/O	Machine	100-1000
Tape	Storage	Machine	32
Disk	Storage	Machine	240-2560

ECE437, Fall 2016

(4)


Disk Trends

- Disk Trends
 - \$/MB decreasing
 - 2010: \$125/2TB = 6.3c/GB
 - 2008: \$149/1.5TB = 10c/GB
 - 2006: \$92/250 GB = 36c/GB
 - 2004: 50c/GB (was 10000c/GB in 1997)
 - OEM prices lower
 - Disk diameter 14" → 1.8" → 1" **
 - Seek time down ~10ms
 - Rotation speed unchanged
 - ~7200 rpm for consumer
 - ~15K rpm for high end disks
 - Xfer rates up
 - Move from parallel buses to Serial bus (Serial ATA)

ECE437, Fall 2016

(5)

Magnetic Disks



- Stack of platters
 - Two surfaces/platter
- Concentric tracks
- Several sectors/track
- Disk Access Time = Queue time + service time


ECE437, Fall 2016

(6)

Magnetic Disks

• Service time

- Seek track
 - ~10-20ms
 - Better with locality
- Rotation time (sector)
 - $\frac{1}{2}$ rotation on average
 - 3600 rpm (8.3ms), 5400 rpm (5.6ms), 7200 rpm (4.1ms)
- Xfer
 - Page size transfers ~4-8 KB
 - Several MB/s



ECE437, Fall 2016

(7)

Disk Access Time

- Disks R' us™ has a **GenX** disk with the following parameters
 - Average seek time = 10ms
 - Rotation speed = 5400 rpm
 - Xfer rate = 160MB/s
 - Disk-block size = 4KB
 - Controller overhead = 3ms
- They introduce a **GenY** disk technology with rotation speed of 7200 rpm which they advertise as being 50% faster
- If average queue time is 0 ms, what is the true speedup of **GenY** over **GenX** when reading a randomly chosen disk-block.

ECE437, Fall 2016

(8)

Disk Access Time

- GenX access time =
 - Queue time + controller overhead + seek time + rotational latency for $\frac{1}{2}$ rotation + Xfer time
 - $0 + 3 + 10 + (0.5 \times 60 / 5400 \times 10^3) \text{ms} + (4\text{K} / 160\text{M} \times 10^3) \text{ms}$
 - $0 + 3 + 10 + 5.6 + 0.025$
 - 18.625ms
- GenY access time
 - $0 + 3 + 10 + (0.5 \times 60 / 7200 \times 10^3) \text{ms} + (4\text{K} / 160\text{M} \times 10^3) \text{ms}$
 - $0 + 3 + 10 + 4.1 + 0.025$
 - 17.125 ms
- Speedup = $18.625 / 17.125 = 1.087 = 8.7\%$

ECE437, Fall 2016

(9)


Exercise

- A disk has the following parameters
 - Access size : 8KB
 - Bandwidth: 80MB/s
 - Controller overhead : 0.1 ms
 - Seek time : 6ms
 - Rotation speed: 7200 rpm
- What is the average access time per transfer?
- If the disk controller has a 8MB cache that results in a 20% hit rate, what is the average access time? (Hits eliminate seek time and rotational latency)

ECE437, Fall 2016

(10)

I/O {Depend/Reli/Avail}-ability



- MTTF Mean time to Failure
- MTTR Mean time to repair
- MTBF Mean time between Failure
 - $MTTR + MTTF$
- Reliability metric: MTTF
- Availability metric: $MTTF / MTBF$

ECE437, Fall 2016

(11)

Redundant Array of Inexpensive Disks

- What if we want to store 100 disks
- MTTF : $5 \text{ yr} / 100 = 18 \text{ days}$
 - RAID 0 = No redundancy (Striping)
 - RAID 1 = mirror = full data redundancy = 100% overhead
 - RAID 3 = bit interleaved parity = small overhead
 - Dedicated parity disk
 - RAID 4 = block interleaved parity = small overhead + small writes
 - RAID 5 = distributed block wise parity = small overhead + small writes
- Other raid levels in book

ECE437, Fall 2016

(12)

Recap

- Disks
 - Components of latency
 - Organization
 - RAID
 - Availability/Dependability/Reliability
- Buses
 - Types
 - One-bus/two-bus/three-bus architectures
 - Transaction timing
 - Synchronous vs. asynchronous

ECE437, Fall 2016

(13)


Connecting I/O to the CPU

- Many I/O devices with vastly different data rates
- I/O device economics
 - Need standards
 - Large number of peripheral device producers
 - Multiple business entities involved favors longer lasting standards
- Let's examine several connection strategies

ECE437, Fall 2016

(14)

Advantages of Buses




- **Versatility:**
 - New devices can be added easily
 - Peripherals can be moved between computer systems that use the same bus standard
- **Low Cost:**
 - A single set of wires is shared in multiple ways
- Manage complexity by partitioning the design

ECE437, Fall 2016

(15)

Disadvantages of Buses



- It creates a communication bottleneck
 - The bandwidth of that bus can limit the maximum I/O throughput
- The maximum bus speed is largely limited by:
 - The **length** of the bus
 - The **number** of devices on the bus
 - The need to support a range of devices with:
 - Widely varying latencies
 - Widely varying data transfer rates

ECE437, Fall 2016

(16)


Types of buses

- **Processor-Memory Bus** (design specific/proprietary)
 - Short and high speed
 - Only need to match the memory system
 - Maximize memory-to-processor bandwidth
 - Connects directly to the processor
 - Optimized for cache block transfers
- **I/O Bus** (industry standard)
 - Usually is lengthy and slower
 - Need to match a wide range of I/O devices
 - Connects to the processor-memory bus or backplane bus
- **Backplane Bus** (standard or proprietary)
 - Backplane: an interconnection structure within the chassis
 - Allow processors, memory, and I/O devices to coexist
 - Cost advantage: one bus for all components

ECE437, Fall 2016

(17)

One Bus (Backplane) Architecture




- A single bus (the backplane bus) is used for:
 - Processor to memory communication
 - Communication between I/O devices and memory
- **Advantages:** Simple and low cost
- **Disadvantages:** slow and the bus can become a major bottleneck
- Example: IBM PC - AT

ECE437, Fall 2016

(18)

Two Bus System




- I/O buses tap into the processor-memory bus via bus adaptors:
 - Processor-memory bus: mainly for processor-memory traffic
 - I/O buses: provide expansion slots for I/O devices
- Apple Macintosh-II
 - NuBus: Processor, memory, and a few selected I/O devices
 - SCCI Bus: the rest of the I/O devices

ECE437, Fall 2016

(19)

Three Bus System




- A small number of backplane buses tap into the processor-memory bus
 - Processor-memory bus is used for processor memory traffic
 - I/O buses are connected to the backplane bus
- **Advantage:** loading on the processor bus is greatly reduced

ECE437, Fall 2016

(20)

Real Stuff: PCI

- Older
 - Dual Independent Bus (DIB) Architecture between processor/memory/L2 cache
- Peripheral Component Interconnect
 - 33 (66) MHz, synchronous, 32 (64) bit data, multiple masters
- Claim: PCI too slow for hi-end graphics.
 - AGP taps directly into FSB

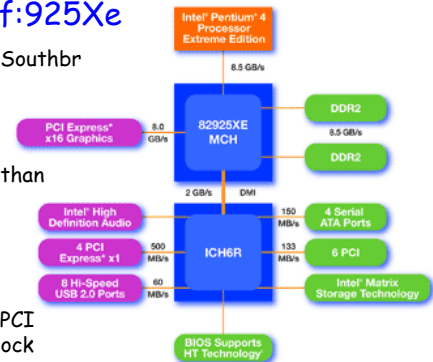


ECE437, Fall 2016

(21)

Real Stuff: 925Xe

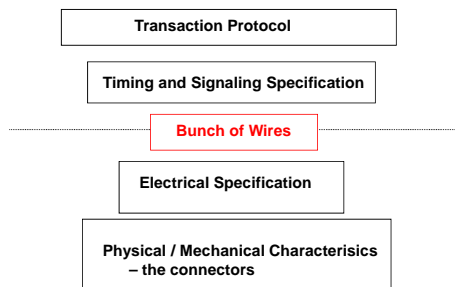
- Northbridge/Southbridge
- Serial faster than parallel
- PCI-e
- PCI-X is just PCI with faster clock



ECE437, Fall 2016

(22)

Bus Definition



ECE437, Fall 2016

(23)

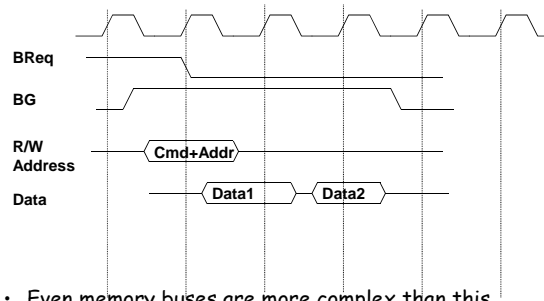
Synchronous vs. Asynchronous

- **Synchronous Bus:**
 - Includes a clock in the control lines
 - A fixed protocol for communication that is relative to the clock
 - Advantage: involves very little logic and can run very fast
 - Disadvantages:
 - Every device on the bus must run at the same clock rate
 - To avoid clock skew, they cannot be long if they are fast
- **Asynchronous Bus:**
 - It is not clocked
 - It can accommodate a wide range of devices
 - It can be lengthened without worrying about clock skew
 - It requires a handshaking protocol

ECE437, Fall 2016

(24)

Simple Synchronous Protocol




- Even memory buses are more complex than this
 - memory may take time to respond
 - need to control data rate

ECE437, Fall 2016

(25)

Typical Synchronous Protocol




- Device indicates when it is prepared for data xfer
- Actual transfer goes at bus rate

ECE437, Fall 2016

(26)

Asynchronous Bus Protocol



Proc asserts READREQ, <address>	Mem asserts ACK after reading <address>
Proc deasserts READREQ, <address> and waits	Mem deasserts ACK
	Mem asserts DATA-RDY, <data>
Proc asserts ACK after reading <data>	Mem deasserts DATA-RDY, <data>
Proc deasserts ACK	

ECE437, Fall 2016

(27)

I/O Interfacing: Outline

- How does the CPU initiate I/O?
- How does I/O device notify CPU when I/O op is complete?
- Who (on the CPU) performs I/O ops?

ECE437, Fall 2016

(28)

I/O Interfacing


- I/O operation needs to be initiated
 - Special opcodes
 - Memory-mapped I/O
- I/O completion must be known
 - Polling
 - Interrupt-driven

ECE437, Fall 2016

(29)

I/O commands

- **Memory-mapped I/O**
 - Special addresses not for memory (\$r9 contains special address for IOC)
 - Commands written (sw) as data (\$r4)
- **I/O commands**
 - Special opcodes
 - Send over I/O Bus



ECE437, Fall 2016

(30)

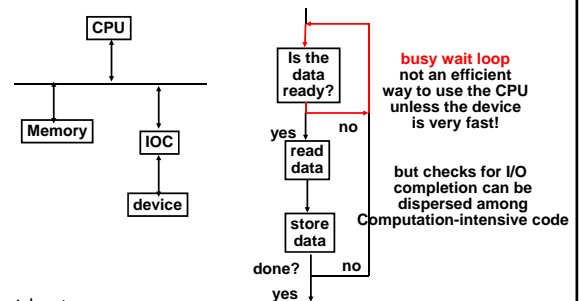
I/O Device Notifying the OS

- The OS needs to know when:
 - The I/O device has completed an operation
 - The I/O operation has encountered an error
- This can be accomplished in two different ways:
 - Polling:
 - The I/O device put information in a status register
 - The OS periodically check the status register
 - I/O Interrupt:
 - Whenever an I/O device needs attention from the processor, it interrupts the processor from what it is currently doing.

ECE437, Fall 2016

(31)

Polling: Programmed I/O




- Advantage:
 - Simple: the processor is totally in control and does all the work
- Disadvantage:
 - Polling overhead can consume a lot of CPU time

ECE437, Fall 2016

(32)

Interrupt Driven Data Transfer



- Advantage:
 - User program progress is halted only during actual transfer
- Disadvantage, special hardware is needed to:
 - Cause an interrupt (I/O device)
 - Detect an interrupt (processor)
 - Save the proper states to resume after the interrupt (processor)

ECE437, Fall 2016

(33)

I/O Interrupts

- An I/O interrupt is just like the exceptions except:
 - An I/O interrupt is **asynchronous**
 - Further **information** needs to be conveyed
- An I/O interrupt is **asynchronous** with respect to instruction execution:
 - I/O interrupt is not associated with any instruction
 - I/O interrupt does not prevent any instruction from completion
 - You can **pick your own convenient point** to take an interrupt
- I/O interrupt is more complicated than exception:
 - Needs to convey the **identity of the device** generating the interrupt
 - Interrupt requests can have different urgencies:
 - Interrupt request **needs to be prioritized**


ECE437, Fall 2016

(34)

Direct Memory Access (DMA)

- Direct Memory Access (DMA):
 - External to the CPU
 - Act as a master on the bus
 - Transfer blocks of data to/from memory without CPU intervention

CPU sends a starting address, direction, and length count to DMAC. Then issues "start".



DMAC provides handshake signals for Peripheral Controller, and Memory Addresses and handshake signals for Memory.

ECE437, Fall 2016

(35)

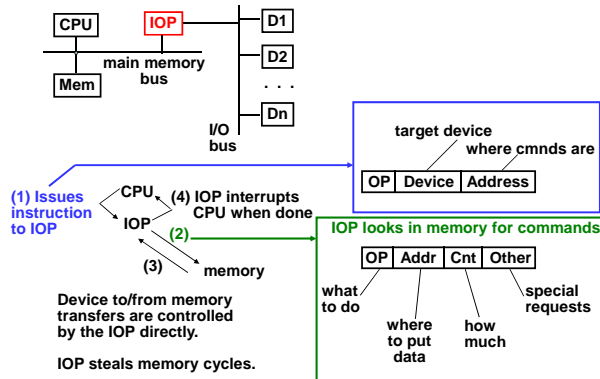
DMA interfacing

- DMA : virtual addresses or physical addresses?
- Challenge: crossing page boundaries
 - Virtual addresses
 - Translation provided by OS
 - Physical addresses
 - One page per transfer
 - OS chains the physical addresses
- No page faults in between
 - "Pin" pages

ECE437, Fall 2016

(36)


Delegating I/O Responsibility



ECE437, Fall 2016

(37)

I/O Coherence Problem



- Where is the cache w.r.t. I/O device (and/or I/O processor)?
 - In front of cache : Slows down CPU
 - Behind Cache : Cache coherence
- Applies to any I/O device that writes autonomously to memory
 - DMA, other
 - Multiprocessors also
- Coherence is a problem when both the following are involved:
 - Multiple writers AND
 - Replication

ECE437, Fall 2016

(38)

CPU involvement

- Low offloading to high offloading
 - Polling
 - CPU does everything including detecting events and subsequent I/O copying to/from memory
 - Interrupts
 - Someone else notifies CPU when processing needed, CPU does the needful
 - DMA
 - CPU describes what needs to be done, DMAC does the rest
 - IOP
 - CPU tells IOP where to find instructions on what needs to be done

ECE437, Fall 2016

(39)

OS Responsibilities

- The operating system acts as the interface between:
 - The I/O hardware and the program that requests I/O
 - Why?
 - Why not let user programs directly read/write to I/O devices?
- Three characteristics of the I/O systems:
 - The I/O system is shared by multiple program using the processor
 - I/O systems often use interrupts (external generated exceptions) to communicate information about I/O operations.
 - Interrupts must be handled by the OS because they cause a transfer to supervisor mode
 - The low-level control of an I/O device is complex:
 - Managing a set of concurrent events
 - The requirements for correct device control are very detailed

ECE437, Fall 2016

(40)

OS Responsibilities

- Provide **protection** to shared I/O resources
 - Guarantees that a user's program can only access the portions of an I/O device to which the user has rights
- Provides **abstraction** for accessing devices:
 - Supply routines that handle low-level device operation
- Handles the interrupts generated by I/O devices
- Provide **equitable access** to the shared I/O resources
 - All user programs must have equal access to the I/O resources
- **Schedule accesses** in order to enhance system throughput

ECE437, Fall 2016

(41)

Interfacing: Summary

- **Multiprogramming**
 - Program invokes syscall (i.e., invokes OS)
 - OS checks protection
 - OS runs device drivers
 - Suspends current process and switches process
 - I/O interrupt fielded by OS
 - OS completes I/O and wakes up suspended process (i.e. make it runnable)
 - Run next ready process
- **Ch I/O done!**

ECE437, Fall 2016

(42)

The end!

- ECE 437 done!
- Remember that you learnt the 4 key pillars of modern computer systems
 - Pipelining
 - Caches
 - Coherence
 - Virtual memory
 - And a few other cool things

ECE437, Fall 2016

(43)