

ECE437: Introduction to Digital Computer Design

Chapter 5c (virtual memory)

Fall 2016


Virtual memory : The real reason

- **Application's view of memory**
 - Large: ~4GB in Pentium (32b address)
 - Exclusive: Only program in memory
- **System view**
 - Smaller: 512MB-2GB
 - Multiple programs share memory
 - Run in a protected manner (memory is private **)
 - Address range is fixed (starts at 0x0)
 - Do not bring in entire program
 - Bring in relevant parts as needed
- **VM reconciles these conflicting "views"**
 - Not just the classical benefits of hierarchy
 - Illusion of
 - speed of expensive level
 - capacity and cost of cheaper level
 - Sharing, Protection, memory mapping

ECE437, Fall 2016

(105)

Virtual Memory



- Data movement between Disk and Main memory
- We know how layers of hierarchy interact
 - Cache and main memory
- Can we apply all the same techniques?
 - Similarities and differences

ECE437, Fall 2016

(106)

Virtual Memory


- **Similarities:**
 - Many-to-few mapping
 - larger address space to smaller memory
 - Used for data movement between layers of the memory hierarchy
- **Differences:**
 - Miss-penalty : 10-100 cycles vs. ~1 million cycle
 - Block size : 16-64 bytes vs. 4 KB - 8 KB
 - Full associativity (But no associative search!)
 - Software handling so more sophisticated schemes than hardware

Fundamental

ECE437, Fall 2016

(107)

VM Operation




- Programs use virtual address
- The data/code resides elsewhere (physical address)

ECE437, Fall 2016

(108)

VM Operation



- Assume 4KB pages
- 32-bit VA and 30-bit PA
- System responsible for translation
- E.g.
 - `lw $r2, 0xffffc004`
 - `0xffffc -> 0x20000` # VPN-> PPN translation
 - `PA = 0x20000004`

ECE437, Fall 2016

(109)

VM Terminology

- Blocksize ~ 4KB - 8KB+
 - Block called as Page
- Miss : Page-fault
 - Fetch from disk
- Derivative properties:
 - Fundamental constraint: high access latency

ECE437, Fall 2016

(110)

Back to the 4 questions

- Recall the 4 questions in caches
- Q1. Where does a block go?
 - Fully associative
 - Block offset and tag
 - NO INDEX
 - Why?
 - $AMAT = \text{hit time} + \text{miss-rate} * \text{miss-penalty}$
 - Miss-penalty = ~1 million cycles
 - Have to minimize miss-rate
 - full- associativity minimizes miss rate

ECE437, Fall 2016

(111)

Q2. Block Identification

- Q2. Block identification
 - Fully associative search??
 - 30-bit physical address (1GB)
 - 4 KB pages
 - Number of frames = $2^{30}/2^{12} = 2^{18} = 256K$
 - Compare 256 K frames in parallel? Whoa!
 - Reframe question:
 - Cache (previously): Is this VA in any given frame? => Parallel search
 - VM (now): Where is this VA? => Table lookup
 - Page table
 - Fully associativity via table lookup
 - » each table entry can point to ANY frame

ECE437, Fall 2016

(112)

Replacement

- Q3. Block replacement
 - LRU and/or LRU approximation (NRU with reference/use bits)
 - Sophisticated mechanisms possible (handling in OS software)
- Page-fault : Exception
 - Save instruction that causes fault
 - OS service the fault, i.e., brings in the relevant page from disk (VA-> disk address??)
 - OS knows service is slow; schedule other program
 - When disk access is complete
 - Restart at offending instruction

ECE437, Fall 2016

(113)

Write handling

- Q4. What happens on a write
 - Write-through or write-back?

ECE437, Fall 2016

(114)

Virtual Memory

- CPU and memory are shared by processes and VM ensures protection but the sharing is slightly different
- CPU is time shared - processes run for a while before being switched out by OS
 - OS saves their registers and PC in its data structures
- Memory is space shared - processes occupy pages and their page tables sit in memory
 - This occurs even when a process is switched out of the CPU

ECE437, Fall 2016

(115)

Exercise: Virtual Memory Design

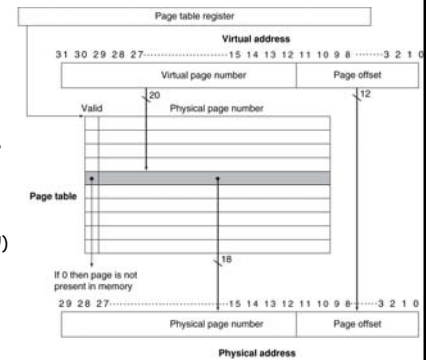
- Page size = 8KB
- Virtual Memory
 - $64b \rightarrow 2^{64} = 2^{34} \text{ GB!}$
- Physical memory = 8GB
- What is the size of the per-process page table?
- What happens on a context switch?

ECE437, Fall 2016

(116)

Page Table

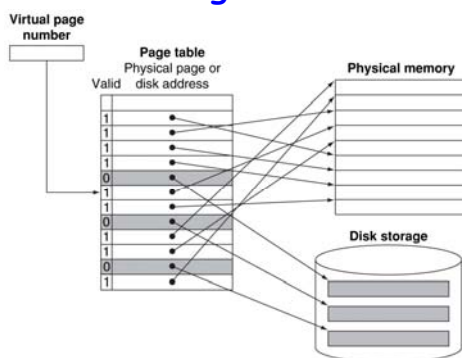
- Virtual page number
 - Tag in caches
- Physical address
 - Frame-number in caches
- PTBR (ONE per CPU)
 - Change value on context switch
 - Per-process page table



ECE437, Fall 2016

(117)

Page Table



ECE437, Fall 2016

(118)

Page Table Entries

- What does a page table entry contain
 - Physical page number (18 bits)
 - Access control (read/write permissions)
 - Valid bit (1 bit)
 - Misc
 - Use bit for replacement
 - Dirty bit for write-back
 - ~ 4 bytes (1 word)

ECE437, Fall 2016

(119)

Size of Page Table

- What is the size of the page table for a system with
 - 32 bit addresses
 - 1 GB physical memory
 - 4KB pagesize
- Virtual pages = $2^{32}/2^{12} = 2^{20}$
- Physical pages = $2^{30}/2^{12} = 2^{18}$
- PT size (per process) = (Entry size) * (# entries)
 - = 4 bytes * $2^{20} = 2^{22}$ bytes = 4 MB
- # of processes? ~60 on my WinXP Pro machine
 - 240 MB for page tables?
 - "Big government" : 25% consumed for administration!!

ECE437, Fall 2016

(120)

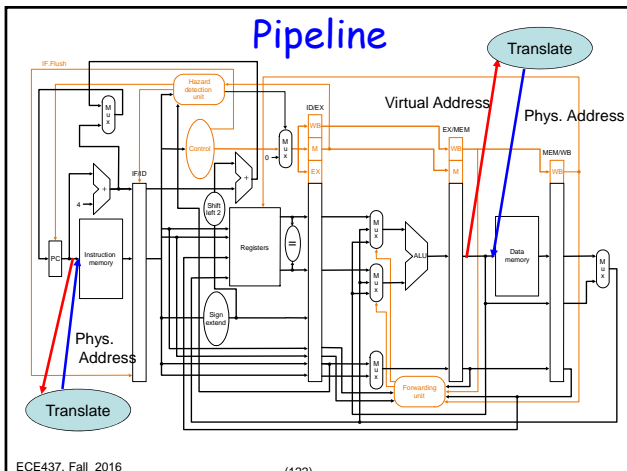
Size of Page Table

- Techniques to reduce Page Table Size
 - Gradual growing -
 - Not all processes use all of their VA space
 - only small fraction of all virtual pages actually used, especially in 64-bit machines
 - so VA space is sparse - not contiguous but has holes in it
 - Previous slide calculation assumes flat table for page table
 - what data structure would you use for a sparse space?
 - Inverted PT : entries per physical page
 - Hash table with as many buckets as physical pages
- Know that these options exist - explored in later courses

ECE437, Fall 2016

(121)

Pipeline



ECE437, Fall 2016

(122)

Translation overhead

- Where does table reside?
 - Main memory
 - Whoa! Wait a minute - this is 100% overhead!
 - Each memory reference now generates two memory references
 - lw \$r2, 0xffff0004
 - Access page table entry for 0xffff0, get PPN
 - Access PPN:004
 - But we busted our heads to put caches and 100 other tricks to avoid main memory access and now VM adds one WHOLE extra main memory access?

ECE437, Fall 2016

(123)

Translation overhead

- We want to minimize this overhead!
 - Page table entries can be cached like ordinary data
 - Why would it work?????
 - Tricky ***!!
 - Special cache for Page table entries

ECE437, Fall 2016

(124)

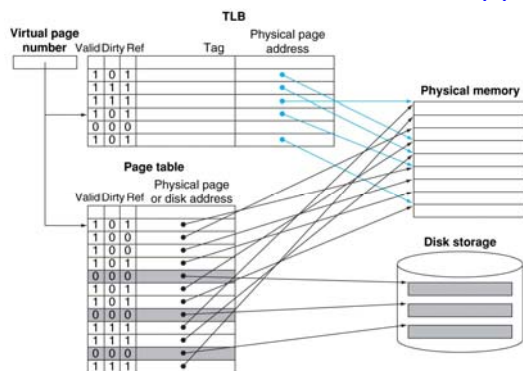
Faster Translation

- 100% overhead with VM system
- Eliminate memory access for translation
 - Caching
 - Translation lookaside buffer (TLB)
 - Also DTB in some literature
 - A cache of translations
 - 64-128 entries
 - Covers 256 KB ~512 KB
 - Organization
 - 64 entry fully associative
 - 256 entry, 16-way set associative
 - May be multi-level
- TLB is a cache for page table

ECE437, Fall 2016

(125)

Translation Lookaside Buffer



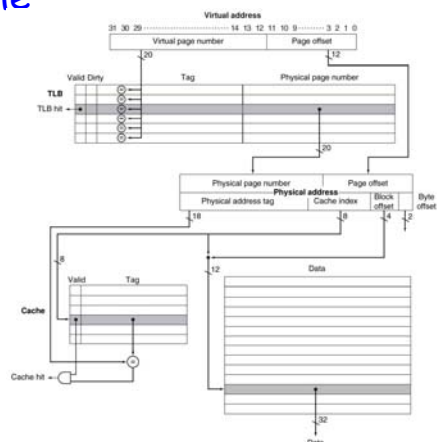
- TLB is a cache for page table

ECE437, Fall 2016

(126)

TLB+Cache

- Cache operation
 - With physical addresses
 - Translation on critical path



ECE437, Fall 2016

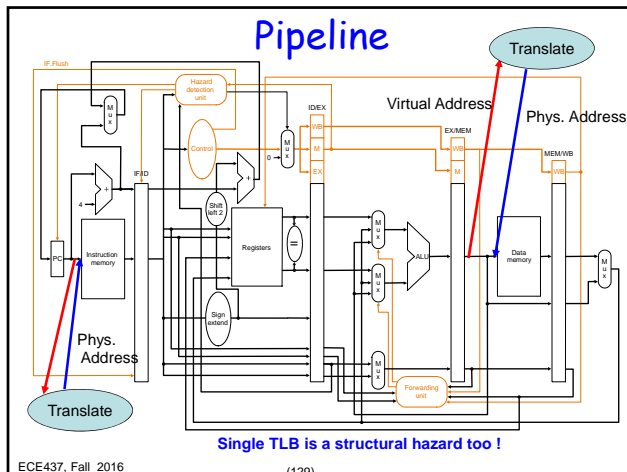
TLB design

- TLB's input address is virtual page number
- TLB is a cache for page table entries (PTEs) which already exploits spatial locality thru pages so TLB's block size = 1 PTE → TLB block offset is 0 bits
- TLB could be set assoc. or fully assoc.
 - VM is fully assoc. (VM miss penalty is huge) but TLB miss penalty is memory access (not huge) so TLB need not be fully assoc.
- So depending on #TLB entries and associativity, compute index & tag

ECE437, Fall 2016

(128)

Pipeline



ECE437, Fall 2016

(129)

TLB access latency

- TLB reduces access latency FROM cache (1-2 cycles) + memory for page table (300 cycles) TO cache (1-2 cycles) + TLB (~1-2 cycles) but still 100% overhead over cache
- One suggestion: use VA for L1, translate on L1 miss to use PA for L2
- Makes sense BUT has 2 problems due to modern OS

ECE437, Fall 2016

(130)

TLB access latency

- Problem1: two processes' VAs look alike
 - so context switch from process1 to process2 and process2 will get cache hits on process1's data!
 - Flush cache at context switch - many cache misses
 - Any other solution?
- Even if we solve this problem2 is harder

ECE437, Fall 2016

(131)

TLB access latency

- Problem2: Modern OS allow a strange thing: the SAME data may use DIFFERENT virtual addresses in different parts of SAME process but accesses to these VAs must access the same data
 - Such VAs are called aliases/synonyms
 - Provide key flexibility for applications
 - Key for DLLs, memory-map I/O,

ECE437, Fall 2016

(132)

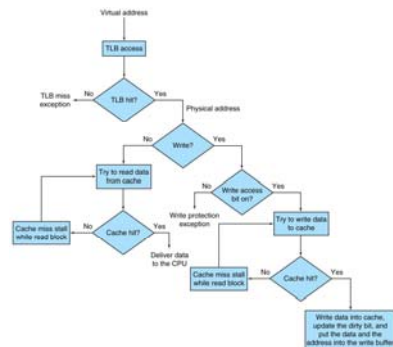
TLB access latency

- Using VA into L1 in the presence of synonyms cause this problem
 - Assume two aliases VA1 and VA2 where a store to VA2 and load from VA1, then VA1 and VA2 will go to different blocks in the cache so the load will not load the stored value!
- Real systems solve this problem so that TLB and L1 are accessed in parallel (come to 565)

ECE437, Fall 2016

(133)

Putting it all together



- Memory access flowchart See Fig 5.25 for full chart

ECE437, Fall 2016

(134)

Full Hierarchy design

- 32b VA, 30 bit PA (i.e. 1GB Phys. Mem)
 - 4 KB pages
 - TLB
 - 256 entry, 4-way associative
 - Cache
 - 32KB, direct mapped, 32B blocks

ECE437, Fall 2016

(135)

Summary

- **4Q on VM**
 - Placement : fully associative
 - Identification : Page Table lookup
 - Replacement : LRU / LRU-approx
 - Writes : Writeback
- **4Q on TLB**
 - Placement: small and fully associative, larger and set-associative
 - Identification: Associative search (CAM)
 - Replacement : NRU, random**
 - Writes : ?? Writes to TLB??

ECE437, Fall 2016

(136)

VM Miscellanea

- On a context switch:
 - Change contents of PTBR for appropriate page table
 - What do we do with TLB contents?
 - What's the problem? All VAs look alike!
 - How will TLB distinguish between VAs of previous and new process
 - Flush all entries
 - Simple, but inefficient
 - Associate Process ID with address
 - Flush required only when process IDs are reused

ECE437, Fall 2016

(137)

VM Miscellanea

- Memory efficiency of page tables
 - Limit register
 - Only region between PTBR and Limit is valid
 - Grow as needed
 - Segmented
 - Two page tables and two limit registers (Stack and Heap)
 - Inverted Page table
 - Hashing to map VA to a number within PA range
 - Hash 32-bit VA to 28 bit PA
 - Lookup complications
 - Collision
 - Multilevel page tables
 - Paging page tables

ECE437, Fall 2016

(138)

Real stuff

Characteristic	ARM Cortex-A8	Intel Core i7
Virtual address	32 bits	48 bits
Physical address	32 bits	44 bits
Page size	Variable: 4, 16, 64 KiB, 1, 16 MiB	Variable: 4 KiB, 2/4 MiB
TLB organization	1 TLB for instructions and 1 TLB for data Both TLBs are fully associative, with 32 entries, round robin replacement TLB misses handled in hardware	1 TLB for instructions and 1 TLB for data per core Both L1 TLBs are four-way set associative, LRU replacement L1 I-TLB has 128 entries for small pages, 7 per thread for large pages L1 D-TLB has 64 entries for small pages, 32 for large pages The L2 TLB is four-way set associative, LRU replacement The L2 TLB has 512 entries TLB misses handled in hardware

- Ch 5 done!

ECE437, Fall 2016

(139)