


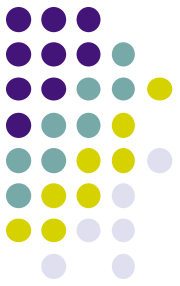
Processes

ECE469, Jan 17

Yiying Zhang

- 1. program process difference
 - 17. process define
 - 18. process state
 - 19. pcb process control block
- 21. Process Creation/Termination
 - 40. Context Switch



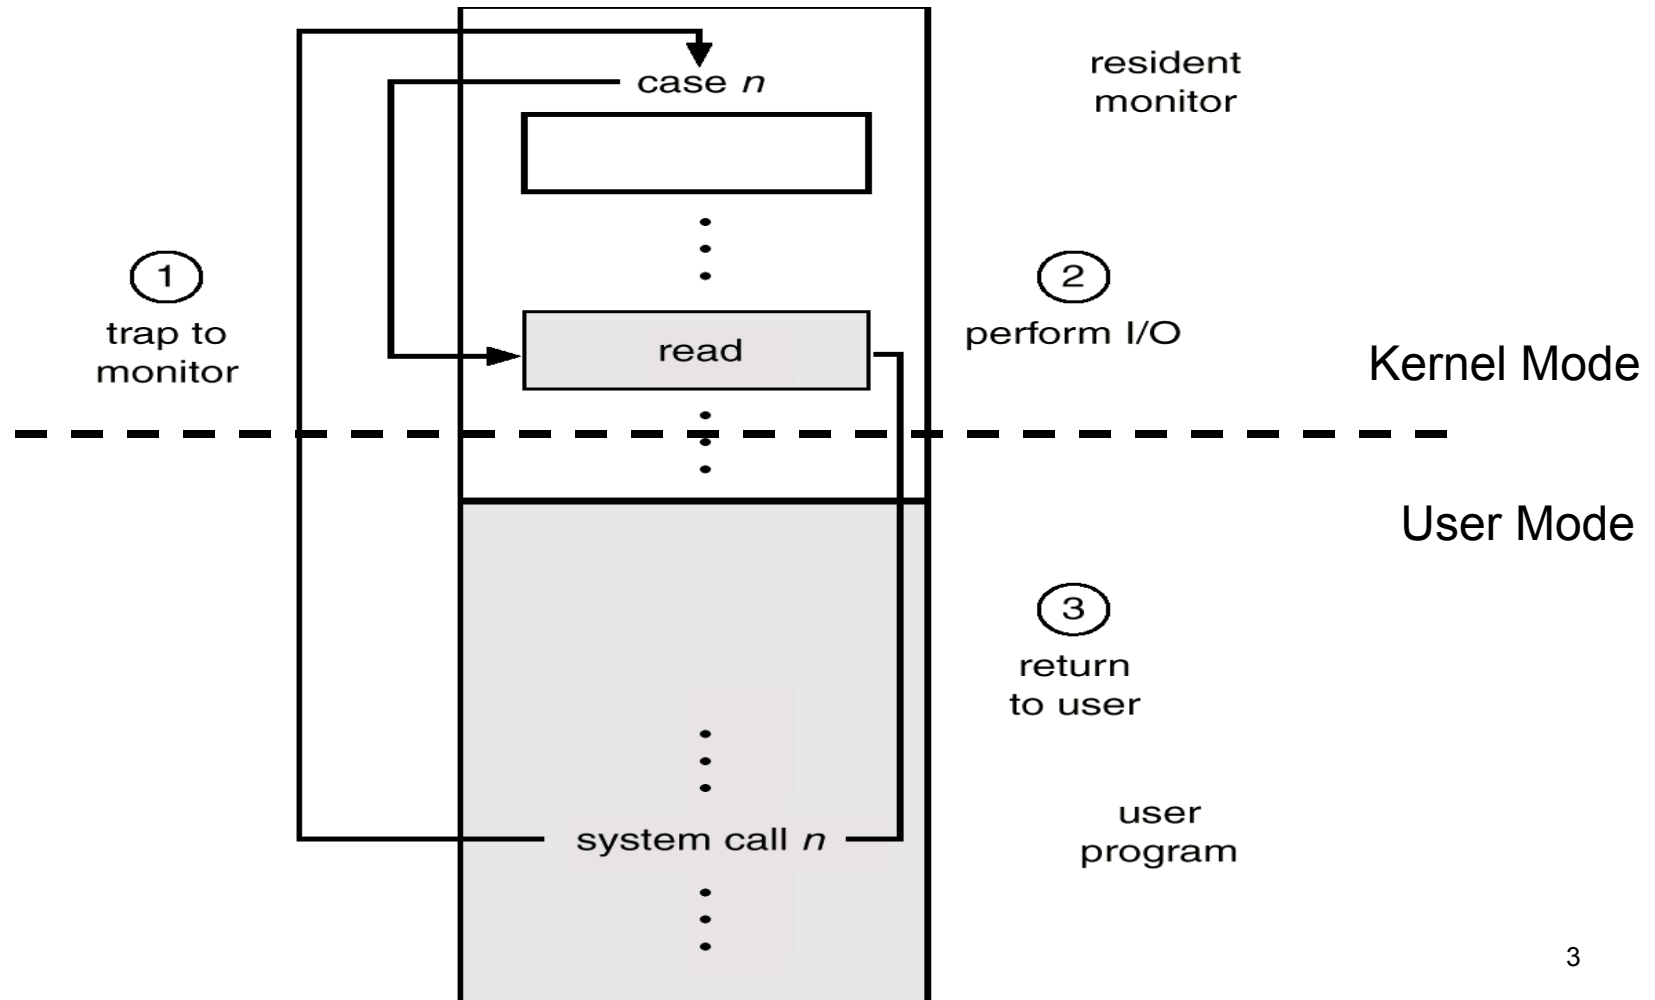


Review

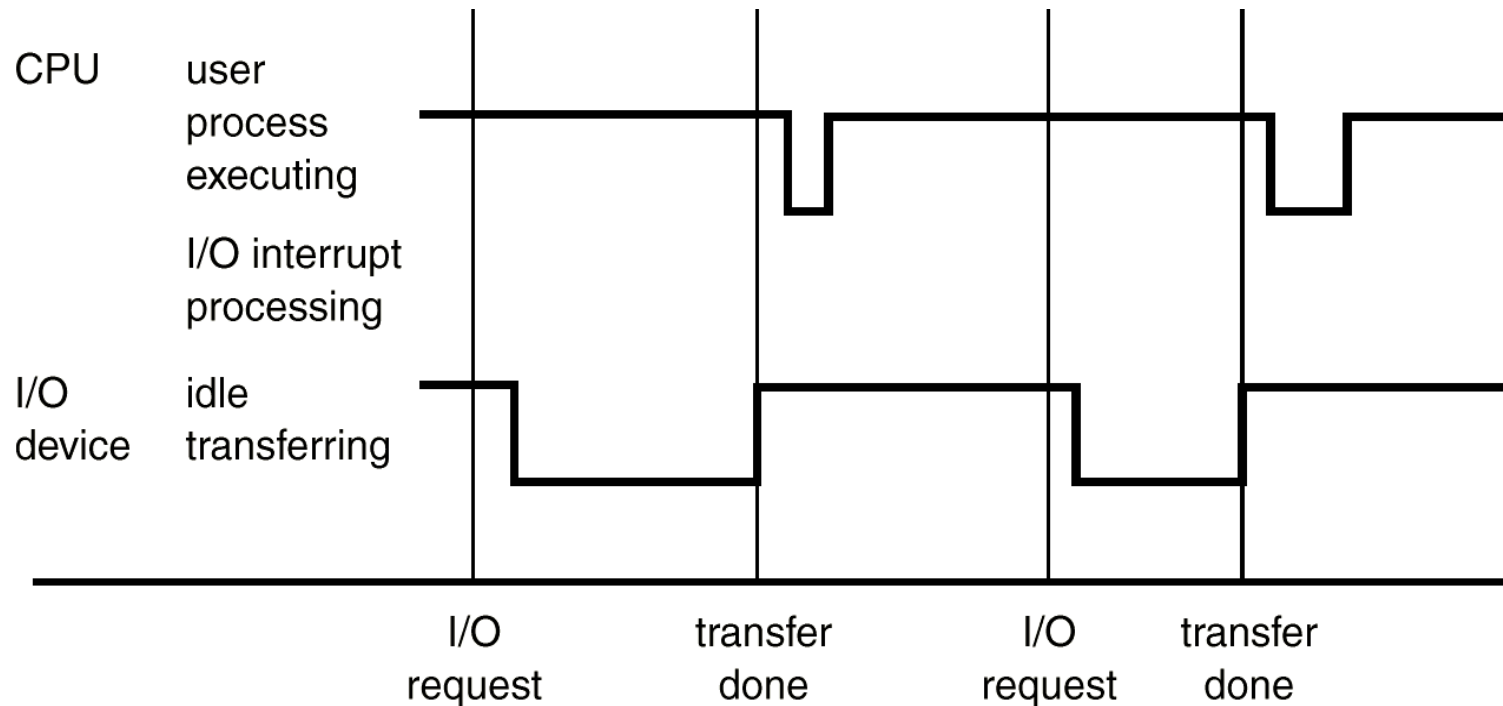
- OS is a resource manager
- OS presents an extended machine
- OS is a “giant interrupt handler”
- System calls
- Interrupt




Use of a system call to perform I/O



Interrupt time line for a single process doing I/O



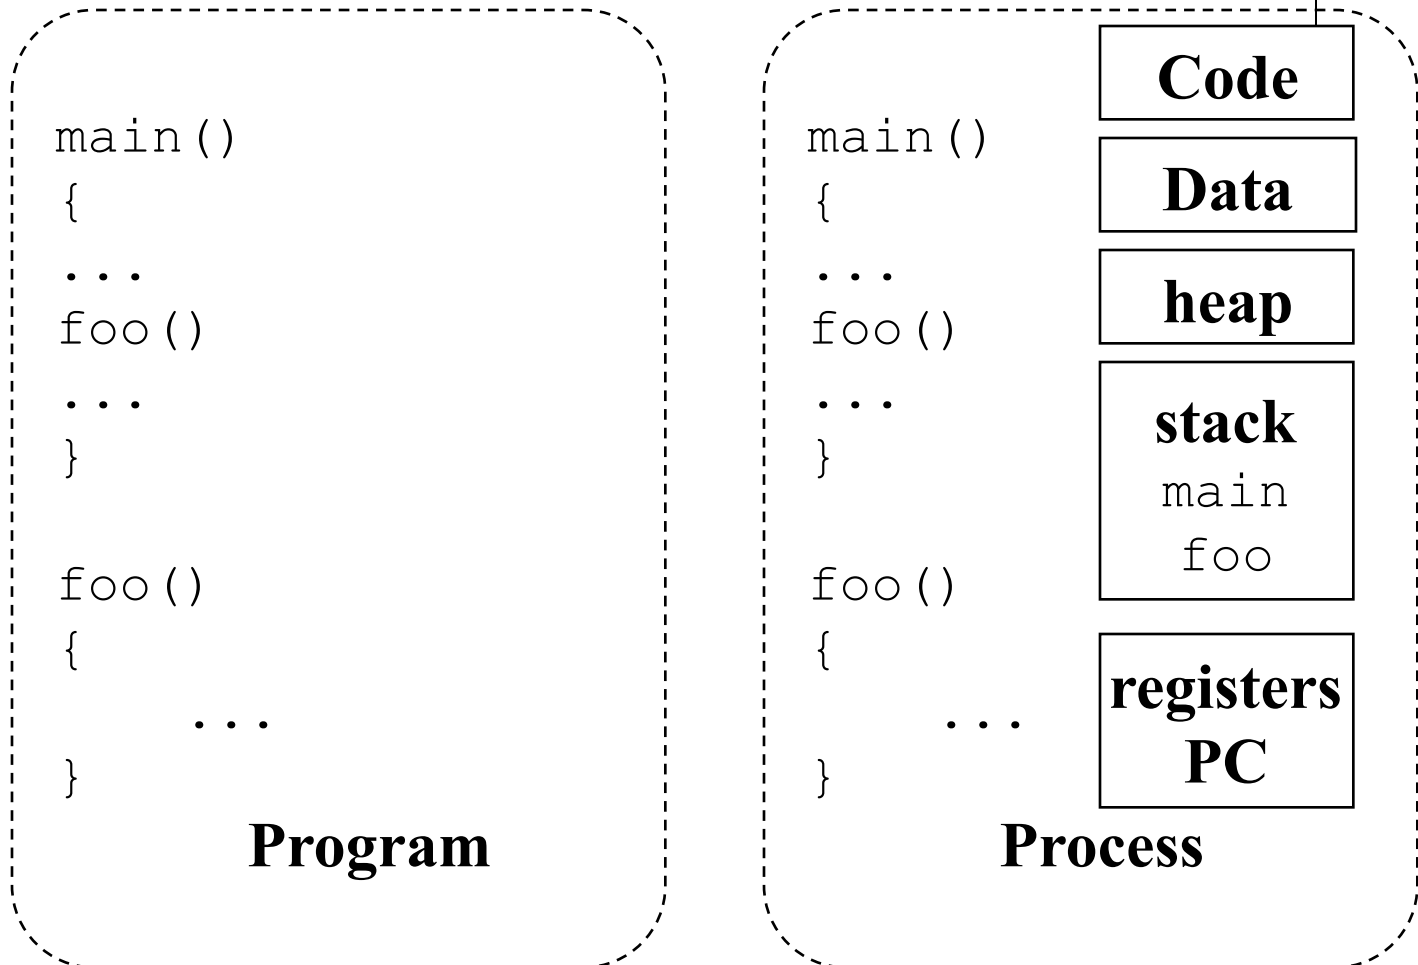


Users, Programs, Processes

- Users have accounts on the system
- Users launch programs
 - Can many users launch the same program?
 - Can one user launch many instances of the same program?

→ A process is an “instance” of a program

Program vs. Process



Program as A Process Collection, or a Program Collection

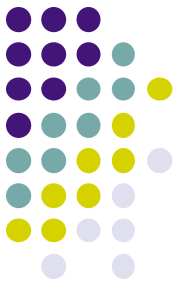


- Firefox (output of “ps ux | grep firefox”)
 - [~]\$ ps ux | grep firefox
 - yiying 529 0.0 0.2 326756 9836 ? SI 2012 1:43 /usr/lib64/xulrunner-2/plugin-container /usr/lib64/mozilla/plugins-wrapped/nswrapper_32_64.libflashplayer.so -greomni /usr/lib64/xulrunner-2/omni.ja -appomni /usr/lib64/firefox/omni.ja 25752 plugin
 - yiying 25752 3.9 6.6 1240212 250328 ? SI 2012 8039:58 /usr/lib64/firefox/firefox
- gcc (via “gcc -pipe -v”)
 - /usr/libexec/cpp |
 - /usr/libexec/cc1 |
 - /usr/libexec/as, followed by
 - /usr/libexec/elf/ld



Users, Programs, Processes

- Users have accounts on the system
 - Users launch programs
 - Can many users launch same program?
 - Can one user launch many instances of the same program?
- A process is an “instance” of a program
- A “program” can be a set of programs
- A “program” can be a set of processes



So What Is A Process? (1)

- It's **one instance** of a “program”
- Any relationship between two instances?



What Does This Program Do?

```
int myval;
```

```
int main(int argc, char *argv[])  
{  
    myval = atoi(argv[1]);  
    printf("myval is %d, loc 0x%lx\n",  
           myval, (long) &myval);  
}
```



Instances of Programs

- The address of the static variable was always the same!
- The values were different!
- Implications:
 - Do instances think they're using the same address?
 - Are they seeing each other?
- **Conclusion: addresses are not absolute!**
- What are the benefits?
 - Compiler/linker/loader do not have to be concerned
 - Allows address space to be bigger than memory ?

So What Is A Process? (2)



- It is **one instance** of a “program”
- It is **separate** from other instances

Sequential execution of each process




- Assuming **single-threaded** program
- No concurrency inside a process
- Everything happens sequentially
- Often with interleaved CPU/IO operations



Concurrent Processes

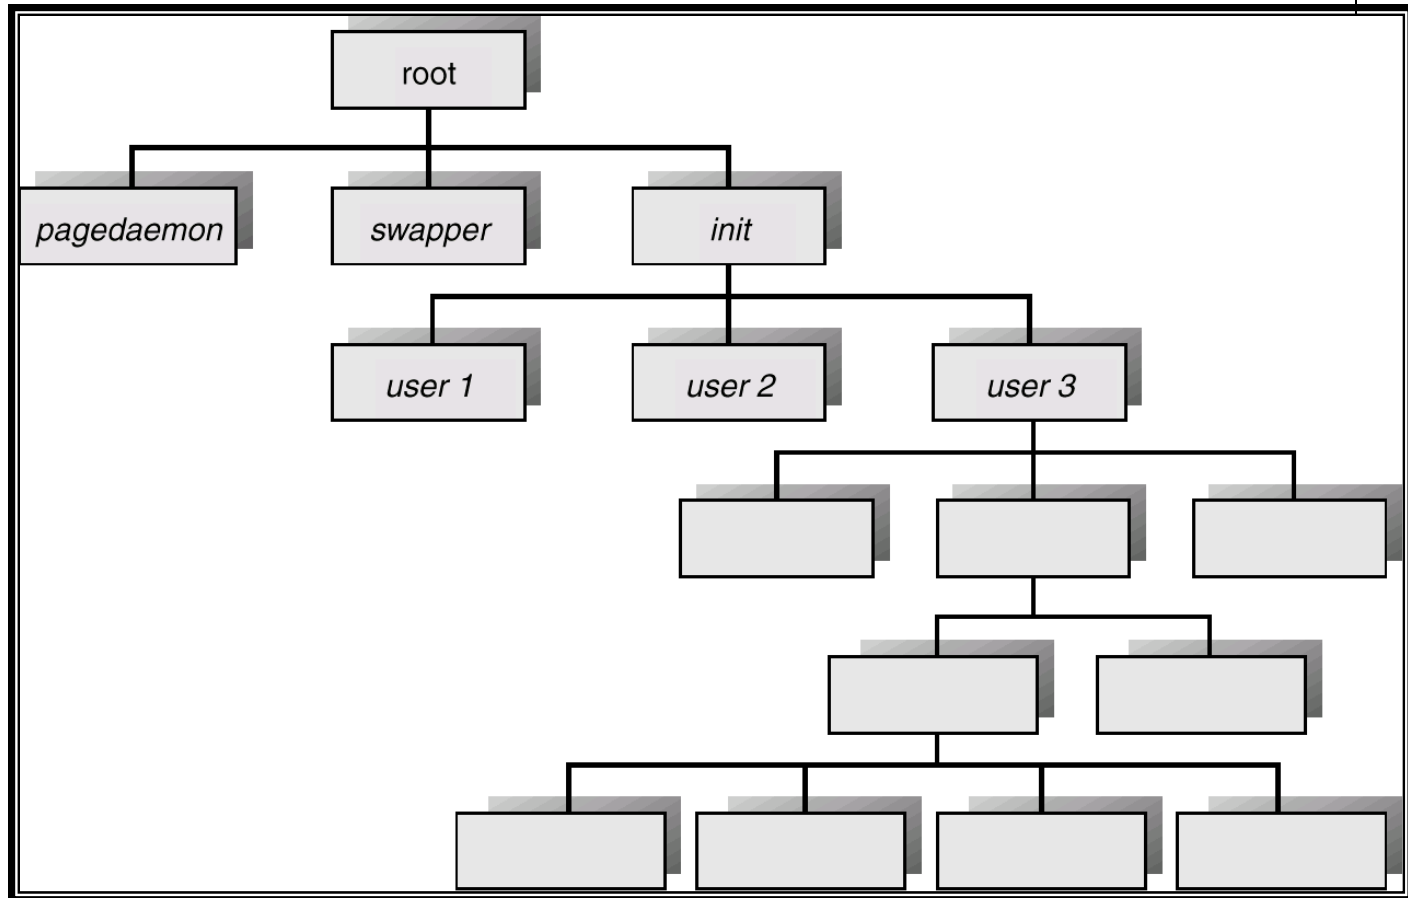

- Processes in a system can execute concurrently (multitasking)
- Motivations for allowing concurrent execution
 - Physical resource sharing (system utilization)
 - Computational speedup – with several CPUs
 - Modularity (firefox)
 - Convenience (desktop: firefox, xclock, emacs, xterm)

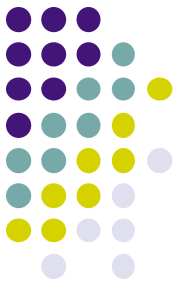


Process Creation

- Who created process cpp?
- Who created cc1?
- Who created gcc?
- Who created shell?

Processes Tree on a UNIX System



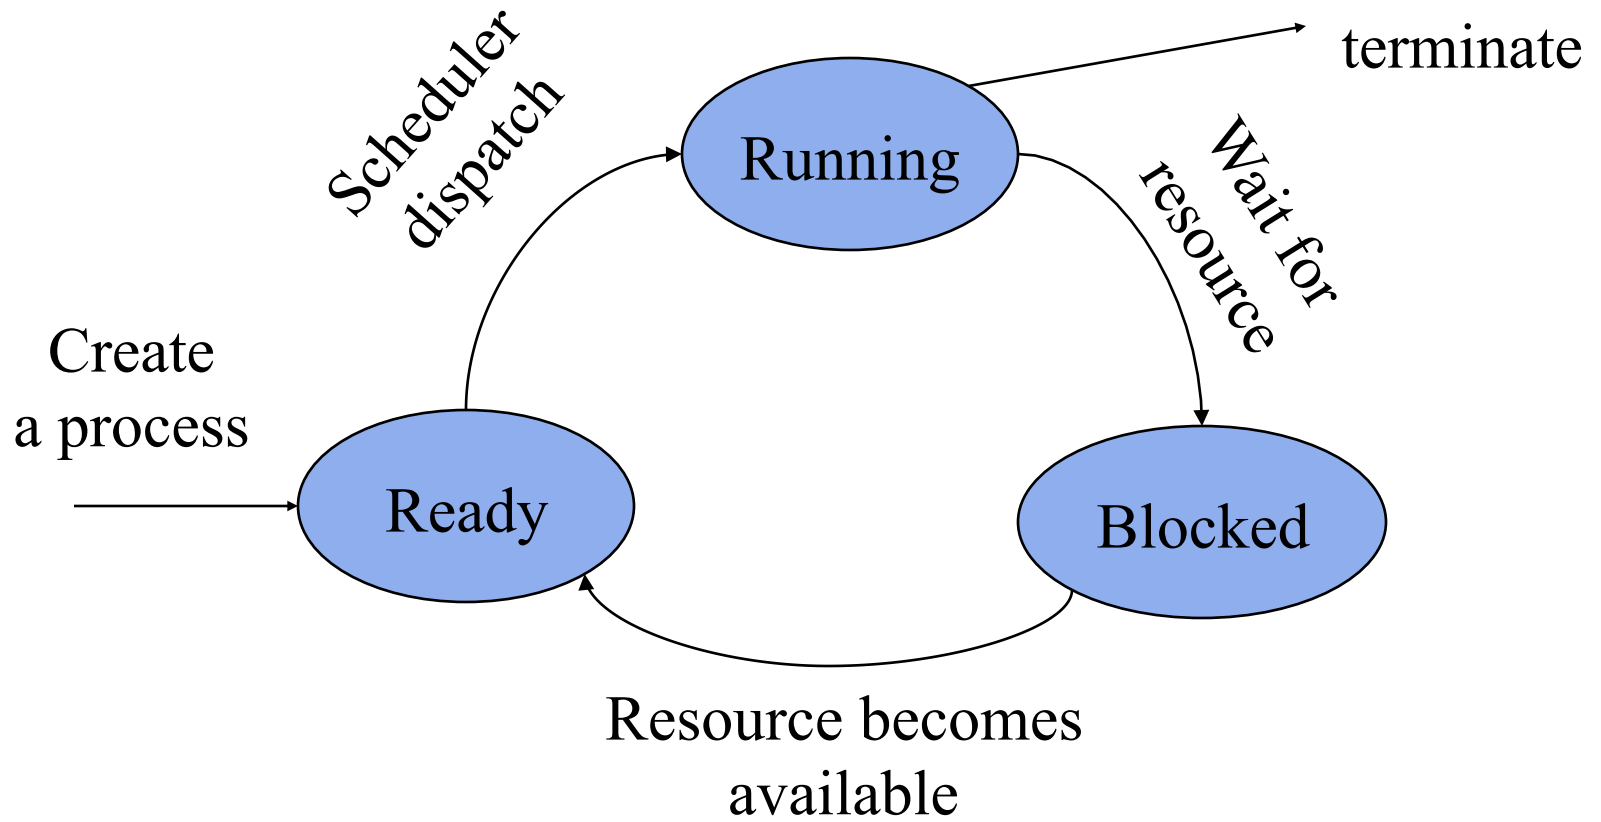



So What Is A Process? (3)

- It's **one instance** of a “program”
- It's **separate** from other instances
- It can start (“launch”) other processes
- It can be launched by them

Life cycle of a process:

Process State Transition

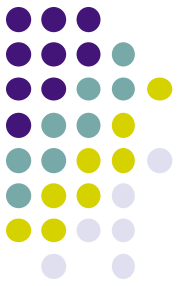


Kernel data structure:

Process Control Block (Process Table)



- Process management info
 - State (ready, running, blocked)
 - PC & Registers, parents, etc
 - CPU scheduling info (priorities, etc.)
- Memory management info
 - Segments, page table, stats, etc
- I/O and file management
 - Communication ports, directories, file descriptors, etc.¹⁹



Process ID

- Every process has an ID – process ID
- Does a program know its process ID?
- When a program is running, how does the process know its ID?

Process Creation/Termination




- Who should *actually* create/terminate processes? (bring program to life)
 - Who manages processes?
- But user process decides when and what
- So what should OS provide user process?
- What should the semantics of the interface be?

Process Creation – interface design options

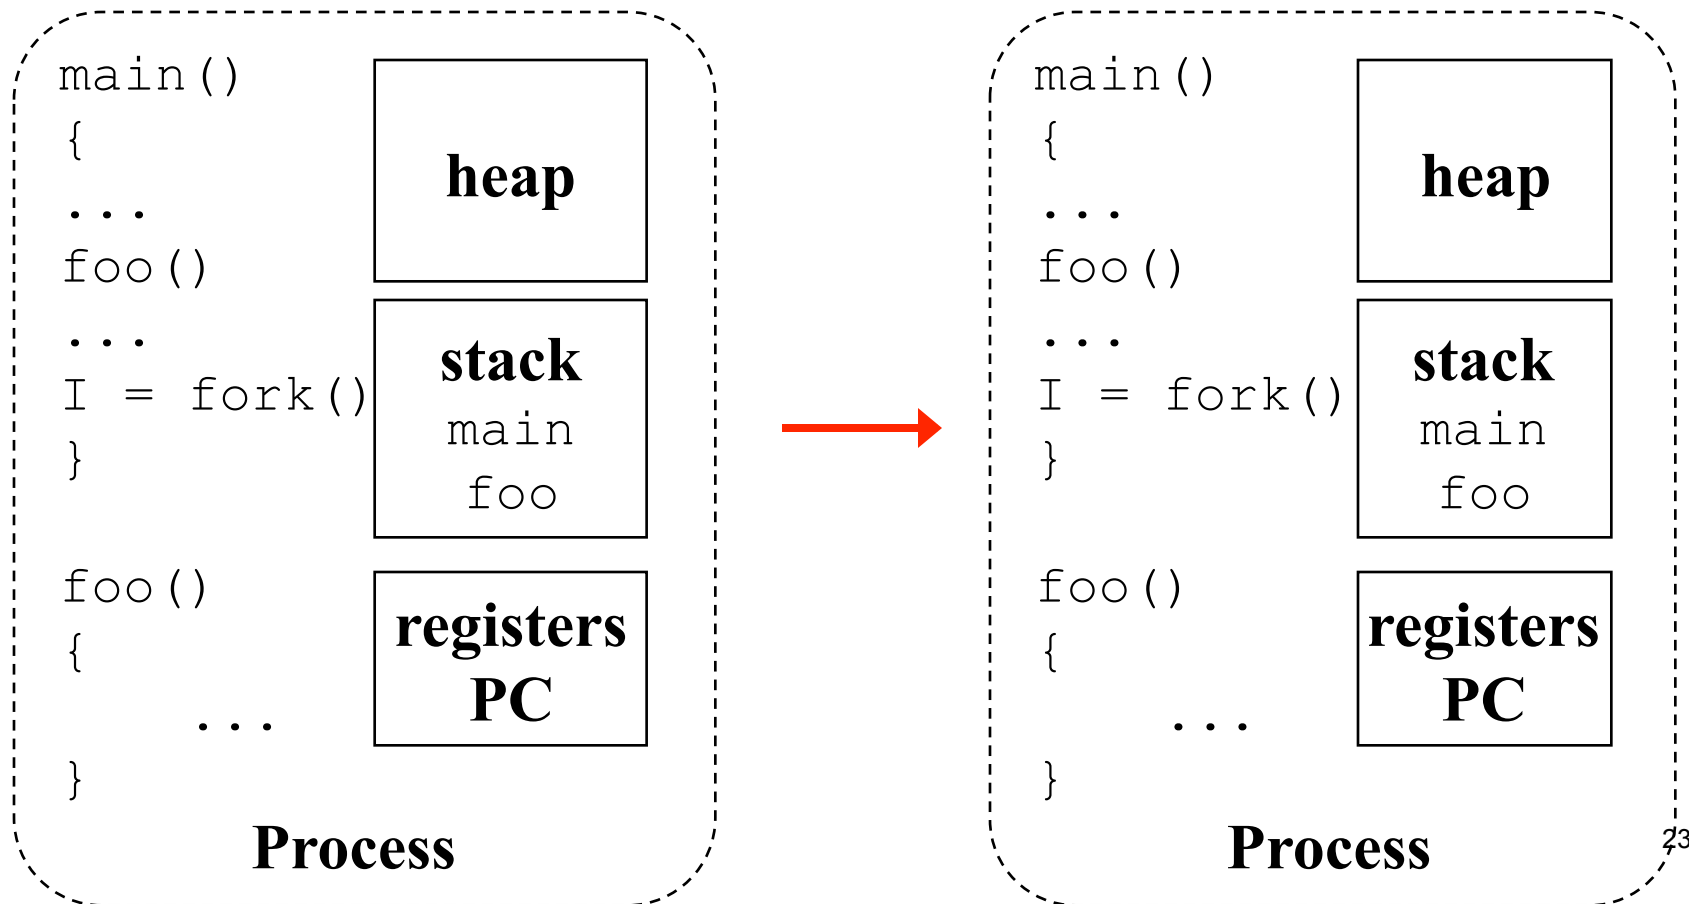


- The system call creates a child process
- Execution possibilities?
 - Parent and child execute concurrently
 - Parent waits till child finishes
- Address space possibilities?
 - Child duplicate of parent (code and data)
 - Child has a new program loaded into it

Process Creation -- UNIX interfaces (1)



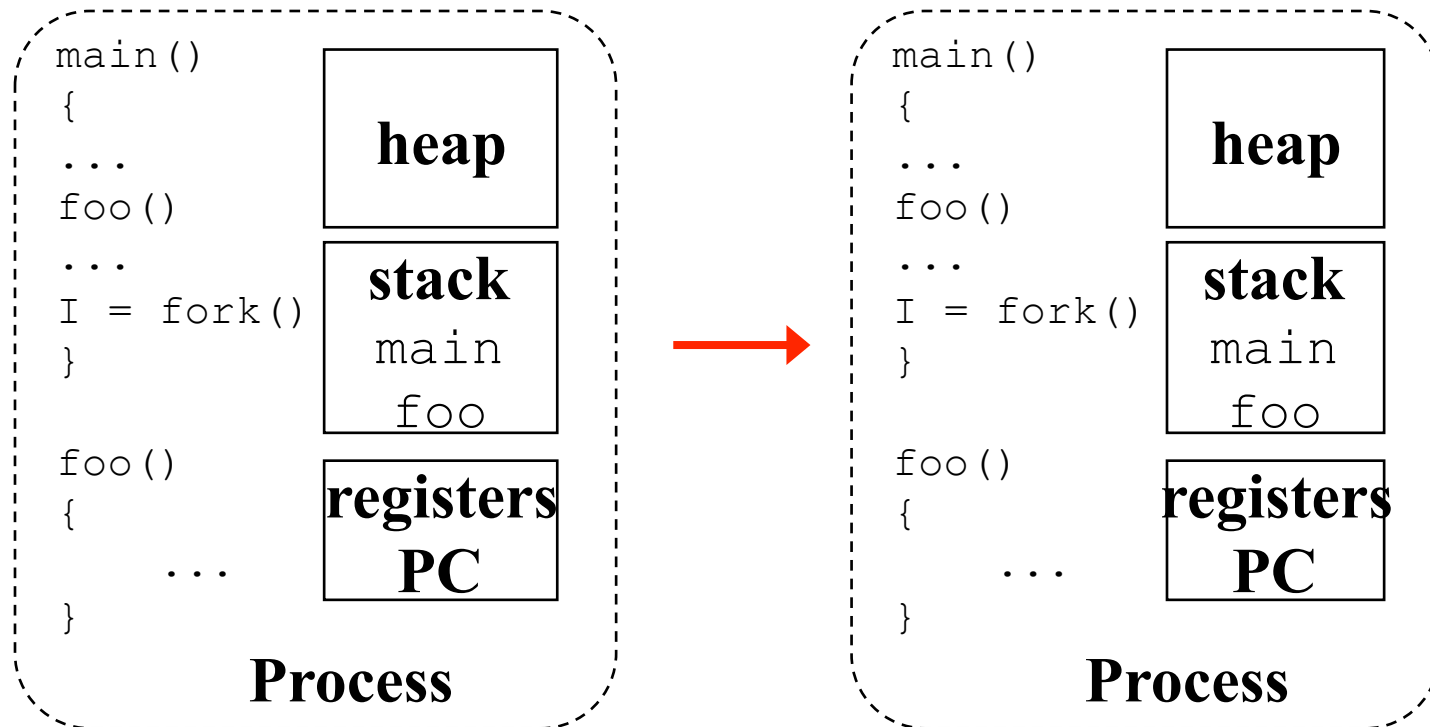
- *fork()* system call creates a duplicate of the original process




Process Creation -- UNIX interfaces (2)



- *fork()* system call creates a duplicate of the original process
 - What is the major benefit?
 - How to disambiguate who is who?



C program forking a new process



```
#include <stdio.h>
```

```
void main()
```

```
{
```

```
    int pid; int was = 3;
```

```
    pid = fork(); /* fork another process */
```

```
    if (pid < 0) { /* error occurred */      exit(-1); }
```

```
    else if (pid == 0) { /* child process */
```

```
        printf("child: was = %d\n", was);
```

```
    else { /* pid > 0; parent process ; pid is child process's */
```

```
        printf("parent: child process id = %d\n", pid);
```

```
        wait(NULL);  exit(0);
```

```
    }
```

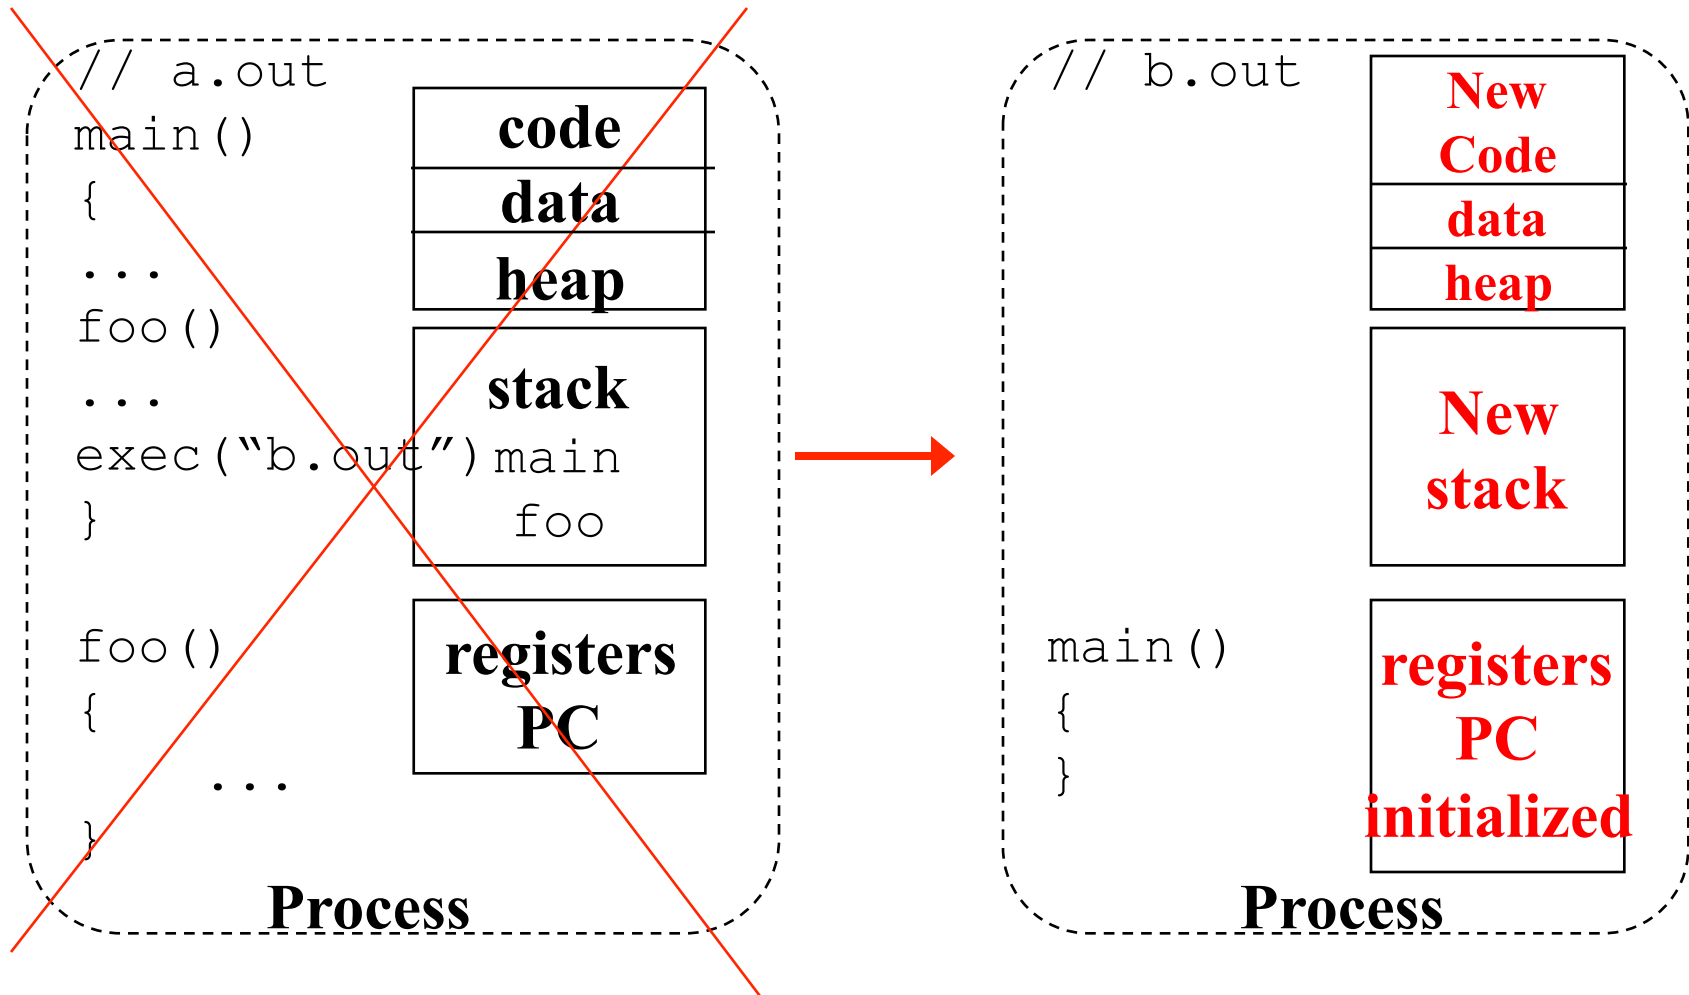
```
}
```

Process Creation -- UNIX interfaces (3)



- *fork()* system call creates a duplicate of the original process
 - What is the major benefit?
 - How to disambiguate who is who?
- *exec()* system call used after a *fork* to replace the process' code/address space with a new program
 - Important: BOTH code and data, i.e., the whole address space is replaced !

exec("b.out")



Afterwards, only one thing about the process was kept, which is?

Process Creation – UNIX interfaces (4)



UNIX has 4 system calls:

- fork – create a copy of this process
 - Clone would have been a better name!
- exec – replace this process with this program
- wait – wait for child process to finish
- kill – (potentially) end a running process



C program forking a new process

```
#include <stdio.h>
void main()
{
    int pid;      int was = 3;
    pid = fork(); /* fork another process */

    if (pid == 0) { /* child process */
        sleep(2); printf("child: was = %d\n", was);
        execlp("/bin/ls", "ls", NULL);}
    else { /* pid > 0; parent process */
        was = 4;
        printf("parent: child process id = %d; was=%d\n", pid, was);
        wait(NULL);  exit(0);
    }
}
```

C program forking a new process



```
#include <stdio.h>
```

```
void main()
```

```
{
```

```
    int pid;        int was = 3;
```

```
    pid = fork(); /* fork another process */
```

```
    if (pid == 0) { /* child process */
```

```
        sleep(2); was=9; printf("child: was = %d\n", was);
```

```
        execlp("/bin/ls", "ls", NULL);    was = 10;
```

```
        printf("It's me, your child was = %d\n", was);}
```

```
    else { /* pid > 0; parent process */
```

```
        was = 4;
```

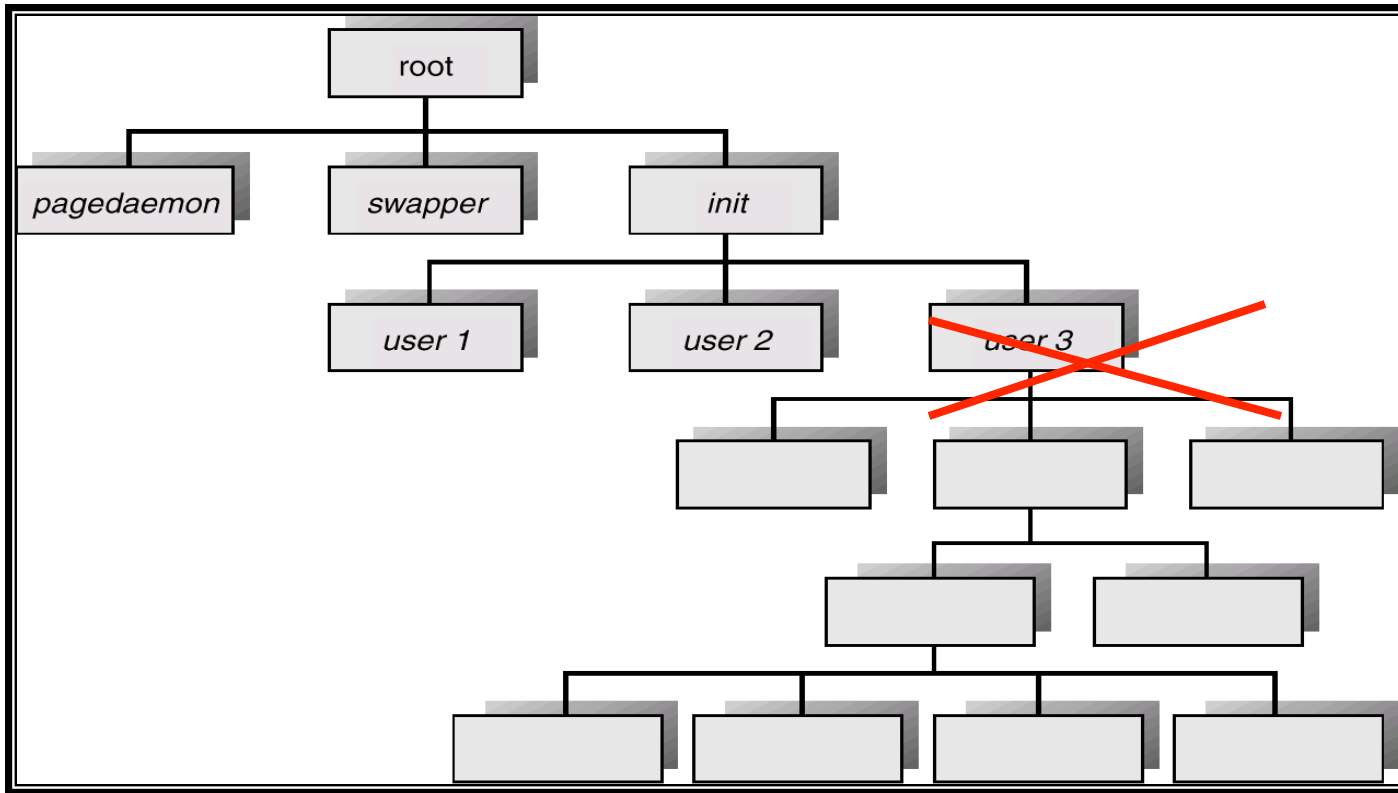
```
        printf("parent: child process id %d was=%d\n", pid, was);
```

```
        wait(NULL);    exit(0);
```

```
    }
```

```
}
```

Processes Tree on a UNIX System



What happens when a parent process disappears?
all child processes are killed by the OS, or
all child processes reset init as their parent

Break



- 100 passengers are boarding an airplane with 100 seats. Everyone has a ticket with his seat number. These 100 passengers boards the airplane in order. However, the first passenger lost his ticket so he just take a random seat. For any subsequent passenger, he either sits on his own seat or, if the seat is taken, he takes a random empty seat. What's the probability that the last passenger would sit on his own seat? There is a very simple explanation for the result.

Multiprogramming needs CPU scheduling



- Without any hardware support, what can the OS do to a running process?

System calls may trigger Scheduler



- Block – wait on some event/resource
 - Network packet arrival (e.g., `recv()`)
 - Keyboard, mouse input (e.g., `getchar()`)
 - Disk activity completion (e.g., `read()`)
- Yield – give up running for now

Non-Preemptive Scheduler



- A non-preemptive scheduler: a scheduler that is only invoked by explicit block/yield calls, or terminations
 - Only method when there is no timer!

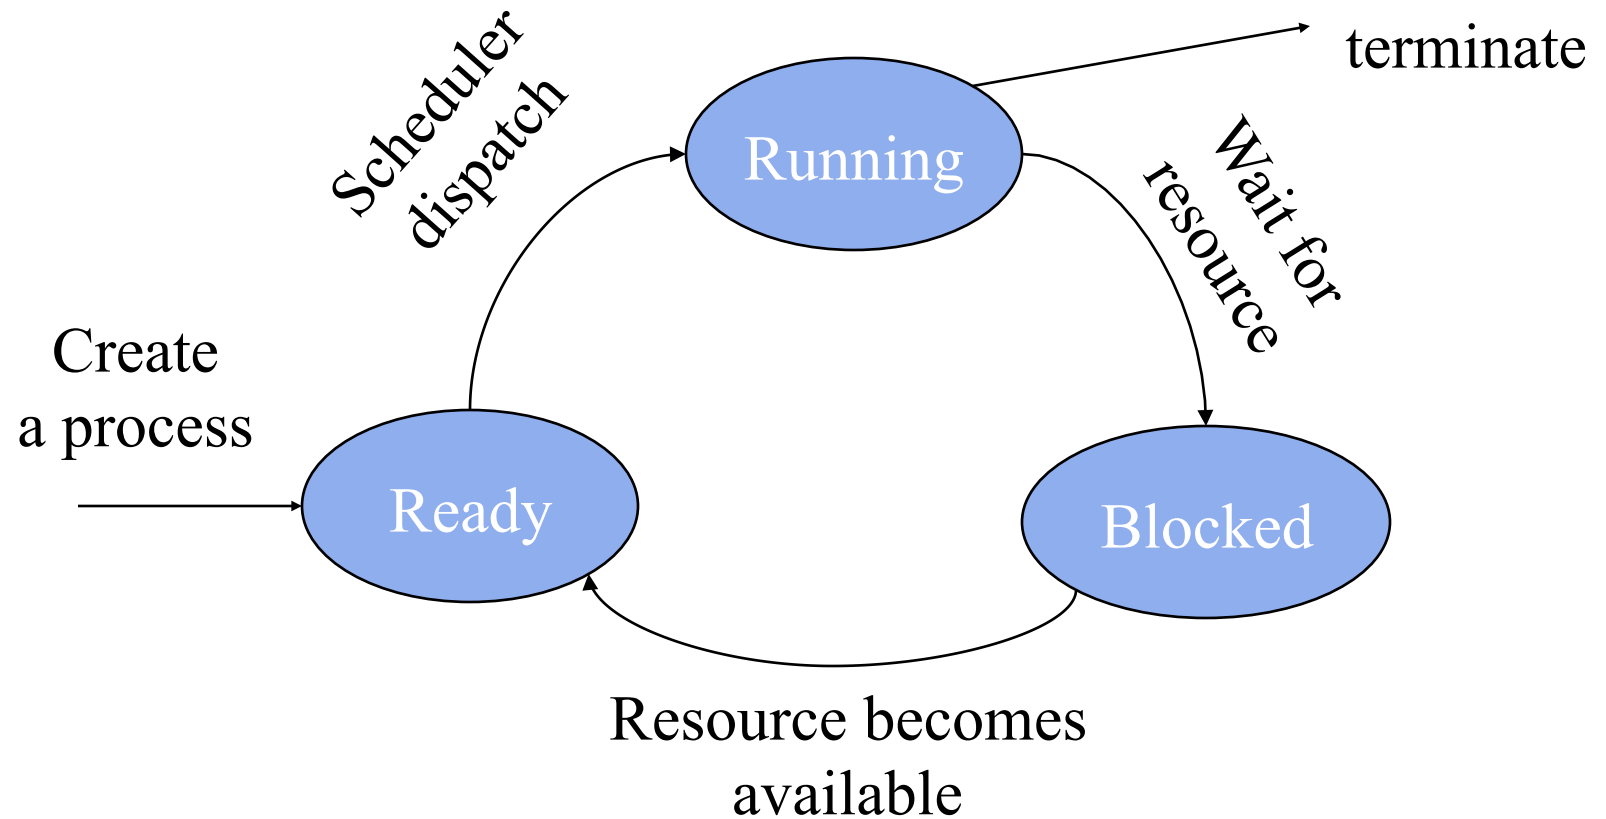
- The simplest form

Scheduler:

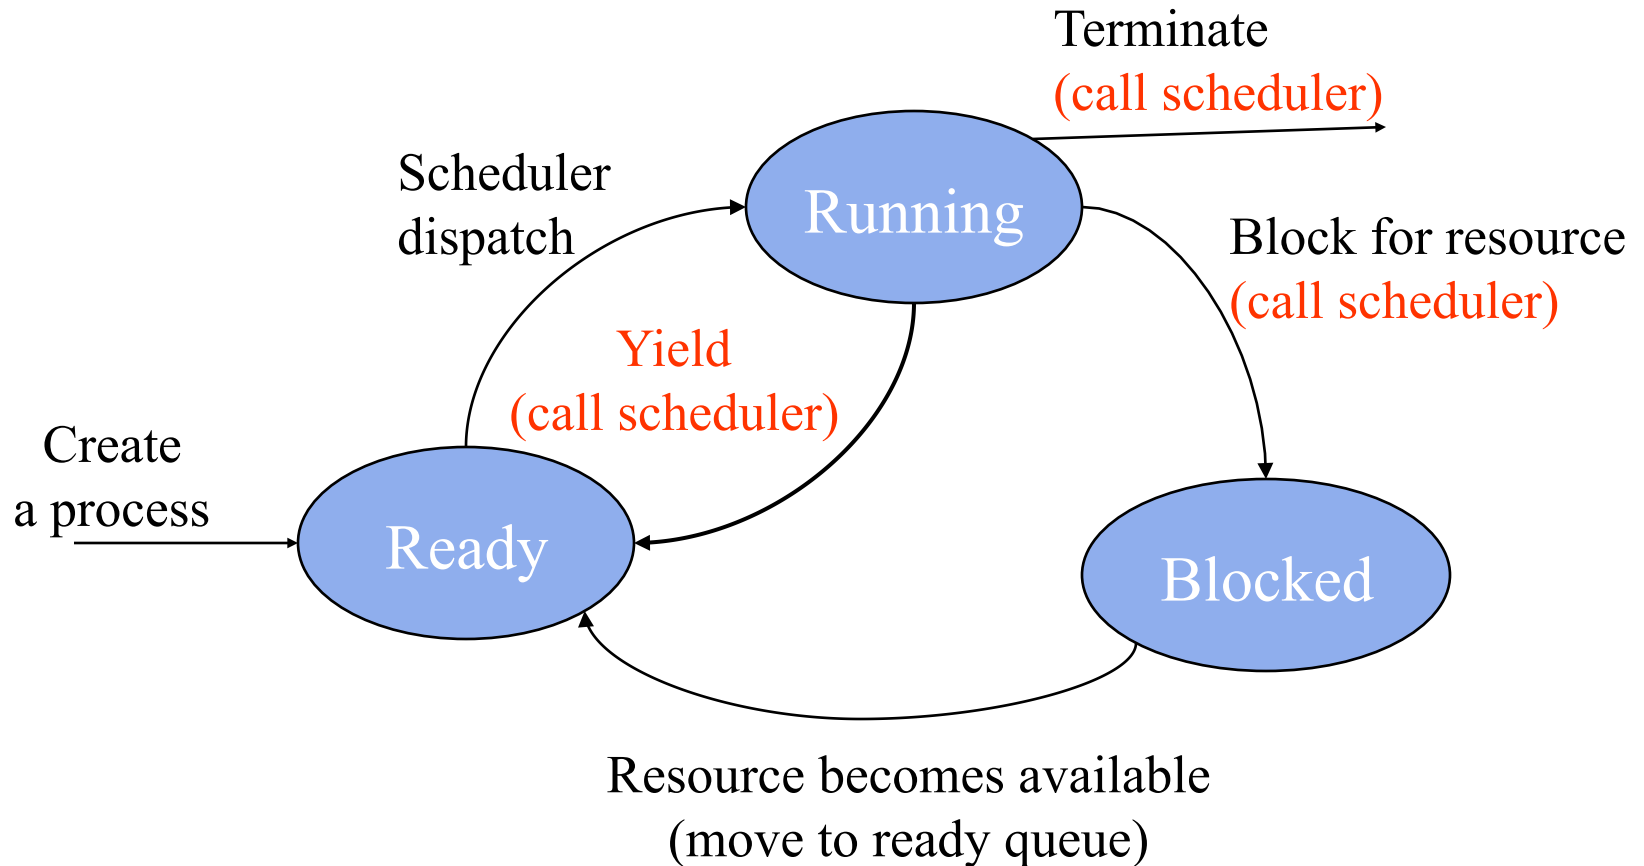
save current process state (into PCB)
choose next process to run
dispatch (load PCB and run)

- Used in Windows 3.1, Mac OS

Our Friend -- the Transition Diagram



Process State Transition of Non-Preemptive Scheduling



How does a process terminate itself?



Context Switch

- Definition:
switching the CPU to run another process, which involves (1) saving the state of the old process and (2) loading the state of the new process
- What state?

[lec3] Program vs. Process



```
main()  
{  
  ...  
  foo()  
  ...  
}  
  
foo()  
{  
  ...  
}
```

Program

```
main()  
{  
  ...  
  foo()  
  ...  
}  
  
foo()  
{  
  ...  
}
```

Process

Code

Data

heap

stack

main
foo

registers
PC



Context Switch


- Definition:
switching the CPU to run another process, which involves (1) saving the state of the old process and (2) loading the state of the new process
- What state?
 - What about L1/L2 cache content?



Context Switch

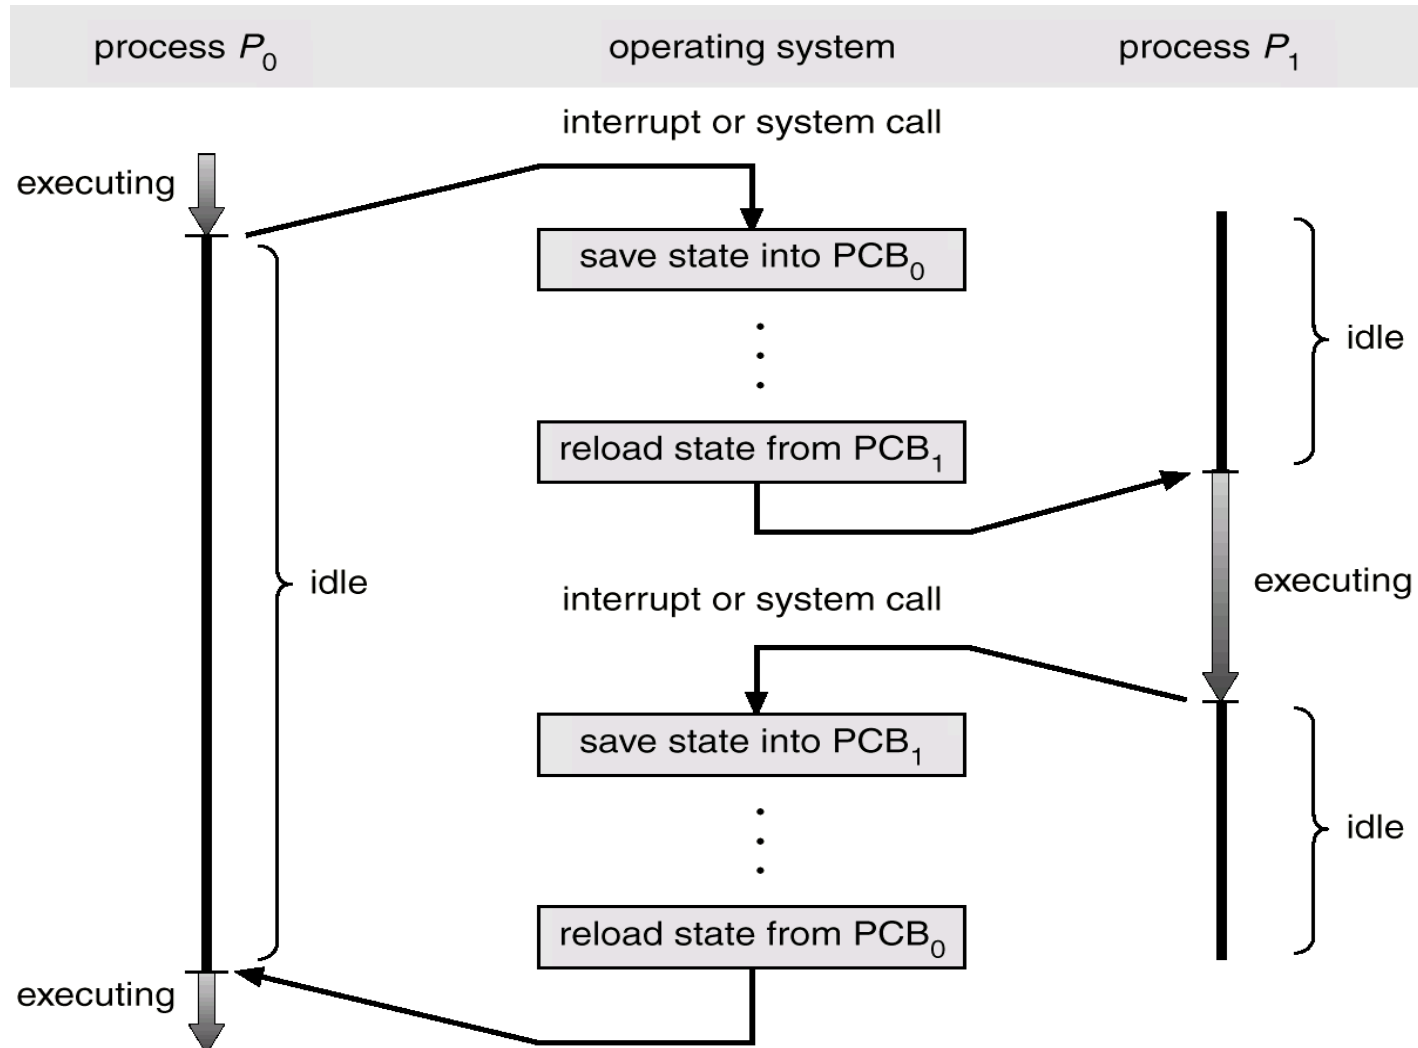
- Definition:
switching the CPU to another process, which involves saving the state of the old process and load the state of the new process
- What state?
- Where to store them?

[lec3] Process Control Block (Process Table)



- Process management info
 - State (ready, running, blocked)
 - PC & Registers, parents, etc
 - CPU scheduling info (priorities, etc.)
- Memory management info
 - Segments, page table, stats, etc
- I/O and file management
 - Communication ports, directories, file descriptors, etc.

Context Switch

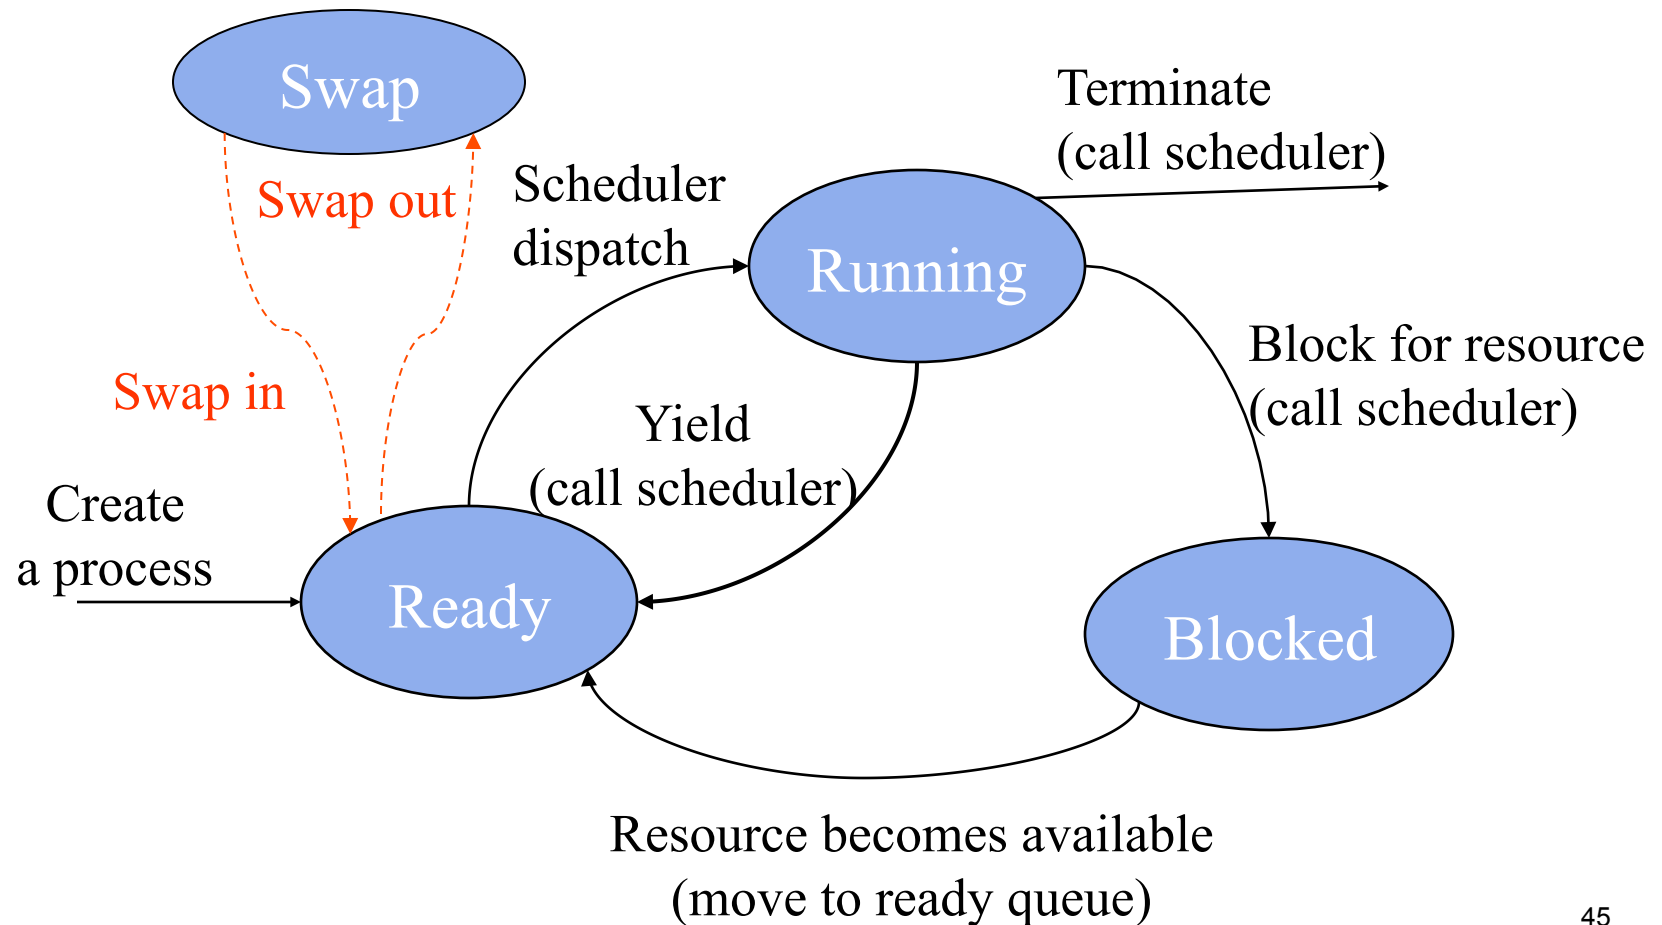


Physical Memory & Multiprogramming



- Want to run many programs
 - Programs need memory to run
 - Memory is a scarce resource
-
- What happens when
 $M(a) + M(b) + M(c) > \text{physical mem?}$


Add Job Swapping to State Transition Diagram





Summary

- What is a process?
- PCB
 - Process management info
 - Memory management info
 - I/O, file management info
- Process state transition
- Concurrent processes
- Process creation and termination
- Multiprogramming and context switch



Reading Assignment

- Dinosaur Chapter 3 by Thursday
- Comet Chapters 4,5,6