

# Semaphore Implementation

ECE469, Jan 26

Yiying Zhang


5. Readers-Writers problem

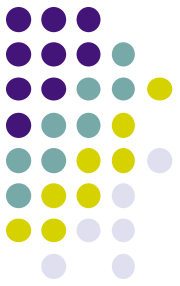
19. Semaphore implementation

32. Using test-and-set(TAS) for mutual exclusion

39. Implementing Locks + Condition Variables

47. Load-Linked and Store-Conditional





# Reading

- Dinosaur Ch 6
- Comet Ch 28
- Lab 2 starts

# Big Picture So Far



- Getting synchronization right is hard **concurrent process**
  - Even some of your esteemed faculty members (e.g., Yiying) can get it wrong.
- How to pick between locks, semaphores, conditional variables, monitors???
- Locks are very simple for many cases.
  - Issues: Maybe not be the most efficient solution
  - For example, can't allow multiple readers but one writer inside a standard lock.
- **Condition variables allow threads to sleep while holding a lock**
  - Just be sure you understand whether they use Mesa or Hoare semantics!
- Semaphores provide pretty general functionality
  - But also make it really easy to botch things up.
- Monitors are a higher-level semantics for using locks and condition variables.

# Semaphore classic examples



1. Producer-consumer problem

2. Readers-writers problem

3. Dining philosophers problem

# [lec5] Semaphore classic problem 2:



## Readers-Writers problem


- A data object is shared among multiple processes
- Allow concurrent reads (but no writes)
- Only allow exclusive writes (no other writes or reads)

# [lec5] Readers-Writers problem (Solution 1)



- Constraints:
  - Writers can only proceed if there are no readers/writers
  - Readers can proceed only if there are no writers
  - use a single semaphore BlockWrite
  - To keep track of how many are reading
  - use a shared variable
  - Only one process manipulates state variable at once
  - use semaphore Mutex
- Initialization:
  - semaphore BlockWrite = 1; // used to allow ONE writer or MANY readers
  - semaphore Mutex = 1; // binary semaphore (basic lock)
  - int Readers = 0; // count of readers reading in critical section

# Writer



`P(BlockWrite); // wait to lock the shared resource for a writer`

`< Do the Writing >`

`V(BlockWrite);`



# Reader


```
P(Mutex);  
Readers++;  
if (Readers == 1) // first reader acquire write lock  
    P(BlockWrite);  
V(Mutex);
```

< Do the Reading >


```
P(Mutex);  
Readers--;  
if (Readers == 0) // last (only) reader releases write lock  
    V(BlockWrite);  
V(Mutex);
```



# What will happen in different scenarios?



1. The first reader blocks if there is a writer; any other readers who try to enter block on mutex.
  2. The last reader to exit signals a waiting writer.
  3. When a writer exits, if there is both a reader and writer waiting, which goes next depends on the scheduler.
  4. If a writer exits and a reader goes next, then all readers that are waiting will fall through.
  5. Does this solution guarantee all threads will make progress?
- Writes can starve
  - => Read preference
- once reader get lock, and many reader following, writer cannot get the lock



# [lec4] What is a good solution

- Only one process inside a critical section
- Processes outside of critical section should not block other processes
- No one waits forever
- No assumption about CPU speeds
- Works for multiprocessors

# Readers-Writers problem (Solution 2)



- How do we let reads yield to writes?
  - `int Readers = 0, Writers = 0; // count of readers and writers in critical section`
  - `semaphore BlockWrite = 1; // used to allow ONE writer or MANY readers`
  - `semaphore BlockRead = 1; // used to block readers`
  - `semaphore RMutex = 1; // binary semaphore for Readers`
  - `Semaphore WMutex = 1; // binary semaphore for Writers`

# Write



```
P(WMutex);  
Writers++;  
if (Writers == 1) // block readers  
    P(BlockRead);  
V(Wmutex);  
  
P(BlockWrite); // ensures only one writer  
< Do the Writing >  
V(BlockWrite);  
  
P(WMutex);  
Writers--;  
if (Writers == 0) // enable readers  
    V(BlockRead);  
V(WMutex);
```

# Reader



```
P(BlockRead); // at most one reader can go before a pending write
```

```
P(RMutex);
```

```
Readers++;
```

```
if (Readers == 1) // first reader acquire write lock
```

```
    P(BlockWrite);
```

```
V(RMutex);
```

```
V(BlockRead);
```

● Any problem?

< Do the Reading >


```
P(RMutex);
```

```
Readers--;
```

```
if (Readers == 0) // last (only) reader releases write lock
```

```
    V(BlockWrite);
```

```
V(RMutex);
```



# Problem of solution 2

- Reader starvation
- Is there a solution that's fair to both reads and writes?
  - Take-home puzzle problem (not graded)

# Reader-writer problem a possible fair solution?



- An idea: use a FIFO queue for all readers and writers


# Synchronization Primitives provided by OS (language/compiler)



- Lock
  - Alone is not powerful enough
- Semaphore (incl. binary semaphore)
  - binary semaphore alone not enough
- Lock and condition variable
- Monitor (hide lock, still use condition variables)



# What are the inner workings of these powerful synchronization tools



# Semaphore



```
wait(S) {  
    while (S<=0) ;  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

- Semantics:
  - S “counts” the number of “resources”
  - wait(S) signifies “consuming” one if there is any, otherwise wait!
  - signal(S) signifies “producing” one

# Semaphore implementation

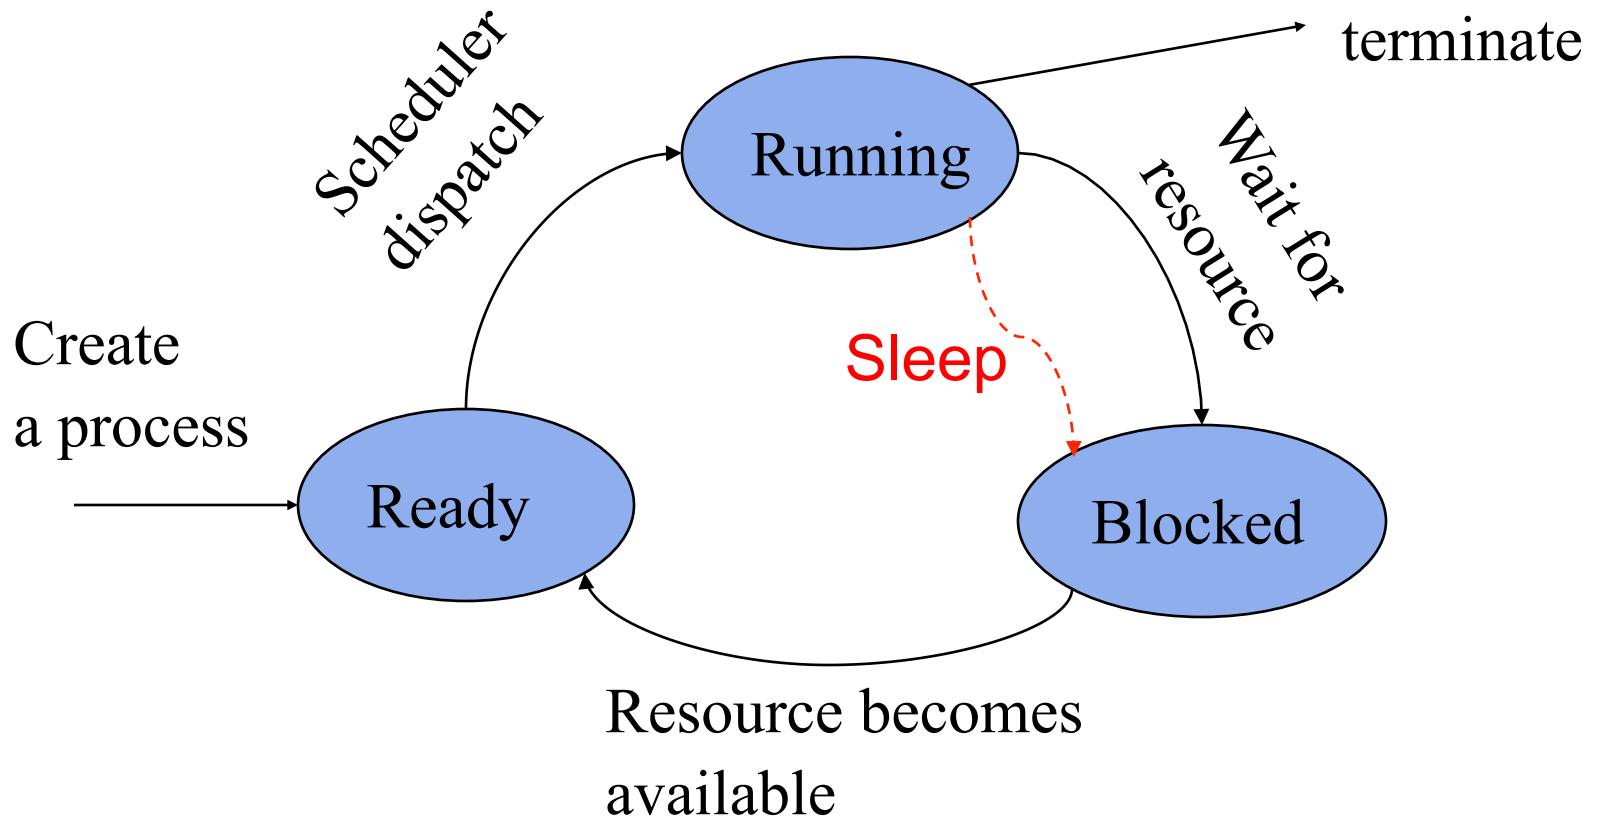


```
wait(S) {  
    while (S<=0) ;  
    S--;  
}
```

```
signal(S) {  
    S++;  
}
```

- Can they be implemented in the user space?
    - An intuitive argument?
  - No existing hardware implements them directly
    - Scheduling/queuing cannot be easily done in HW
- ➔ Semaphore must be done in OS, typically with low-level synchronization support from hardware

# [lec 3] Process State Transition



# Uniprocessor solution



```
typedef struct {  
    int count;  
    queue q;  
} semaphore;
```

```
void wait(semaphore s)  
{
```

```
    if (s->count > 0) {  
        s->count --;
```

```
    // sleep -> do not need while loop
```

```
        return;
```

```
    }
```

```
    add(s->q, current_process);
```

```
    sleep();
```

```
}
```

```
void signal(semaphore s)  
{
```

```
    s->count ++;
```

```
    if (!isEmpty(s->q)) {
```

```
        process = removeFirst(s->q);
```

```
        wakeup(process);
```

```
        /* put process on Ready Q */
```

```
    }
```

```
}
```

# Challenge



- Need to make primitives atomic!
- What Mutex support do we have from HW on uniprocessor?
  - On uniprocessor, reads and writes are atomic

# Uniprocessor solution



- What can cause the few lines to be not atomic?
- What causes context switches?
- Recall -- only way the OS dispatcher regains control is via **interrupts** (incl. explicit requests, i.e. syscalls)
  - E.g. typing -> keyboard interrupt -> handler -> kernel -> user process

# Uniprocessor solution: disable interrupts!




```
void wait(semaphore s)
{
    disable interrupts;
    if (s->count > 0) {
        s->count --;
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    enable interrupts;
    // implying sleep(); // re-dispatch
}
```

```
void signal(semaphore s)
{
    disable interrupts;
    s->count ++;
    if (!isEmpty(s->q)) {
        process = removeFirst(s->q);
        wakeup(process);
        // put process on Ready Q
    }
    enable interrupts;
}
```



# Generic shared-memory multiprocessor



# What about multiprocessors?



- True concurrency – simultaneity!
  - Cache coherence in HW (fairly complicated)
  - For simplicity: a read sees most recent write
- Is turning off interrupts enough?
  - Turning off interrupt on a processor does not prevent other processors from accessing shared memory
  - Turning off interrupt on all other processors?
- Use atomic read/write and busy waiting?

# A Multiprocessor solution?



```
void wait(semaphore s)
{
    disable interrupts;
    while (s->count == 0);

    s->count --;

    enable interrupts;
}
```

```
void signal(semaphore s)
{
    disable interrupts;
    s->count ++;
    enable interrupts;
}
```



# Does this solution work?

- Did we achieve atomicity of the while loop (busy wait) and the decreasing of `s->count`?  
no because `++` is more than one instruction
- Busy waiting on local processor is not efficient
- What if the signaler is scheduled on the same core as wait?
- What about concurrent memory access (`s->count`)?

# What about multiprocessors? (cont)



- Cannot just turn off interrupts
  - Turn off interrupt on a processor does not prevent other processors from accessing shared memory
  - Turn off all other processors?
- Use atomic read/write and busy waiting?
  - Concurrent read can happen
  - “Busy waiting” is not acceptable
- Need more help from HW!
- What is the right HW support?
  - Hot research topic for a long time

# [lec4] A Possible Solution?



```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

```
if ( noMilk ) {  
    if (noNote) {  
        leave note;  
        buy milk;  
        remove note;  
    }  
}
```

- process can get context switched after checking milk and note, but before leaving note
- *Why does it work for human?*



# Read-modify-write on CISC

- Most CISC machines provide atomic *read-modify-write* instruction
  - read existing value
  - store back a new value
  - Example: *test-and-set* by IBM and others

```
int TAS(int *old_ptr, int new {  
    int old = *old_ptr; // fetch old value at old_ptr  
    *old_ptr = new; // store 'new' into old_ptr  
    return old; // return the old value  
}
```

- Using TAS to implement lock (mutex)

# Using test-and-set for mutual exclusion (too-much milk)



- Implement a critical section on multiprocessor
  - Prevents 2 processes doing 0-to-1 transition simultaneously

```
global int lock = 0;  
...  
??????  
...  
critical section  
...  
??????
```



# Using test-and-set for mutual exclusion (too-much milk)



- Implementation with **spin lock** (busy wait)

```
global int lock = 0;  
...  
while (TAS(&lock, 1) == 1);  
...  
critical section  
...  
lock = 0;
```

# Evaluating spin lock implementation with TAS



- Correctness?
  - Yes (assume preemptive scheduler on uniprocessor)
- Performance?
  - No, esp. when critical section is long

# Use TAS to implement semaphores on multiprocessor



- For each semaphore, keep an extra integer (lock)

```
typedef struct {  
    int lock; /* initially 0 */  
    int count;  
    queue q; /* queue of procs waiting on this semaphore */  
} semaphore;
```

# Use TAS to implement semaphores on multiprocessor?



```
void wait(semaphore s)
{
    disable interrupts;
    while (1 == tas(&lock,1));
    if (s->count > 0) {
        s->count --;
        ???
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    ???
    sleep(); /* re-dispatch */
    enable interrupts;
}
```

```
void signal(semaphore s)
{
    disable interrupts;
    ???
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        thread = removeFirst(s->q);
        wakeup(process);
        /* put process on Ready Q */
    }
    ???
    enable interrupts;
}
```

# Use TAS to implement semaphores on multiprocessor?



```
void wait(semaphore s)
{
    disable interrupts;
    while (1 == tas(&lock,1));
    if (s->count > 0) {
        s->count --;
        lock = 0;
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    lock=0;
    sleep(); /* re-dispatch */
    enable interrupts;
}
```

```
void signal(semaphore s)
{
    disable interrupts;
    while (1 == tas(&lock,1));
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        thread = removeFirst(s->q);
        wakeup(process);
        /* put process on Ready Q */
    }
    lock = 0;
    enable interrupts;
}
```



# Deep Thinking

- Why is this busy waiting not a concern?
  - What's our critical section here?
- Can we remove the disable/enable interrupts?
  - With interrupts, when a process that gets the lock (pass the while loop) get context switched out
  - All other wait processes on other cores will be busy waiting

# Implementing Locks+Condition Variables

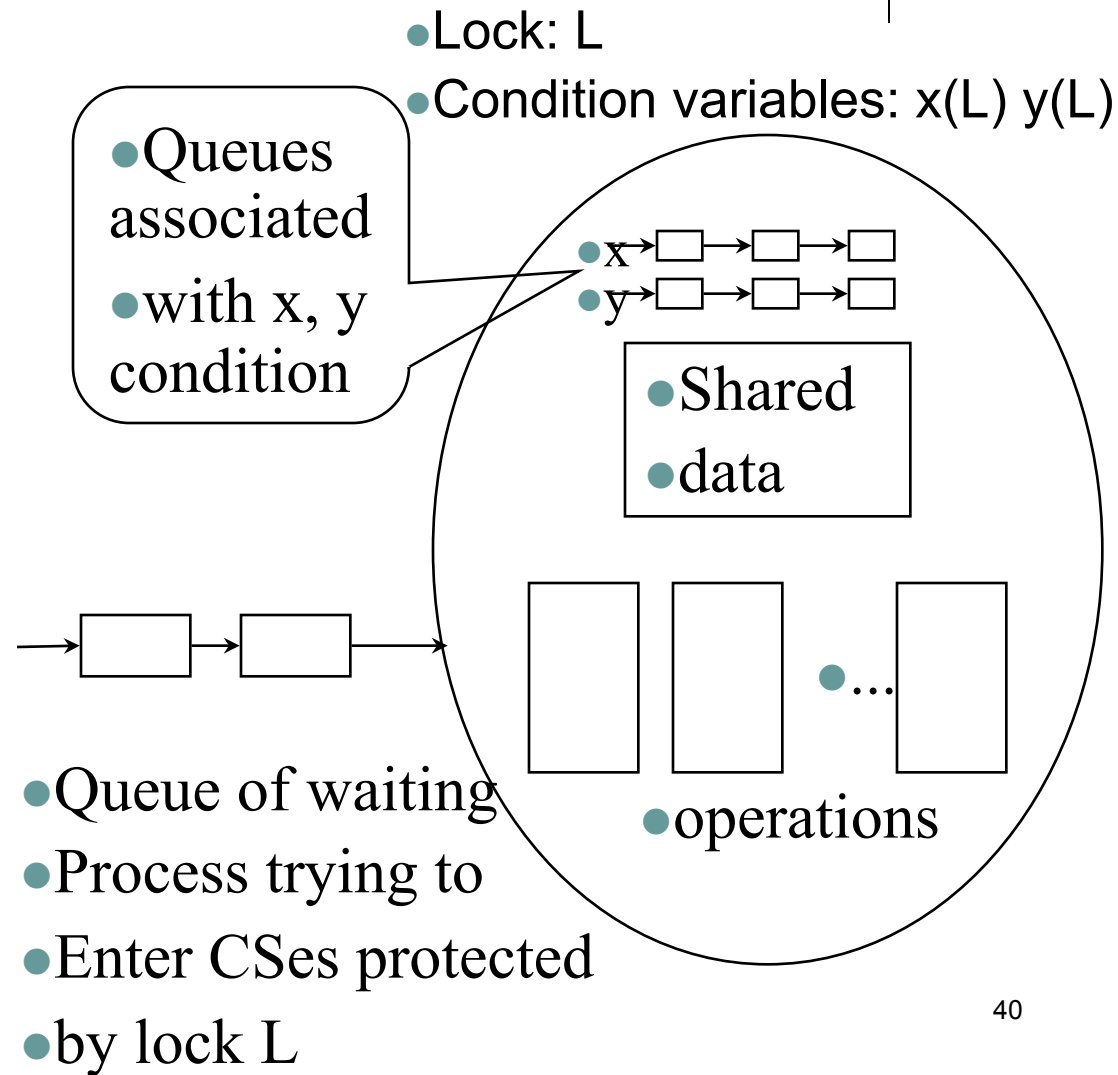


- Use the same mechanisms for
  - achieving atomicity and
  - avoiding busy waitingas in semaphore implementation



# [lec5] Condition Variables

- Wait (condition)
  - Block on “condition”
- Signal (condition)
  - Wakeup a process blocked on “condition”
- Conditions are like semaphores but:
  - signal is no-op if none blocked
  - There is no counting!







# [lec5] Two Options of Signaler

- Relinquishes control to the awoken process; suspend signaler (Hoare-style, early time)
  - Signaler gives up lock, waiter runs immediately
  - Waiter gives back lock and CPU to signaler after critical sec.
  - Complex if the signaler has other work to to
  - In general, easy to prove things about system (e.g. fairness)
- Continues its execution (Mesa-style, modern)
  - Signaler keeps lock and CPU
  - Waiter put on ready queue
  - Easy to implement (e.g., no need to keep track of signaler)
  - But, what can happen when the awoken process gets a chance to run?
    - E.g. producer 1 context switch at line 1, producer 2 wait, consumer 1 signals producer 2

# Roadmap on Process Management



- Processes
- Need for process synchronization
- Mutual exclusion & Critical section
- “Too much milk” problem
- Semaphore (binary semaphore)
- Producer-consumer problem
- Condition variable (Monitor)
- *Time to dive inside OS*

# Busy waiting does not work – We are going back to where we started (in a way)?



- Uniprocessor (read/write atomic)
- Concurrent execution → needs busy waiting
  - Busy waiting in apps not efficient →  
Atomic sync primitive from OS: e.g. semaphore
  - OS has to implement semaphore atomically →  
disable interrupt
- Multiprocessor (read/write concurrent)
  - Concurrent read/write by processes on processors
  - If do not disable interrupt on all processors?
    - Needs to bring back busy-waiting to help
    - Concurrent execution: No busy waiting -> no mutual exclusion -> no atomicity -> No P & V

# Semaphores vs. lock/condition variables




- survey:
  - How many feel lock/cond variables are easier?
- My carefully designed lecture was unintentionally sabotaged by ...
  - The sequence of using lock, using semaphore, and then using lock/cond variables
- We then dived into the inner workings condition variables
  - Which is not a bad because we needed to understand
  - Along way, many engr ideas optimizing perf came up
  - They were all good ideas, showing we engr can opt perf

# Sema vs. lock/condition variables



- But let us not lose sight on the big picture
  - That lock/condition vars are more natural
- Let us decouple
  - How easy it is to program (use)
    - Lock/cond variables fit naturally the programming thought process (evolution from lock())
      - Details of condition left to programmers → more flexible
    - Semaphores are like revolution, clean-slate new design
      - **Bundle condition** (i.e., resource counting) with sync variable (i.e., semaphore) may not be easy in general
  - How easy it is to understand how \*\* works internally

# Backup Slides



# Load-Linked and Store-Conditional

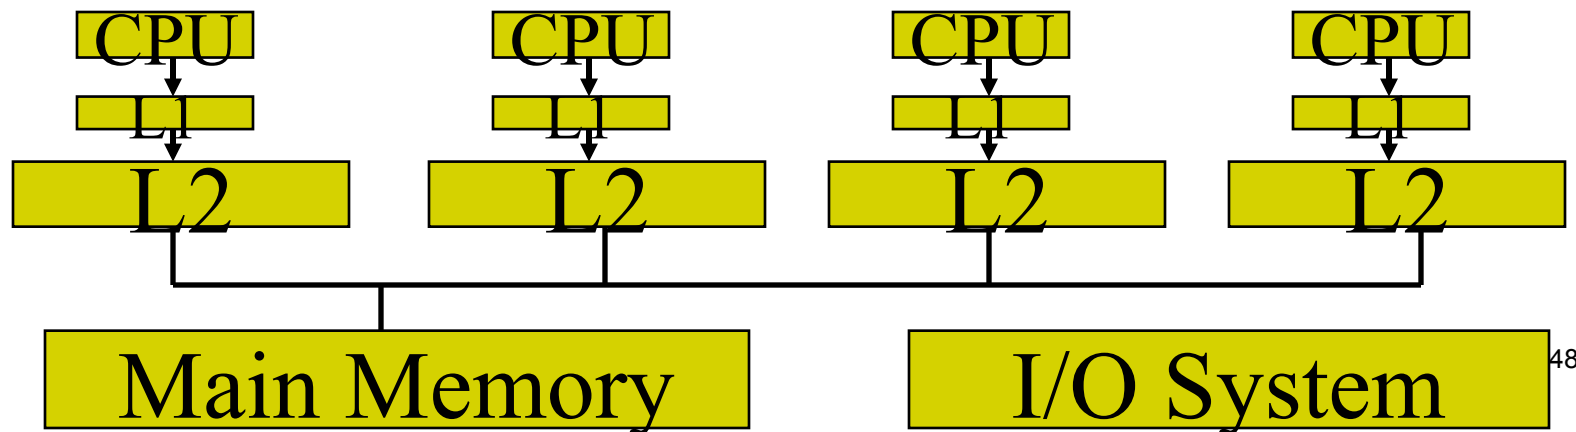


- Load-linked
  - Operates like a typical load instruction, and simply fetches a value from memory and places it in a register.
- Store-conditional
  - Only succeeds (and updates the value stored at the address just load-linked from) if no intervening store to the address has taken place.

# Load-Linked and Store-Conditional on RISC [MIPS R4000 series]




- *Load-linked* instruction: **LDL Rx,y**
  - loads Rx with a word from mem addr y
  - holds y in per-processor lock register
- Store operation to addr y (by any processor) resets all other processors' lock registers if containing addr y
- *Store-conditional* instruction: **STC Rx, y**
  - stores a word iff y matches the processor's lock register
  - indicates success ( 1 ) or failure ( 0 )





# Using load-linked and store-conditional to implement locks



```
int lock=0;
...
?????
...
critical section
...
?????
```

```
int lock=0;
...
while (ldl(&lock) == 1) || stc((&lock, 1) == 0);

...
critical section
...
lock = 0;
```

can ldl,

# Use Idl/stc to implement semaphores on multiprocessor



```
void wait(semaphore s)
{
    disable interrupts;
    ???
    if (s->count > 0) {
        s->count --;
        ???
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    ???
    sleep(); /* re-dispatch */
    enable interrupts;
}
```

```
void signal(semaphore s)
{
    disable interrupts;
    ???
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        thread = removeFirst(s->q);
        wakeup(process);
        /* put process on Ready Q */
    }
    ???
    enable interrupts;
}
```

# Use ldl/stc to implement semaphores



```
void P(semaphore s)
{
    disable interrupts;
    while (ldl(s->lock) == 1 ||
           stc(s->lock, 1) == 0);
    if (s->count > 0) {
        s->count --;
        s->lock = 0;
        enable interrupts;
        return;
    }
    add(s->q, current_process);
    s->lock = 0;
    sleep(); /* re-dispatch */
    enable interrupts;
}
```

```
void V(semaphore s)
{
    disable interrupts;
    while (ldl(s->lock == 1) ||
           stc(s->lock, 1) == 0);
    if (isEmpty(s->q)) {
        s->count ++;
    } else {
        thread = removeFirst(s->q);
        wakeup(thread);
        /* put process on Ready Q */
    }
    s->lock = 0;
    enable interrupts;
}
```