

CPU Scheduling (cont)

ECE469, Feb 9

Yiying Zhang



Readings



- Dinosaur Chapter 5
- Comet Chapter 7, 8, 9



True or false?

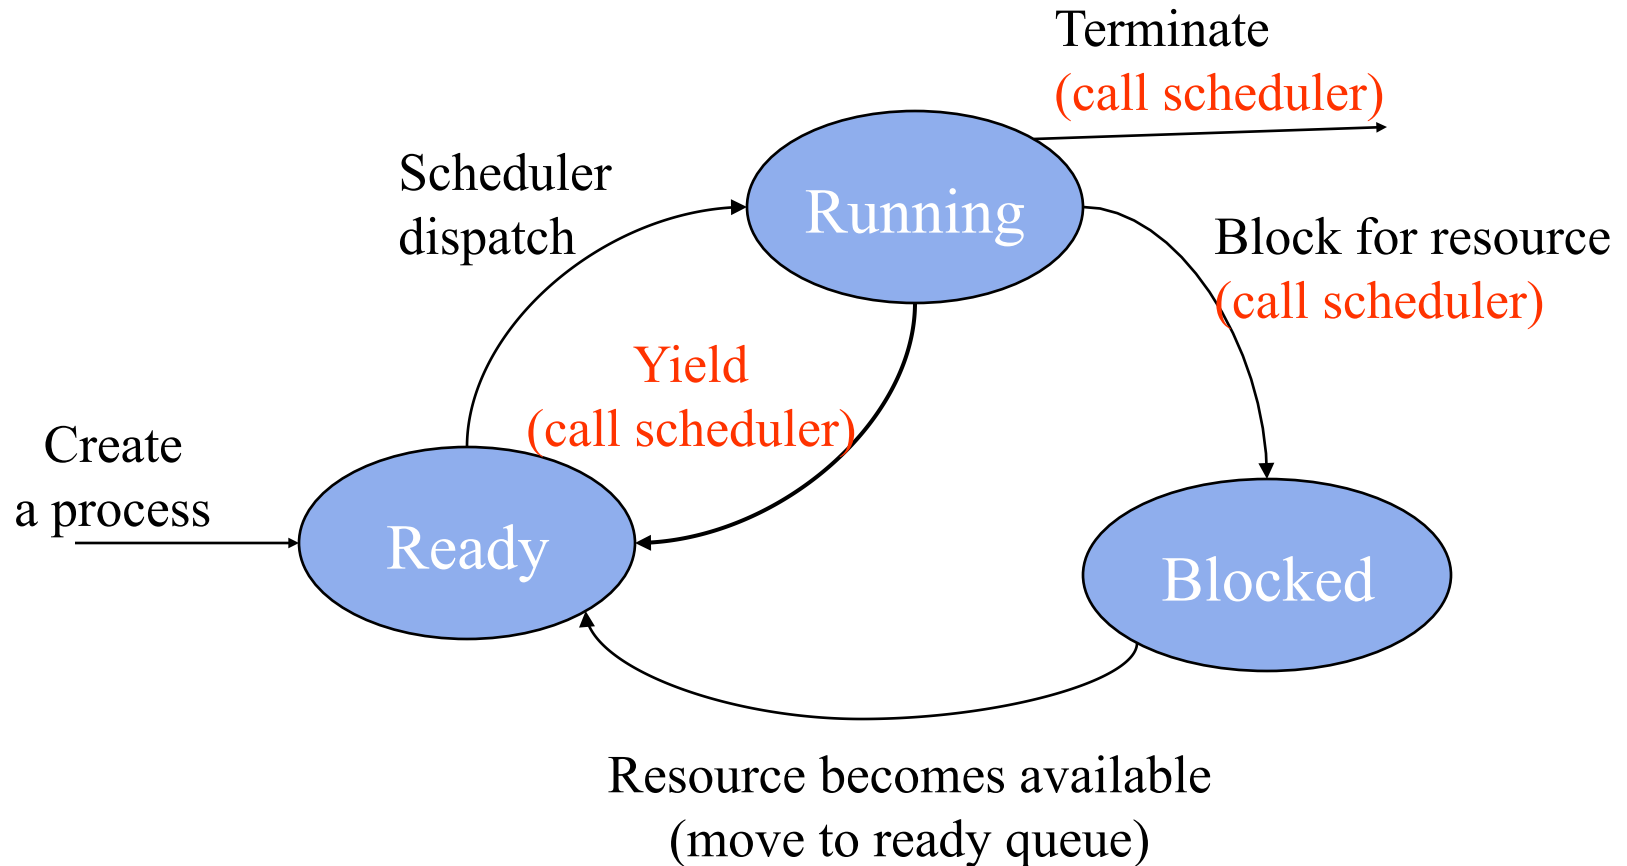
- “A CPU scheduling algorithm that minimizes avg turnaround time cannot lead to starvation.”
- “Among all CPU scheduling algorithms, Round Robin always gives the worse average turnaround time.”

Timesharing Systems

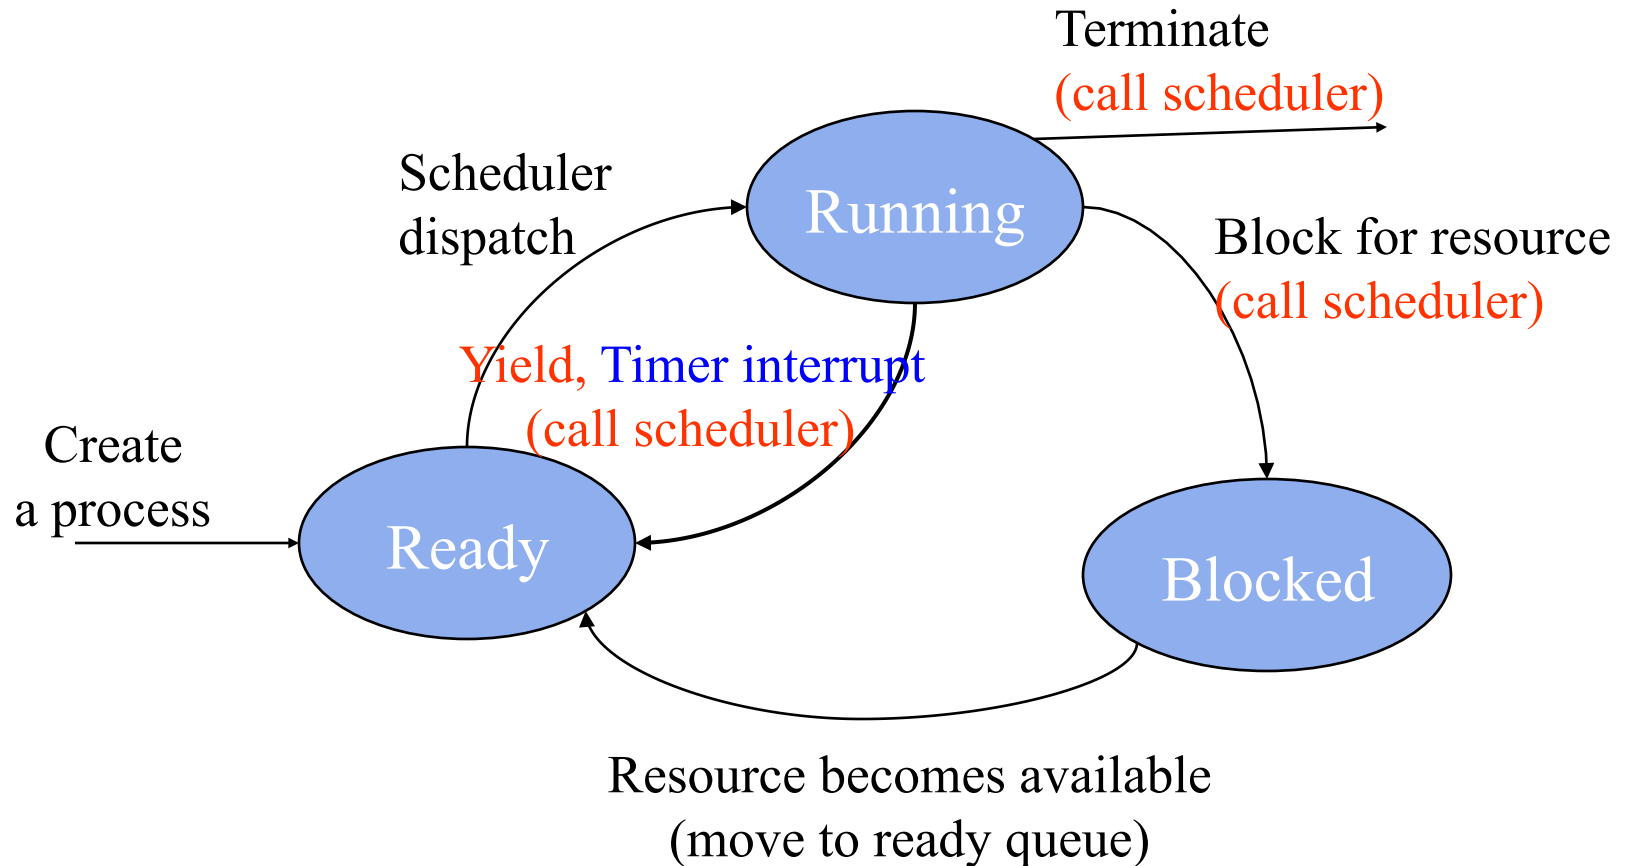


- **Timesharing** systems support interactive use
 - each user feels he/she has the entire machine
 - CPU, memory, I/O devices
- CPU - How?
 - optimize response time
 - based on time-slicing

Process State Transition of Non-Preemptive Scheduling



Preemptive Scheduling



Review on CPU scheduling



- Mechanism is easy, policy is hard
 - Jobs have diverse characteristics
 - 4 performance metrics
 - Don't know about future
- Hard to analyze even when narrowing down metric/job nature

[lec 9] Scheduling policies




- FIFO
 - Round Robin
 - SJCF
 - SRTCF
-
- FIFO vs. Round Robin in average turnaround time
 - 10 procs, 100 sec each, who wins?
 - 10 procs, 100 sec for first, 10 sec other 9, who wins?

STCF (SJF) – Shortest Job First



- What shall we do if we care about turnaround time?

- FIFO can be bad



- STCF/SJF

- schedule shortest (total completion time) job first



SJF – Pros and Cons




- Can we do better than Shortest Job First in terms of average turnaround time?
 - Assume all jobs arrive at the beginning
- In fact, SJF can be proved to be the optimal scheduling algorithm with the above assumption
 - But we are not going to prove it, since this is not a theory class 😊
- SJF Advantage
 - Minimal average turnaround time
- Disadvantage
 - Difficult to know the future, has to run until finish

STCF vs. SRTCF




- Shortest time to completion first (shortest job first)
 - Non-preemptive



- Shortest *remaining* time to completion first

- Preemptive



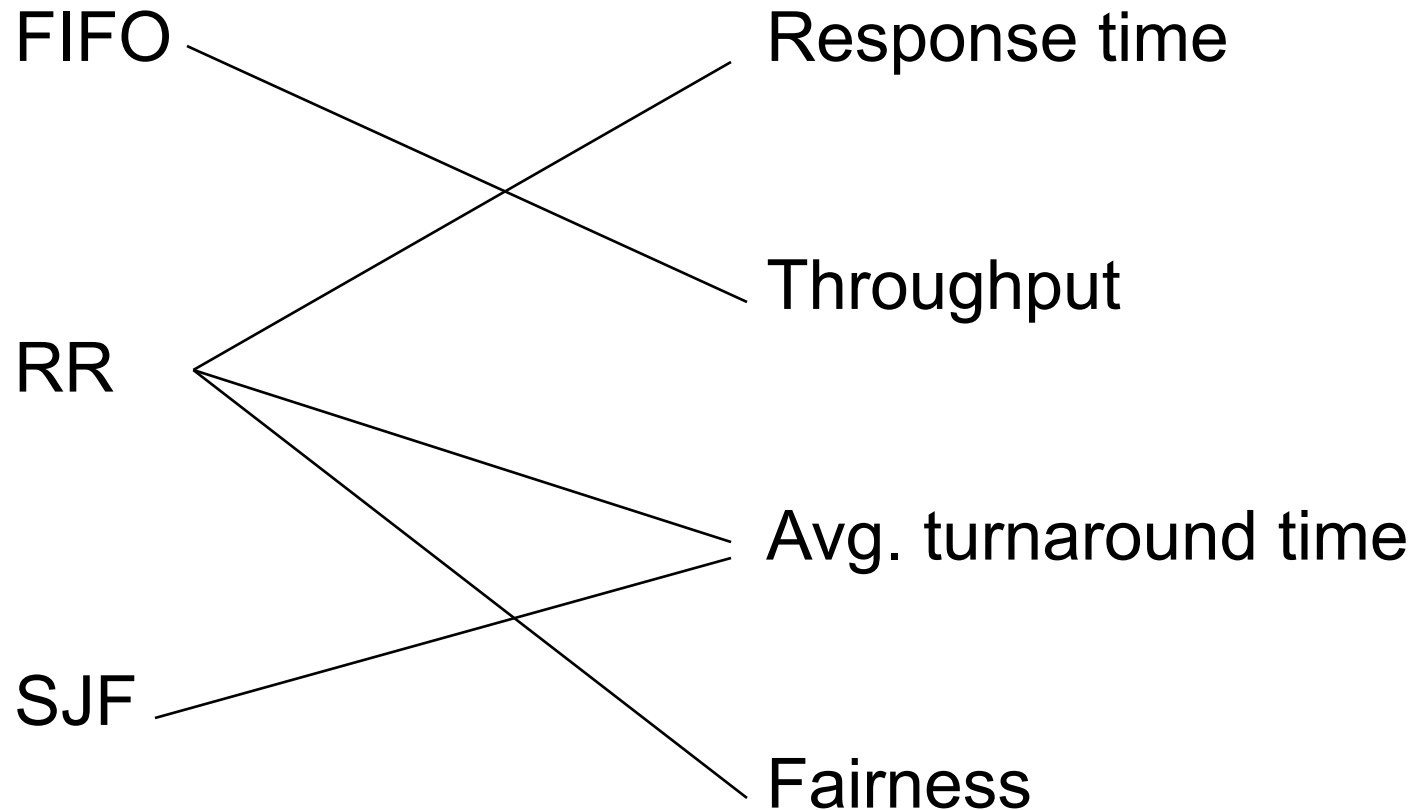
Any potential problems?
- Can cause **starvation!**




Observations so far

- Need to accommodate interactive jobs
 - Need some kind of RR
- Diversity in jobs – job length, I/O mix
 - RR also appears to help
- SJF also has virtue
 - Reduce avg. turnaround time
- Can we accommodate all?

Scheduling policies





What Issues are in Policy?

- High utilization (efficiency)
 - Lots of processes (want diff resources)
 - Lots of resources (want full parallelism)
- Issue?
 - How do you get the most ***useful*** work out of the system? (job throughput)



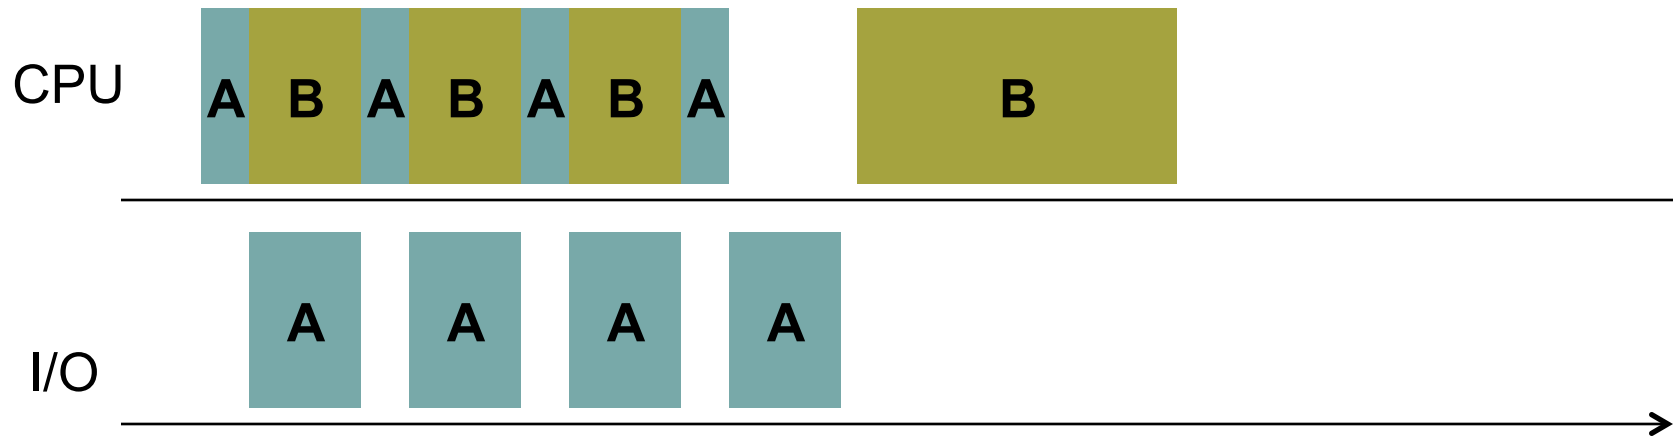

What Issues are in Policy?

- Fairness
- Flexibility
- High utilization (efficiency)
- Good response time
- Good turnaround time


Other issues?

- Smooth media playback, device handling

What about I/Os?

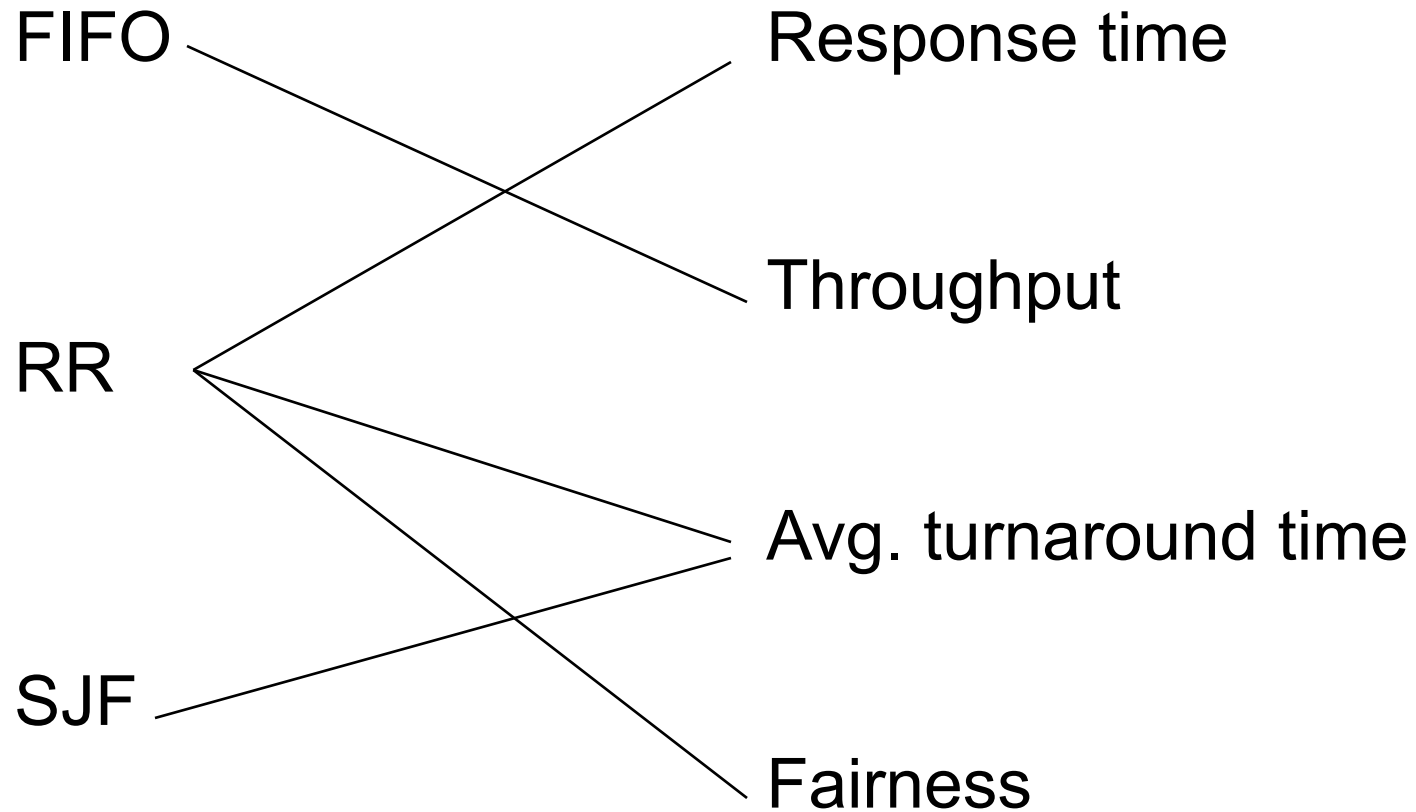


Adding I/O Into the Mix



- Resource utilization example
 - A and B each uses 100% CPU
 - C loops forever (1ms CPU and 10ms disk)
 - Time slice 99ms: nearly 30% of disk utilization with Round Robin and nearly 70% of CPU utilization
 - Time slice 1ms: nearly 90% of disk utilization with Round Robin and nearly 100% of CPU utilization
- What do we learn from this example?
 - *Small time slice can improve utilization / fairness to I/O jobs*

Scheduling policies



Priority Scheduling



- To accommodate the spirits of SJF/RR/FIFO
- The method
 - Assign each process a *priority*
 - Run the process with highest priority in ready queue first
 - Use FIFO for processes with equal priority
 - Adjust priority dynamically
 - To deal with *all* issues: e.g. aging, I/O wait raises priority
- Advantage
 - Flexibility: Not all processes are “born” equal
- Challenge?



Priority Scheduling (cont)

- Who sets the priorities
 - Internally by OS
 - I/O to computation ratio (can be dynamic)
 - Memory requirement (can be dynamic)
 - Time constraints (e.g. real-time systems)
 - Externally by users/sysadm
 - Importance
 - Funds paid for
 - Being nice
- How -- Dynamically adjustment is tricky



Break

- Brain teaser time

- There is a person who have two numbers, he tells sum to the person S and product of those numbers to P. Now there is this conversation between S and P.
 - S: I don't know what are the numbers.
 - P: I also don't know what are the numbers.
 - S: Now I know what are the numbers.
 - P: Now I also know what are the numbers.
- Assuming S and P to be very wise and good in mathematics, What are those two numbers? Note: Numbers are greater than 0.

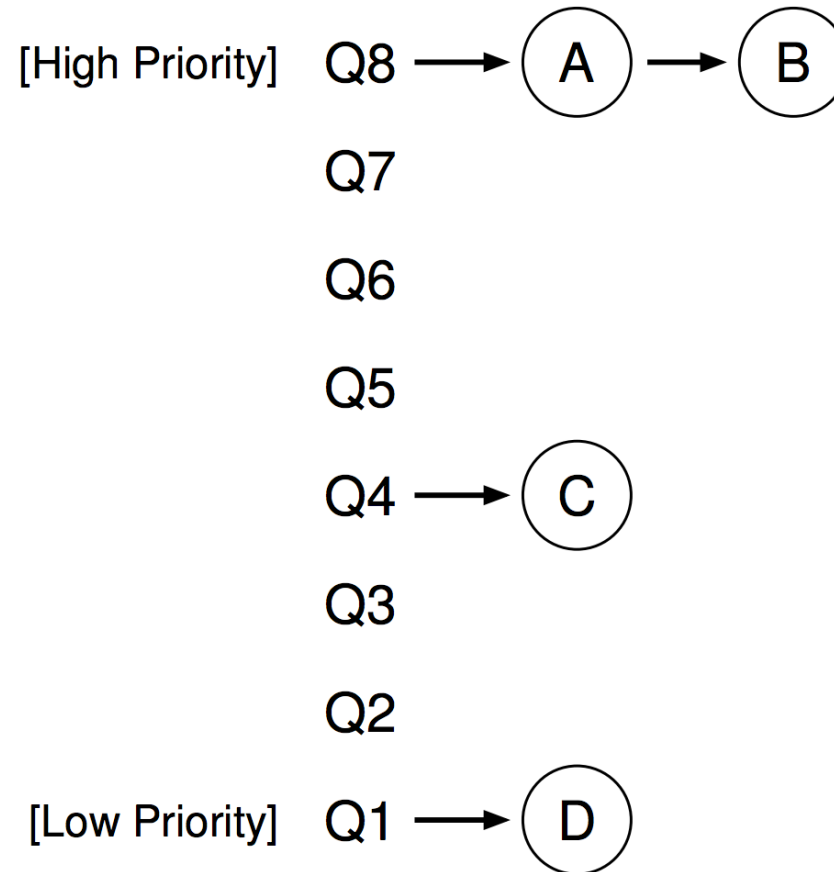

Approach 1:

Multiple Queue Scheduling



- Motivation: processes may be of different nature and can be easily classified
 - e.g. foreground jobs vs. background jobs
- The method:
 - Processes permanently assigned to one queue, based on processes priority / type
 - Preference to jobs with higher priorities
 - Each queue can have its own scheduling algorithm
 - e.g. RR for foreground queue, FCFS for background queue
 - Need a scheduling among the queues
 - e.g. fixed priority preemptive scheduling (high-pri queue trumps other)
 - e.g. time-slice between queues

Multiple Queue Example



Pros/Cons of Multiple Queue Scheduling



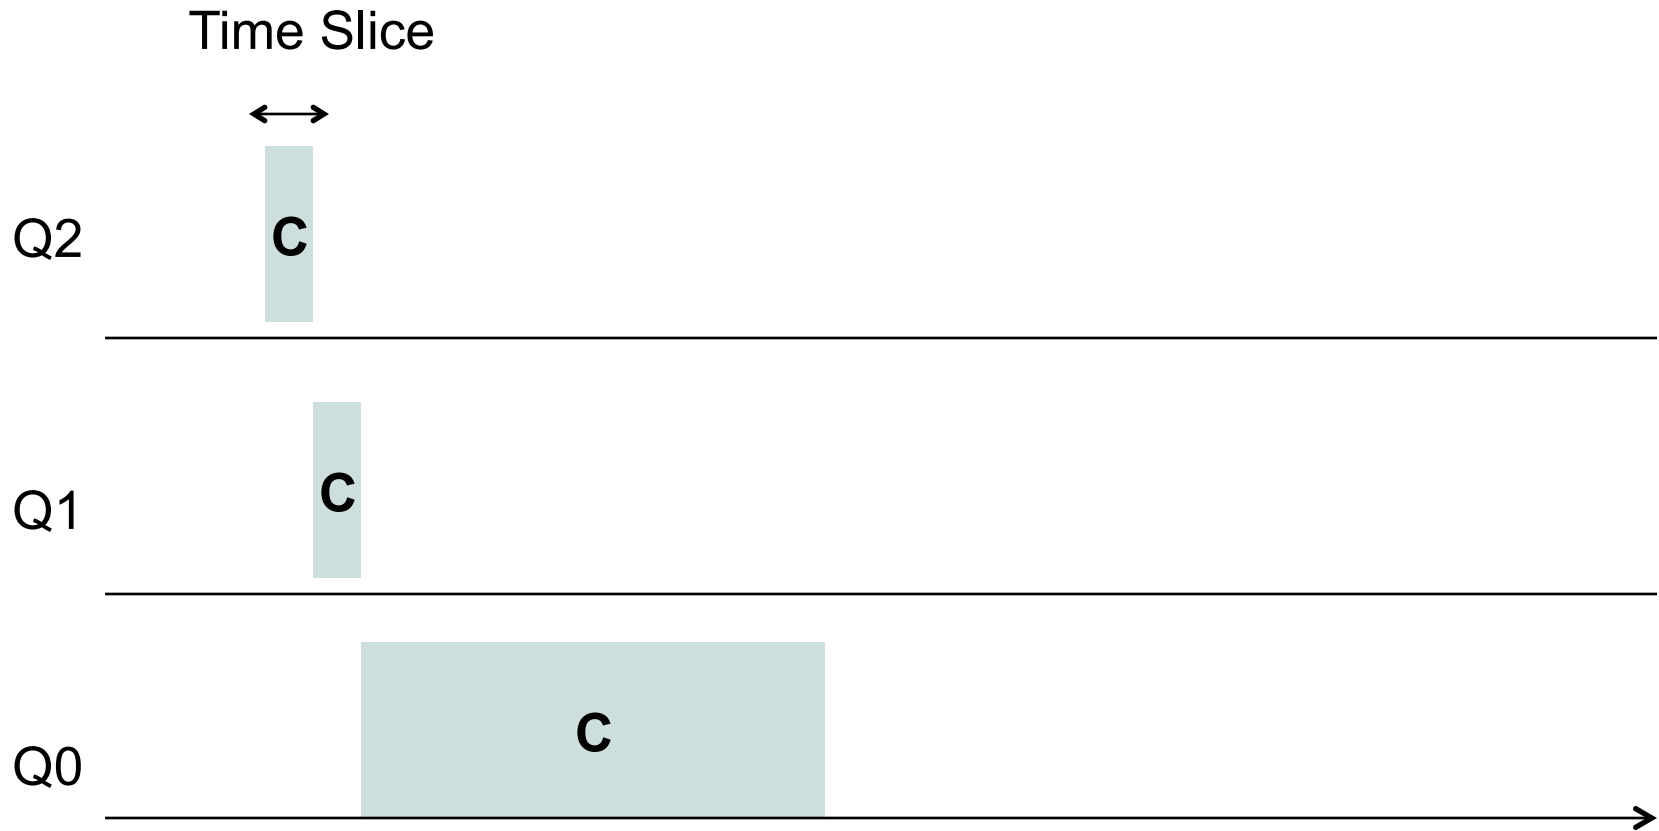
- Pros:
 - Low scheduling overhead
 - Jobs do not move across queues
- Cons:
 - Processes permanently assigned to one queue – not flexible
 - Program behavior may change
 - E.g. can switch between I/O bound and CPU bound
 - Need some learning/adaptation at runtime
 - Starvation cannot be easily handled
 - Need some learning/adaptation at runtime

Approach 2: Multilevel Feedback Queue (MLFQ)

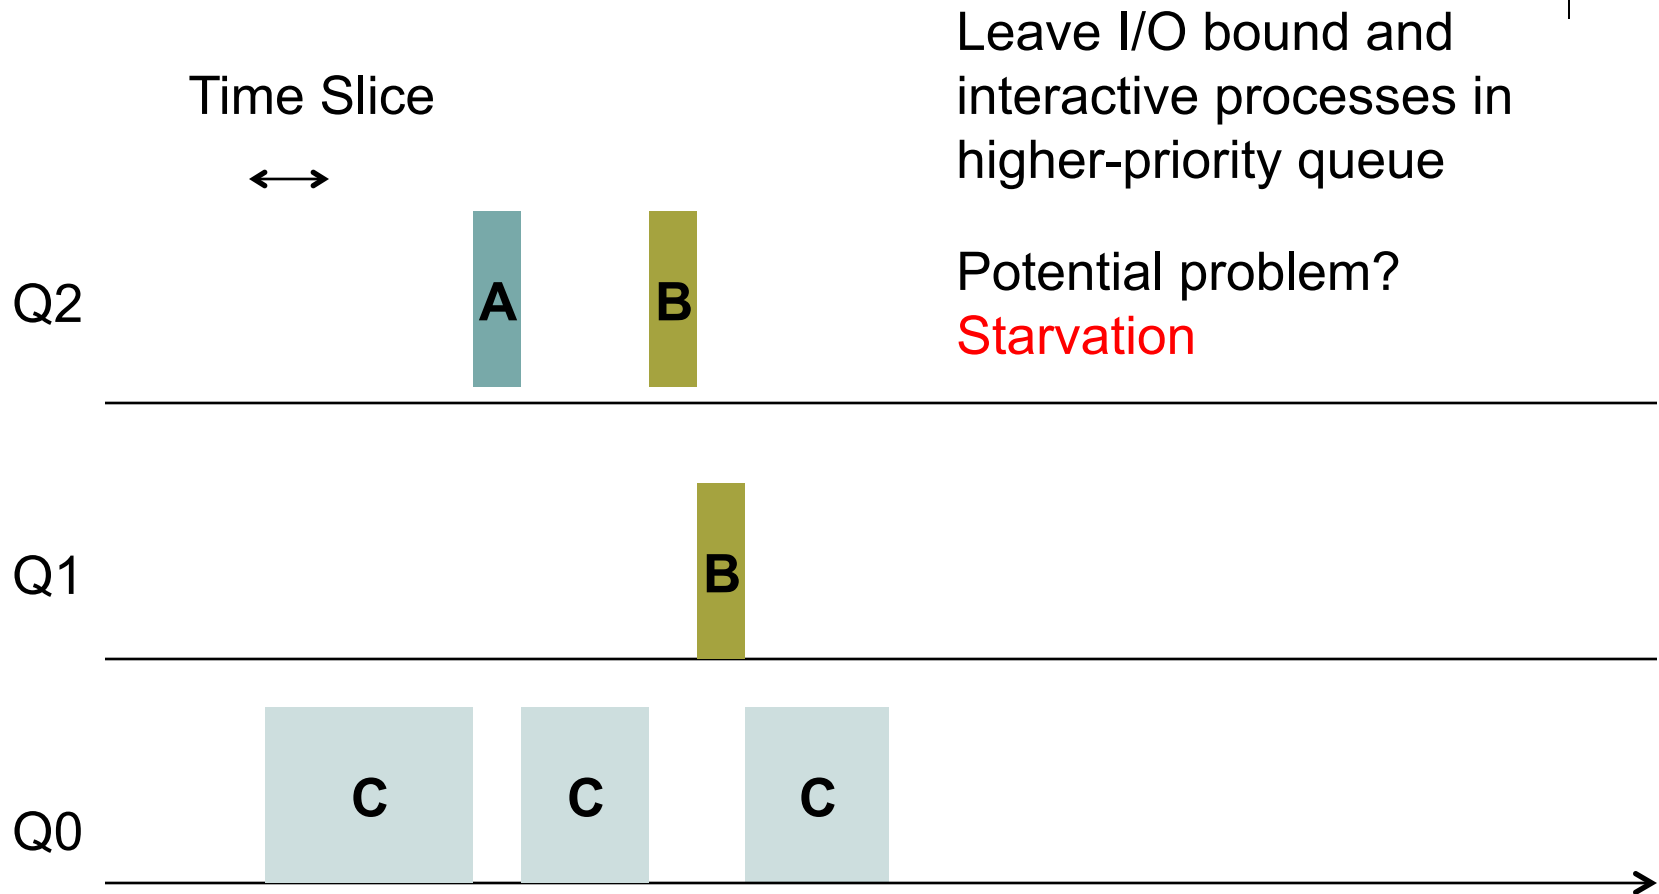


- Problem: how to change priority?
- Jobs start at highest priority queue
- Feedback
 - If a job uses up an entire time slice while running, its priority is reduced (i.e., it moves down one queue).
 - If a job gives up the CPU before the time slice is up, it stays at the same priority level.
 - After a long time period, move all the jobs in the system to the topmost queue (aging)


MLFQ Example – Single long job



MLFQ Example – a long job + short jobs in between



MLFQ Example – a long job + short jobs, with boost



Scheduling Algorithms in OSes



Operating System	Preemption	Algorithm
Windows 3.1x	None	Cooperative Scheduler
Windows 95 , 98 , Me	Half	Preemptive for 32-bit processes, Cooperative Scheduler for 16-bit processes
Windows NT (2000, XP, Vista, 7, and Server)	Yes	Multilevel feedback queue
Mac OS pre-9	None	Cooperative Scheduler
Mac OS 9	Some	Preemptive for MP tasks, Cooperative Scheduler for processes and threads
Mac OS X	Yes	Multilevel feedback queue
Linux pre-2.6	Yes	Multilevel feedback queue
Linux 2.6-2.6.23	Yes	O(1) scheduler
Linux post-2.6.23	Yes	Completely Fair Scheduler
Solaris	Yes	Multilevel feedback queue
NetBSD	Yes	Multilevel feedback queue
FreeBSD	Yes	Multilevel feedback queue



Cooperative Scheduling

- Early multitasking systems used applications that voluntarily ceded time to one another.
- This approach, which was eventually supported by many OSes, is known today as cooperative multitasking.
- Cooperative multitasking was once the scheduling scheme employed by [Microsoft Windows](#) (prior to [Windows 95](#) and [Windows NT](#)) and [Mac OS](#) (prior to [Mac OS X](#)) to enable multiple applications to be run simultaneously
- Now rarely used in larger systems



Let's look at fairness again

- Proportional share: another view of fairness
 - Each job gets a (fair) proportional of CPU time
 - Goals here are not turnout time or response time
- How to share CPU proportionally?
 - Idea: proportional => probabilistic

Lottery Scheduling [OSDI 94]



- Motivations
 - SJF does well with avg turnaround time, but unfair
 - Priority scheduling is implemented by adjusting priorities, adjusting priority is a bit ad hoc.
- Lottery method: using **probabilistic** to assign CPU time
 - Give each job a number of tickets
 - **Randomly** pick a winning tickets => jobs with more tickets have higher chance to win (get CPU)
 - To approximate priority scheduling, high priority jobs get more tickets
 - To approximate SRTCF, short jobs get more tickets
 - To avoid starvation, give each job at least one ticket

Best thing about lottery scheduling



- Easy to implement!



Real-Time Scheduling


- Two types of real-time
 - Hard deadline: must meet, otherwise can cause fatal error
 - Soft headline: meet most of the time, but not mandatory
- Characteristics
 - User control: provide users with abilities to control and specify
 - Deterministic: upper bound on when to get services on an I/O
 - Responsive: how long does OS delay before ack an interrupt



Deadline Scheduling


- Admission control
 - Take a job only if the system can guarantee real-time
- Information needed
 - Ready time: time at which task becomes ready
 - Starting deadline: time by which a task must begin
 - Completion deadline: time by which a task must complete
 - Processing time: time required to execute the task to completion
 - Resource requirements
 - Priority
 - Subtask structure

Multiprocessor and Cluster



Multiprocessor architecture


- L2 cache coherence
- A single “image” OS



Cluster/Multicomputer

- Distributed memory
- An OS on each box

Multiprocessor/Cluster Scheduling



- New design issue: process/thread inter-dependence
 - Threads of the same process may synchronize
 - Processes of the same job may send/recv messages

Multiprocessor/Cluster Scheduling: Example Approach




- Gang scheduling (coscheduling)
 - Threads of same process will run together on multiprocessor
 - Processes of same application run together on cluster
- Dedicated processor assignment
 - Threads will be running on specific processors to completion
 - Pros / cons?
 - Good for reducing cache misses
 - Bad for load balance / fairness

Case Study: Windows NT Scheduling



- Classes and priorities
 - Real time: 16 static priorities
 - Variable: 16 variable priorities, start at a base priority
 - If a process has used up its quantum, lower its priority
 - If a process waits for an I/O event, raise its priority
- Priority-driven scheduler
 - For real-time class, do round robin within each priority
 - For variable class, multiple queue feedback
- Multiprocessor scheduling
 - For N processors, run N-1 highest priority threads on N-1 processors and run remaining threads on a single processor
 - A thread will wait for processors in its affinity set, if there are other threads available (for variable priorities)

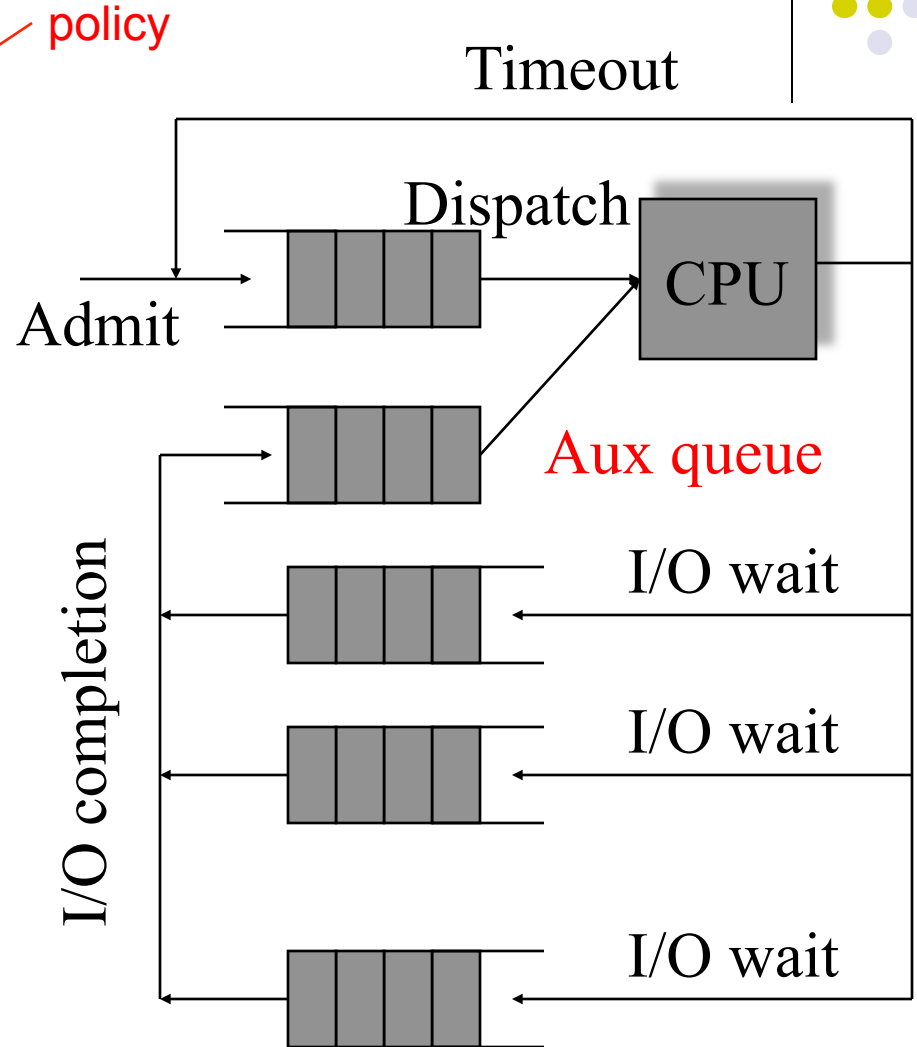
Backup






Virtual Round Robin

- To improve fairness for I/O bound processes
- Aux queue is FIFO
- I/O bound processes go to aux queue (instead of ready queue) to get scheduled
- Aux queue has preference over ready queue



mechanism

UNIX System V



- One of the first commercial versions of the Unix
- Originally developed by AT&T
- First released in 1983
- 4 major versions, termed R1 R2, R3, and 4
- SVR4, commercially most successful version, the result of Unix System Unification (with collaborations of major UNIX vendors)
 - IBM's AIX is based on SVR3
 - Sun 's Solaris and HP's HP-UX are based on SVR4
- 80s-early 90s, System V and BSD were the two major "flavors" of UNIX
 - Large multi-user systems vs. Desktop workstations

Case Study: Traditional Unix Scheduling (SVR3, 4.3 BSD)



Multilevel Feedback Queue with 1 sec preemption

- “1 sec” preemption
 - Preempt if a process doesn't block or complete within 1 sec
- Priority is recomputed every sec (low # is higher)

$$P_i = \text{base} + \text{CPU}_{i-1} / 2 + \text{nice}, \text{ where } \text{CPU}_i = (U_i + \text{CPU}_{i-1}) / 2$$

base is the base priority of the process

U_i is process utilization in interval i , nice in $[-20, 20]$

- Base priorities
 - Virtual memory Swapper
 - Block I/O device control
 - Character I/O device control
 - User processes