


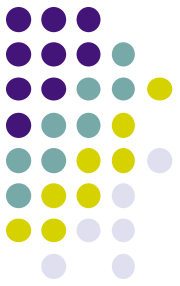
# Sharing Main Memory, Segmentation

---

ECE469, Feb 21

Yiying Zhang

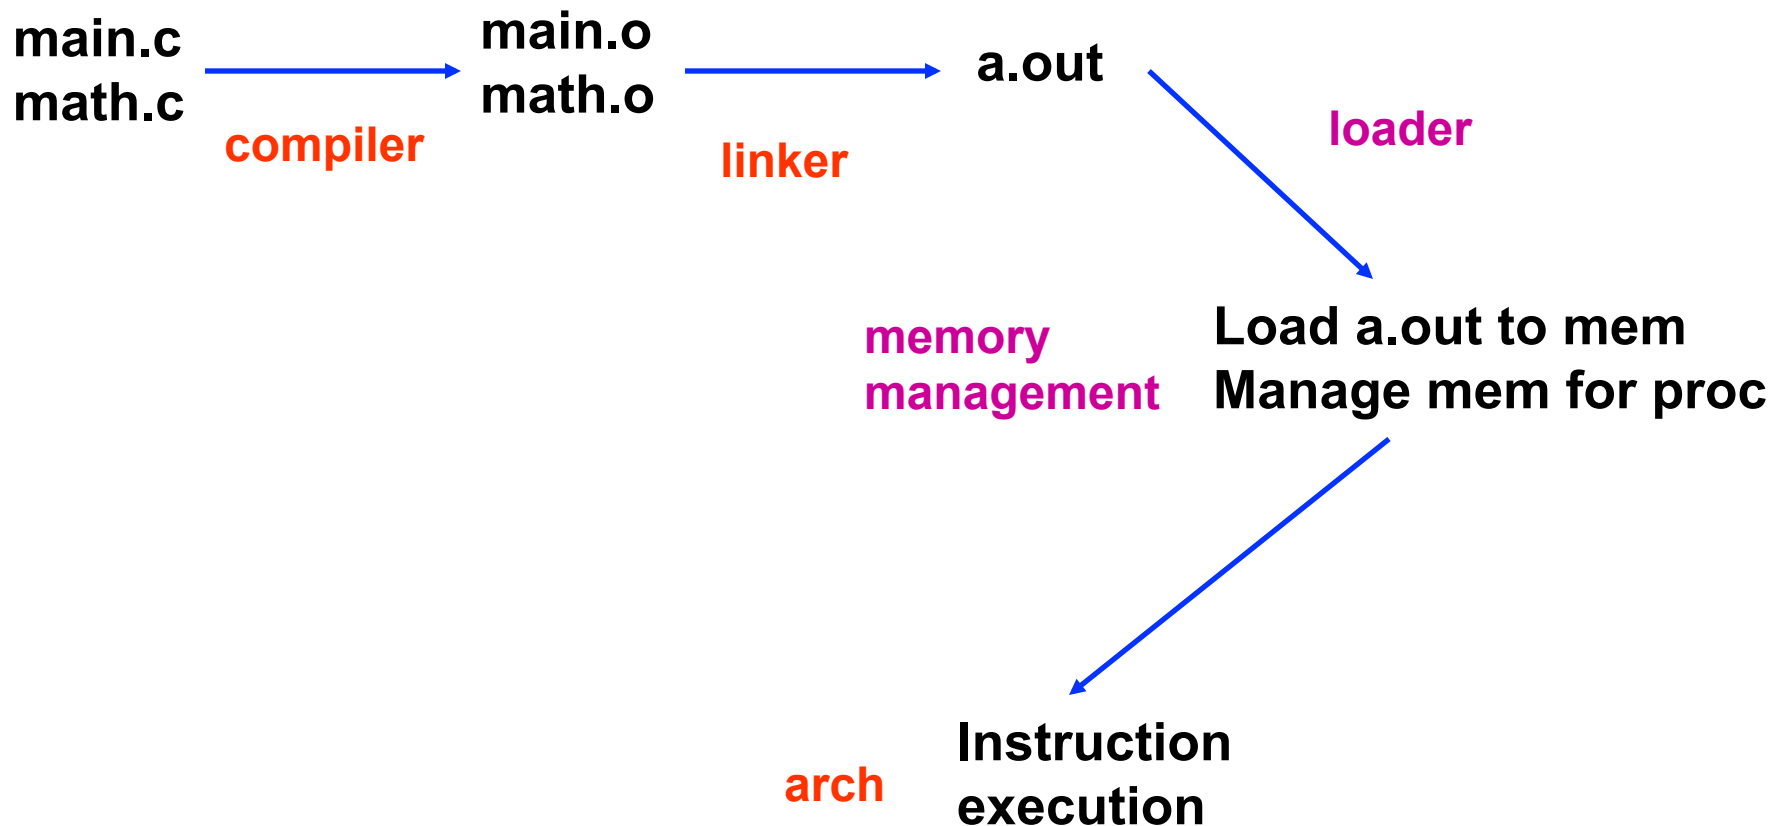




# Reading assignment

- Dinosaur Chapter 8
- Comet Chapter 13, 15, 16

# Connecting the dots





# The big picture

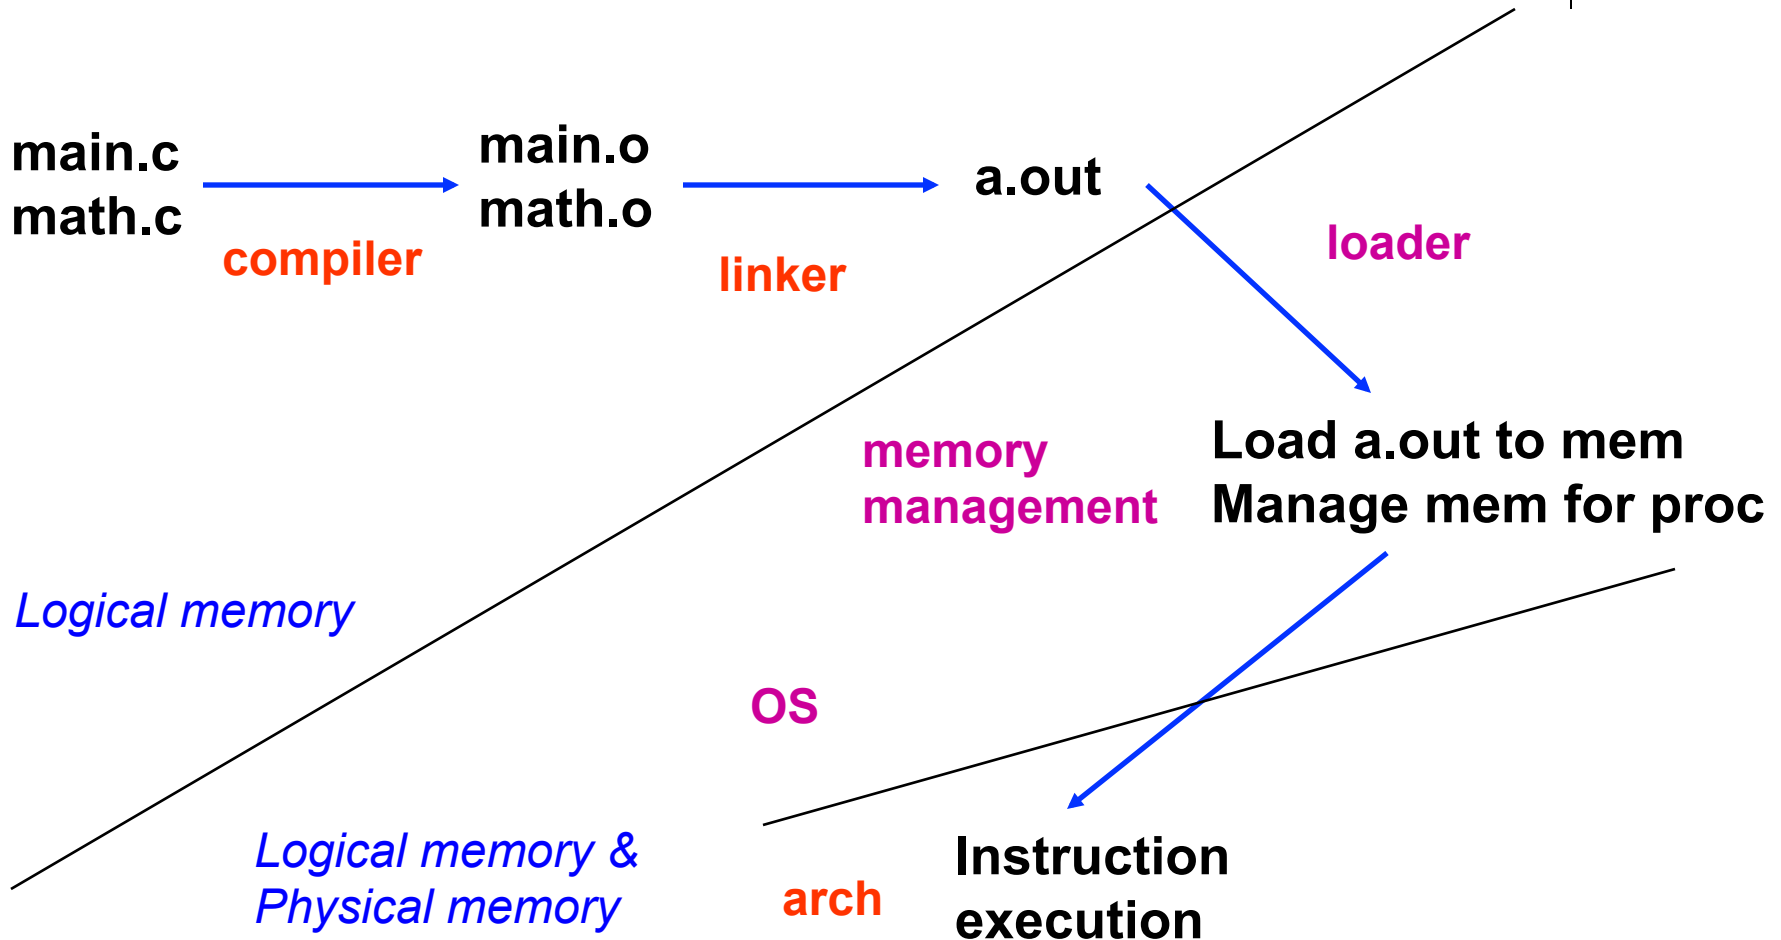
- a.out needs address space for
  - text seg, data seg, and (hypothetical) heap, stack
- A running process needs phy. memory for
  - text seg, data seg, heap, stack
- But no way of knowing where in phy mem at
  - Programming time, compile time, linking time
- **Best way out?**
  - Make agreement to divide responsibility
    - Assume address starts at 0 at prog/compile/link time
    - OS needs to work hard at loading/runing time



# Big picture (cont)

- OS deals with physical memory
  - Loading
  - Sharing physical memory between processes
  - Dynamic memory allocation

# Connecting the dots

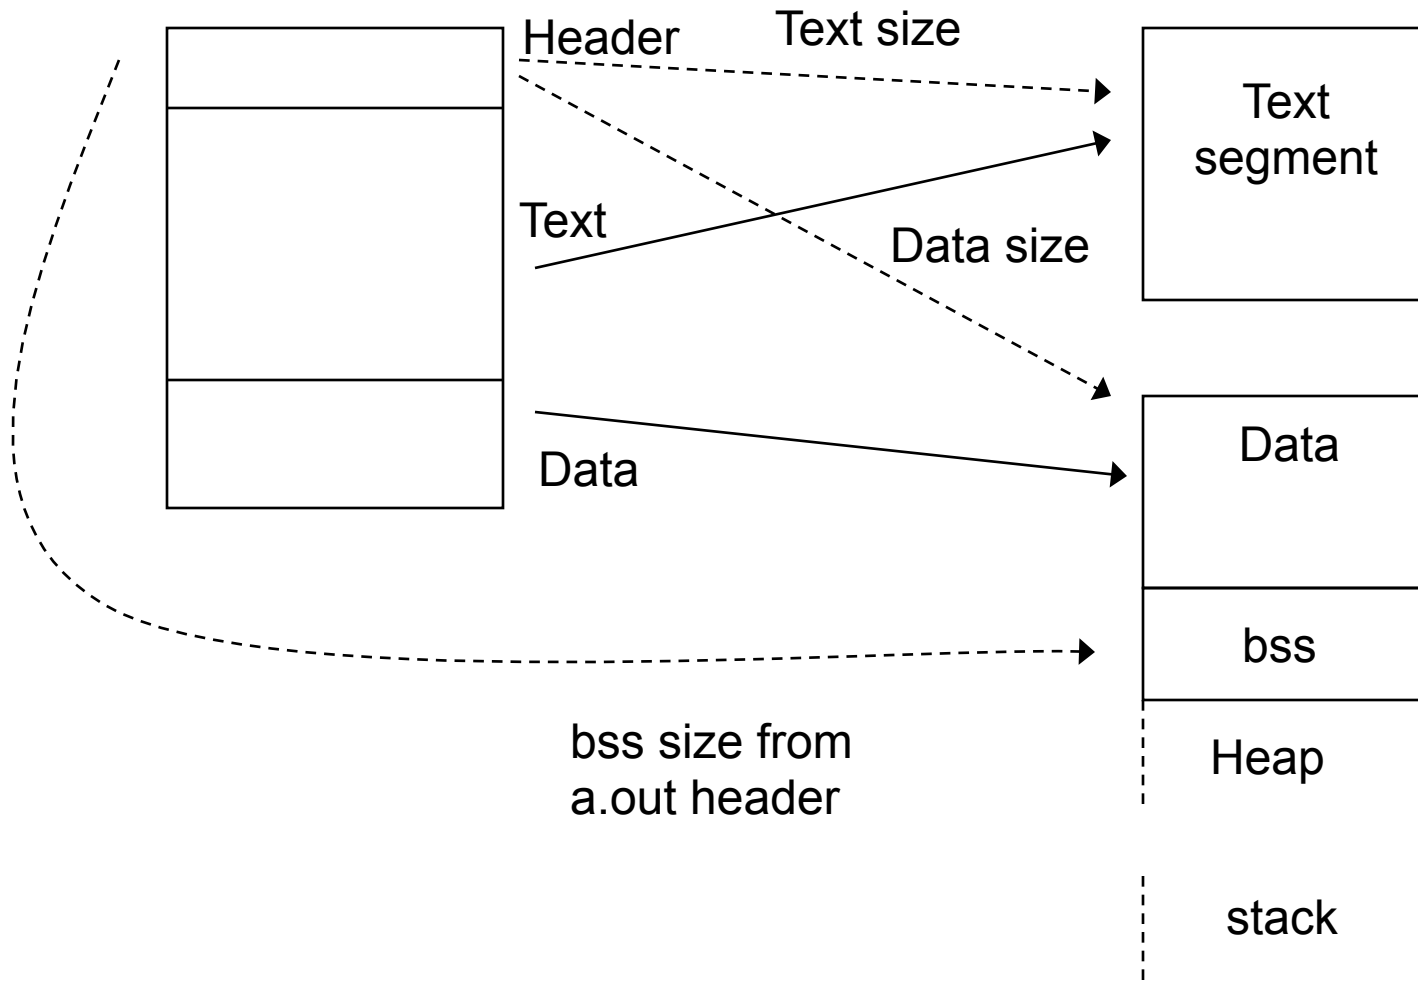


# Loading



**a.out file**

**Process**



bss = block started by symbols (un-initialized data segment)

# Dynamic memory allocation during program execution




- Stack: for procedure calls
- Heap: for malloc()
- Both dynamically growing/shrinking
- Assumption for now:
  - Heap and stack are fixed size
  - OS has to worry about loading 4 segments per process:
    - Text
    - Data
    - Heap
    - stack




# 1. Simple uniprogramming:

Single segment (code, data, stack heap) per process



Physical memory




# Simple uniprogramming:

## Single segment per process

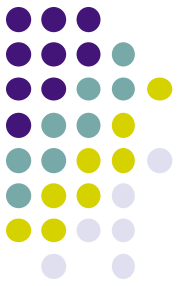


- Highest memory holds OS
- Process is allocated memory starting at 0, up to the OS area
- When loading a process, just bring it in at 0
  - virtual address == physical address!
- Examples:
  - early batch monitor which ran only one job at a time
    - if the job wrecks the OS, reboot OS
  - 1<sup>st</sup> generation PCs operated in a similar fashion
- Pros / Cons?



# Multiprogramming

- Want to let several processes coexist in main memory



# Issues in sharing main memory

- **Transparency:**
  - Processes should not know memory is shared
  - Run regardless of the number/locations of processes
- **Safety:**
  - Processes mustn't be able to corrupt each other
- **Efficiency:**
  - Both CPU and memory utilization shouldn't be degraded badly by sharing



## 2. Simple multiprogramming


With **static software memory relocation**, no protection, **1 segment per process**:

- Highest memory holds OS
- Processes allocated memory starting at 0, up to the OS area
- When a process is loaded, **relocate** it so that **it can run in its allocated memory area**
  - How? (use symble table and relocation info)
- Analogy to linking?

# Simple multiprogramming: Single segment per process, static relocation



# Simple multiprogramming: Single segment per process, static relocation



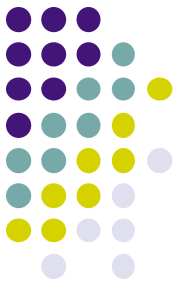
# Simple multiprogramming:

## Single segment per process, static relocation




- four drawbacks
  1. No protection
  2. Low utilization -- Cannot relocate dynamically
    - Binary is fixed (after loading)
    - Cannot do anything about holes
  3. No sharing -- Single segment per process
    - Cannot share part of process address space (e.g. text)
  4. Entire address space needs to fit in mem
    - Need to swap whole, very expensive!





# What else can we do?

- Already tried
  - Compile time / linking time
  - Loading time
- Let us try execution time!




### 3. Dynamic memory relocation

- Instead of changing the address of a program before it's loaded, change the address dynamically *during every reference*
  - Under dynamic relocation, **each program-generated address** (called a *logical address* or *virtual address*) is translated in hardware to a *physical* or *real address*


*Can this be done in software?*

# Translation overview




- Actual translation process is usually performed by hardware
- Translation table is set up by software
- CPU view
  - what program sees, virtual addresses
- Memory view
  - physical memory addresses

# 3.1 Base and bound




- Built in Cray-1 (1976)
- A program can only access physical memory in  $[base, base+bound]$
- On a context switch: save/restore base, bound registers
- Pros:
  - simple, fast translation, cheap
  - Can relocate segment

# 3.1 Base and bound



- The essence:
  - A level of (static) indirection
  - $\text{Phy. Addr} = \text{Vir. Addr} + \text{base}$

# 3.1 Base and bound



- Cons:
  - Only one segment
  - How can two processes share code while keeping private data areas (shared editors)?
    - Can it be done safely with a single-segment scheme?



# What have we solved?

## ● four drawbacks



1. No protection



2. Low utilization -- Cannot relocate dynamically

- Cannot do anything about holes


3. No sharing -- Single segment per process

- Cannot share part of process address space (e.g. text)

4. Entire address space needs to fit in mem

- Need to swap whole, very expensive!

## 3.2 Multiple Segments



- Have a table of (seg, size)
- Further protection: each entry has (nil, read, write, exec)
- On a context switch: save/restore the table (or a pointer to the table) in kernel memory



# How does this allow two processes to share code segment?



# Segmentation example



text segment [0x0000, 0x04B0]

foo:

019A: LD R1, 15DC

01C2: jmp 01F4

01E0: call 0320

01F4: X:

bar procedure

0320: bar:

Data segment [0x1000, 0x16A0]

15DC: \_Y:

## 2-bit segment number, 12-bit offset

Segment	Base	Bounds	RW
---------	------	--------	----

0	4000	4B0	10
---	------	-----	----

1	0	6A0	11
---	---	-----	----

2	3000	FFF	11
---	------	-----	----

3	--	--	00
---	----	----	----

# Segmentation example



text segment [0x0000, 0x04B0]

foo:

019A: LD R1, 15DC

01C2: jmp 01F4

01E0: call 0320

01F4: X:

bar procedure

0320: bar:

Data segment [0x1000, 0x16A0]

15DC: \_Y:

## 2-bit segment number, 12-bit offset

Segment	Base	Bounds	RW
---------	------	--------	----

0	4000	4B0	10
---	------	-----	----

1	0	6A0	11
---	---	-----	----

2	3000	FFF	11
---	------	-----	----

3	--	--	00
---	----	----	----

➔ Where is 01F4 in physical memory?

# Segmentation example



text segment [0x0000, 0x04B0]

foo:

019A: LD R1, 15DC

01C2: jmp 01F4

01E0: call 0320

01F4: X:

bar procedure

0320: bar:

Data segment [0x1000, 0x16A0]

15DC: \_Y:

## 2-bit segment number, 12-bit offset

Segment	Base	Bounds	RW
---------	------	--------	----


0	4000	4B0	10
---	------	-----	----

1	0	6A0	11
---	---	-----	----

2	3000	FFF	11
---	------	-----	----

3	--	--	00
---	----	----	----

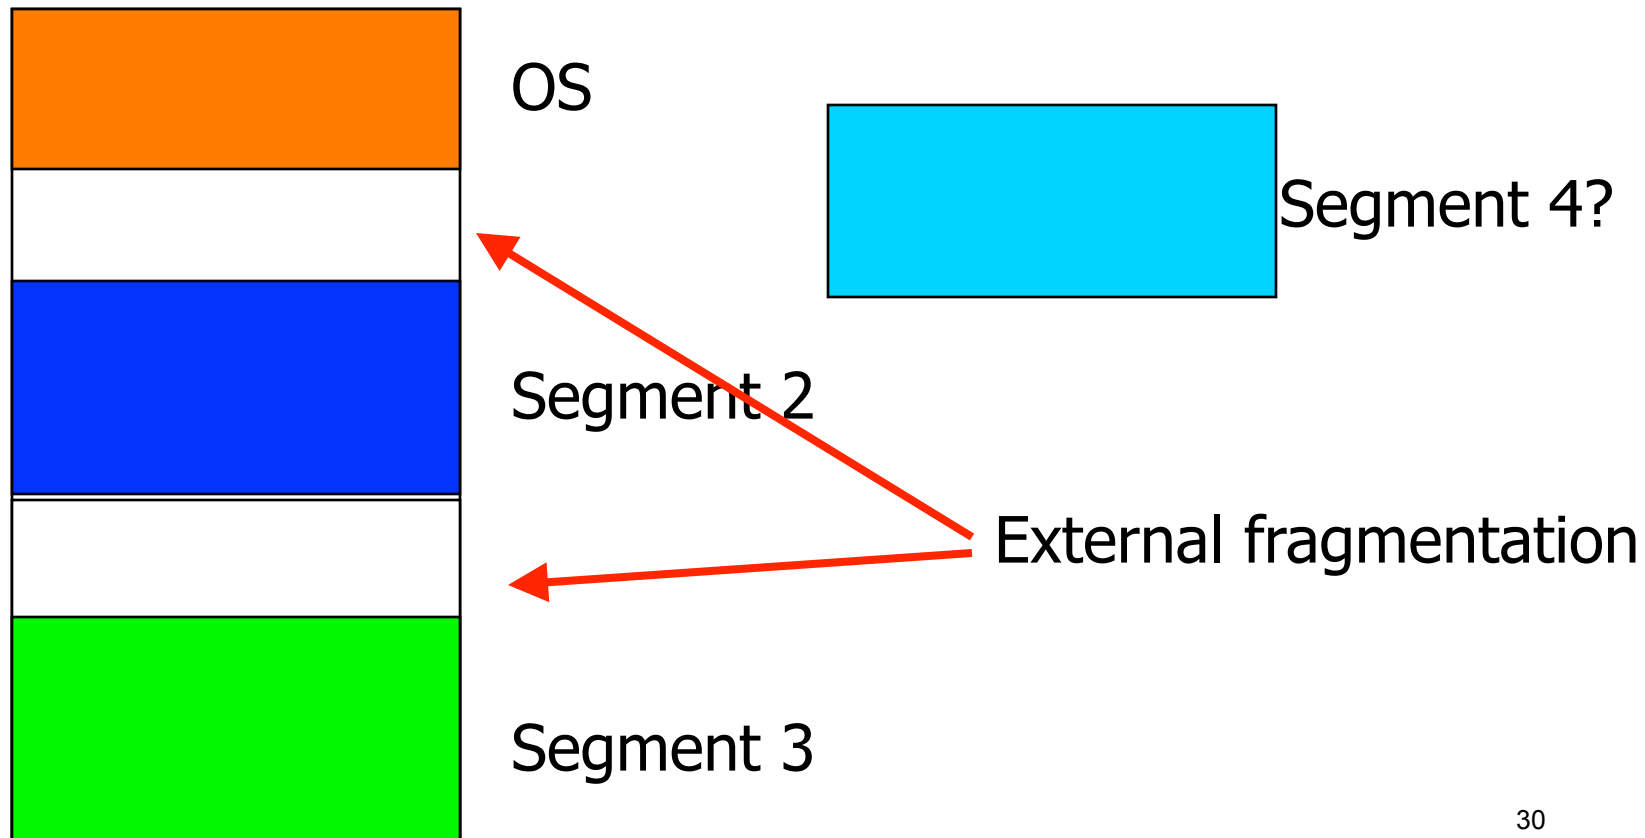
➔ Where is 15DC in physical memory?



# Pros/cons of segmentation

- Pros:
  - Process can be split among several segments
    - Allows sharing
  - Segments can be assigned, moved, or swapped independently
- Cons:
  - **External fragmentation**: many holes in physical memory
    - Also happens in base and bound scheme

# Simple multiprogramming: Single segment per process, static relocation



# What fundamentally causes external fragmentation?



1. Segments of many different sizes
2. Each has to be allocated contiguously

# Dynamic memory allocation problem



- Problem: External fragmentation caused by holes too small
- How much can a smart allocator help?
  - The allocator maintains a free list of holes
  - Allocation algorithms differ in how to allocate from the free list



# Dynamic allocation algorithms



- Best fit: allocate the smallest chunk big enough
- First fit: allocate the first chunk big enough
  - Rotating first fit
- Is best fit necessarily better than first fit?
  - Example: 2 free blocks of size 20 and 15
  - If allocation ops are 10 then 20, which one wins?
  - If ops are 8, 12, then 12, which one wins?

# Dynamic allocation algorithms



- Analysis shows
  - First fit tends to leave average-size holes
  - Best fit tends to leave some very large holes, very small holes
- Knuth claims that if storage is close to running out, it will run out regardless of which scheme is used
  - Pick the easiest or most efficient (e.g. first fit)


# Segmentation: OS implementation



- Keep segment table in PCB
- When creating process, allocate space for segments, fill in PCB bases/bounds
- When process dies, return physical space used by segments to free pool
- Context switch?
  - Saves old segment table / Loads new segment table
  - What about context switch of threads?
  - True-or-false: CS between threads of same process cheaper than CS between processes

# [lec2] Kernel data structure:

## Process Control Block (Process Table)




- Process management info
  - State (ready, running, blocked)
  - PC & Registers, parents, etc
  - CPU scheduling info (priorities, etc.)
- Memory management info
  - Segments, page table, stats, etc
- I/O and file management
  - Communication ports, directories, file descriptors, etc.



# Managing segments (cont)

To enlarge a segment:


- See if space above segment is free. If so, just update the bound and use that space
- Or, move this segment to disk and bring it back into a larger hole (or maybe just copy it to a large hole)




# Managing segments (cont)

- When there is no space to allocate a new segment:
  - Compact memory – how?

# Summary: Evolution of Memory Management (so far)



Scheme	How	Pros	Cons
Simple uniprogramming	1 segment loaded to starting address 0	Simple	1 process 1 segment No protection
Simple multiprogramming	1 segment relocated at loading time	Simple, Multiple processes	1 segment/process No protection External frag.
Base & Bound	Dynamic mem relocation at runtime	Simple hardware, Multiple processes Protection	1 segment/process, External frag.
Multiple segments	Dynamic mem relocation at runtime	Sharing, Protection, multi segs/process	More hardware, External frag.



# Dynamic storage allocation

- Why isn't static allocation sufficient?
- Need dynamic memory allocation for
  - both main memory
  - file space on disk



# Dynamic storage allocation: two basic operations



- Allocate
- Free

# Dynamic storage allocation: two general ways



- Stack
  - Restricted
  - Simple and efficient
- Heap
  - More general
  - Less efficient
  - More difficult to implement



# Stack organization

- Memory allocation & freeing are predictable
  - “*We do better when we can predict future*”
  - Example:
    - Procedure call
- Allocation is LIFO
- Stack keeps all the free space together



# Heap organization

- Allocation & freeing are unpredictable
- For arbitrary, complex data structures
  - Example: payroll system
    - Don't know when employee will join or leave the company
    - Must keep track of all of them
- Memory consists of allocated areas and free areas (holes) → lots of holes inevitable
- Goal: keep # of holes small, size large



# Heap organization

- *Fragmentation*: inefficient use of memory due to holes too small
  - What happens in stack? Do we have fragmentation there?
- Typically, heap allocation uses a *free list* of holes
- Allocation algorithms differ in how to manage the free list




# Heap allocation

- Best fit
- First fit
  - Rotating first fit
- Is best fit necessarily better than first fit?
  - Example: 2 free blocks of size 20 and 15
  - If allocation ops are 10 then 20, which one wins?
  - If ops are 8, 12, then 12, which one wins?



# Heap allocation

- First fit tends to leave average-size holes
- Best fit tends to leave some very large holes, very small holes
- Knuth claims that if storage is close to running out, it will run out regardless of which scheme is used  
→ Pick the easiest or most efficient (first fit)



# Implementation

- Bit map
  - For fixed-size chunks (e.g., disk blocks)
- Pools
  - A separate allocation pool for each popular size
  - Fast, no fragmentation
  - But some pools may run out faster than others





# Reclamation

- When can dynamically-allocated memory be freed?
  - Easy if a chunk is used in one place
  - Hard when a chunk is shared
  - Sharing is indicated by presence of pointers to the data



# Reclamation techniques

- Reference counts:
  - Keep track of the number of outstanding pointers to each chunk of memory
  - When this goes to 0, free the memory
  - Example:
    - File descriptors in UNIX
  - Works fine with hierarchical structures
  - What about circular structures?



# Reclamation techniques

- Garbage collection
  - Storage isn't freed explicitly (using free), rather implicitly, i.e., by deleting pointers
  - When the system needs storage, it scans through all pointers (all of them!!!) and collects things not used
  - For circular structures, this is the only way
  - Makes life easier on programmers, but GCs are hard to implement



# Reclamation techniques

- Garbage collection – implementation
  - Must be able to find all objects
  - Must be able to find all pointers to objects
  - Pass1: mark
    - Go through all statically-allocated and procedure local variables, looking for pointers
    - Mark each obj pointed to, and recurs
    - Compiler has to help by saving info about pointers with structures
  - Pass 2: sweep
    - Go through all objs, free up those that aren't marked