

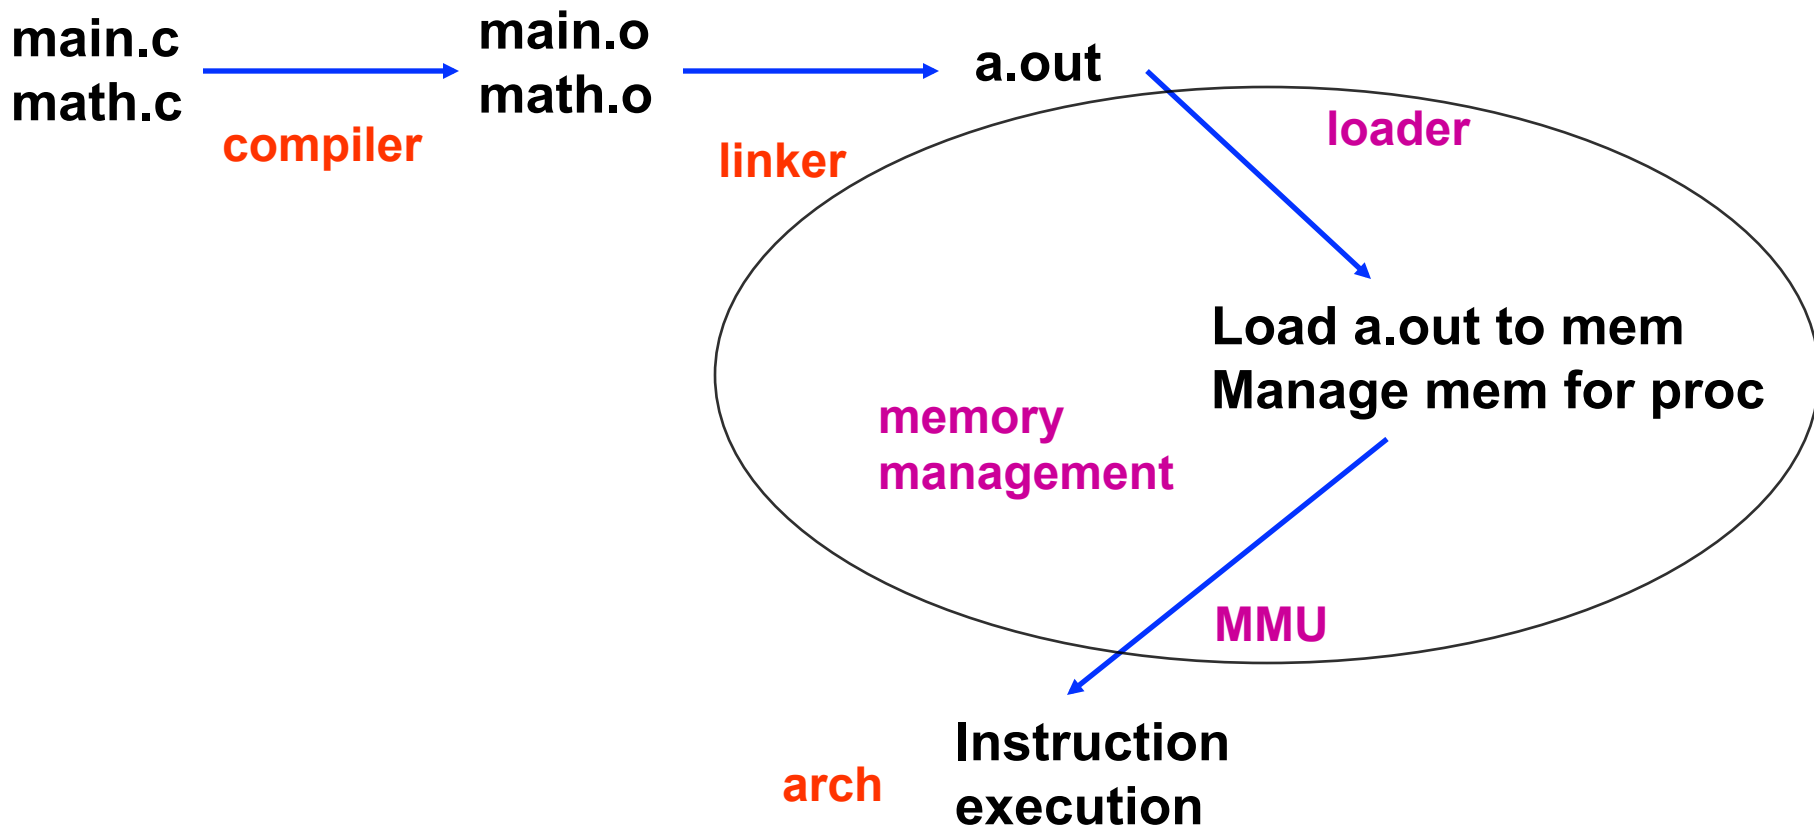
Paging, Inverted Page Table, TLB

ECE469, Feb 28

Yiying Zhang




The big picture




Today's topics

- Basic paging [ctnd]
- Inverted page table
- TLB




[lec13] Segmentation




- ◆ A table of (seg, size)
- ◆ Protection: each entry has
 - (nil, read, write, exec)
- ◆ On a context switch: save/restore the table or a pointer to the table to/from PCB

- 5

[review] Virtual pages / physical pages




Virtual address




Virtual pages

Physical memory

physical pages



[review] Paging




- Context switch
 - similar to the segmentation scheme
- Pros:
 - easy allocation, keep a free list
 - easy to swap
 - easy to share

[Review] How many PTEs do we need?

(assume page size is 4096 bytes)

- Worst case for 32-bit address machine
 - # of processes $\times 2^{20}$
- What about 64-bit address machine?
 - # of processes $\times 2^{52}$




Deep thinking



- In segmentation, why does each segment need to be contiguous in physical memory?
- In segmentation, what to do with heap/stack?
 - What happens when they grow/shrink?
- In paging, do pages belonging to the same “segment” (e.g. heap) need to be contiguous in physical memory?
 - What made this possible?
 - What to do with heap/stack growing/shrinking now?



Two-level page tables



What does this buy us?

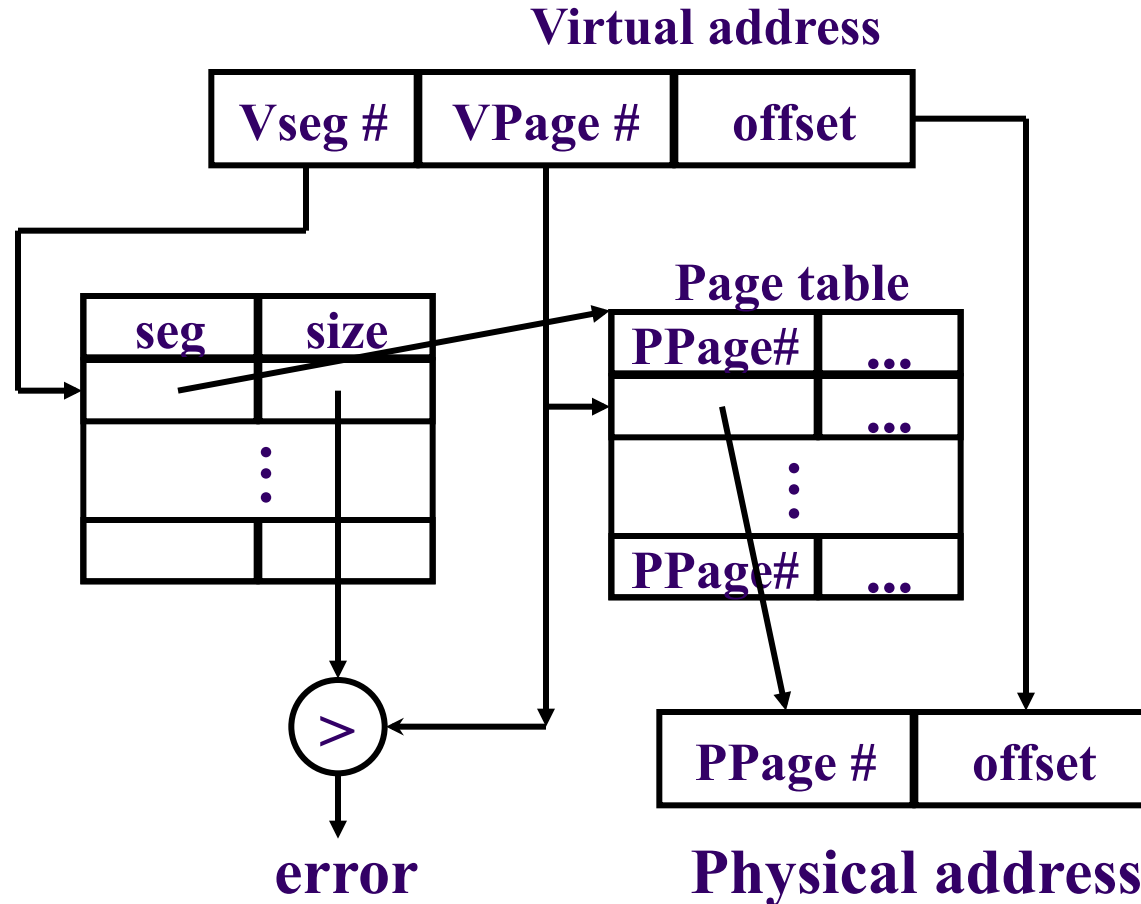


Multi-level page tables

- 3 Advantages?
 - L2 page tables do not have to be consecutive
 - They do not have to be allocated before use!
 - They can be swapped out to disk!

The power of an extra level of indirection!

Segmentation with paging



Ex: IBM System 370 (24-bit, 4-bit segment #, 8-bit page #)



Segmentation + paging

- Use **two levels of mapping** to make tables manageable:
 - Each segment contains one or more pages
 - Segments correspond to **logical units**: code, data, stack
 - Segments vary in size and are often large
 - Pages are for easy of management by OS: fixed size -> easy to allocate/free
- Going from P to P+S is like going from single segment to multiple segments, except at a high level
 - One page table -> many page tables with bases/bounds

[lec1] Separating Policy from Mechanism



Mechanism – tool to achieve some effect

Policy – decisions on how to use tool

examples:

- All users treated equally
- All program instances treated equally
- Preferred users treated better

Separation leads to flexibility

Segmentation + paging vs. multi-level paging




- *Mechanisms* are similar
- Difference lies in *policy*
 - Segmentation + paging still maintains notion of **segments**
 - Multi-level paging deals **the whole, uniform address space** (like one-level paging)

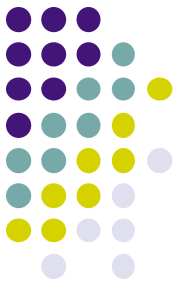
The Intel Pentium (1993)

(pro, II, III, 4) (Ch 8.7, fig 8.22, 8.23)



- Supports both pure segmentation and segmentation with 1-level paging (page size=4M) or 2-level paging (page size=4k)
- CPU generates logical addresses
 - (selector, offset), 16 bits and 32 bits
 - As many as 16K segments
 - Up to 4GB per segment






Linux on Pentium

- Linux uses 3-level paging
 - For both 32-bit and 64-bit architectures
- On Pentium, degenerates to 2-level paging
 - Middle-level directory has zero bits

Today's topics

- Basic paging
- Inverted page table
- TLB



[review] How many PTEs do we need?



- Worst case for 32-bit address machine
 - # of processes $\times 2^{20}$ (if page size is 4096 bytes)
- What about 64-bit address machine?
 - # of processes $\times 2^{52}$

Hmm, but my PC only has 1GB, 256K PTEs should be enough?!




Inverted Page Table

- Motivation
 - Example: 2 processes, page table has 1M entries, 10 phy pages
- Is there a way to save page table space?



Ideally,



- One PTE for each physical page frame, disregarding how many processes
 - Assuming rest virtual addressed not allocated/used
 - i.e., linear inverted page table (an array of phy pages mapped to virtual addresses)






But,

- How do we do lookups with linear inverted page table?
 - Has to go through the entire array and compare!



(Hashed) Inverted page tables



- Main idea
 - One PTE for each physical page frame
 - Hash (Vpage, pid) to Ppage#
- Pros
 - Small page table for large address space
- Cons
 - Lookup is difficult
 - Overhead of managing hash chains, etc
- Ex: 64-bit UltraSPARC₂₃, PowerPC



Deep thinking

- How can two processes share memory under inverted page table?
 - Since the inverted page table provides only one forward mapping, it is very difficult to share memory among processes. For this reason most modern OSs use multi-level page tables.




break

- There is four digit number in aabb form and it is a perfect square. Find out the number.

Today's topics

- Basic paging
- Inverted page table
- TLB

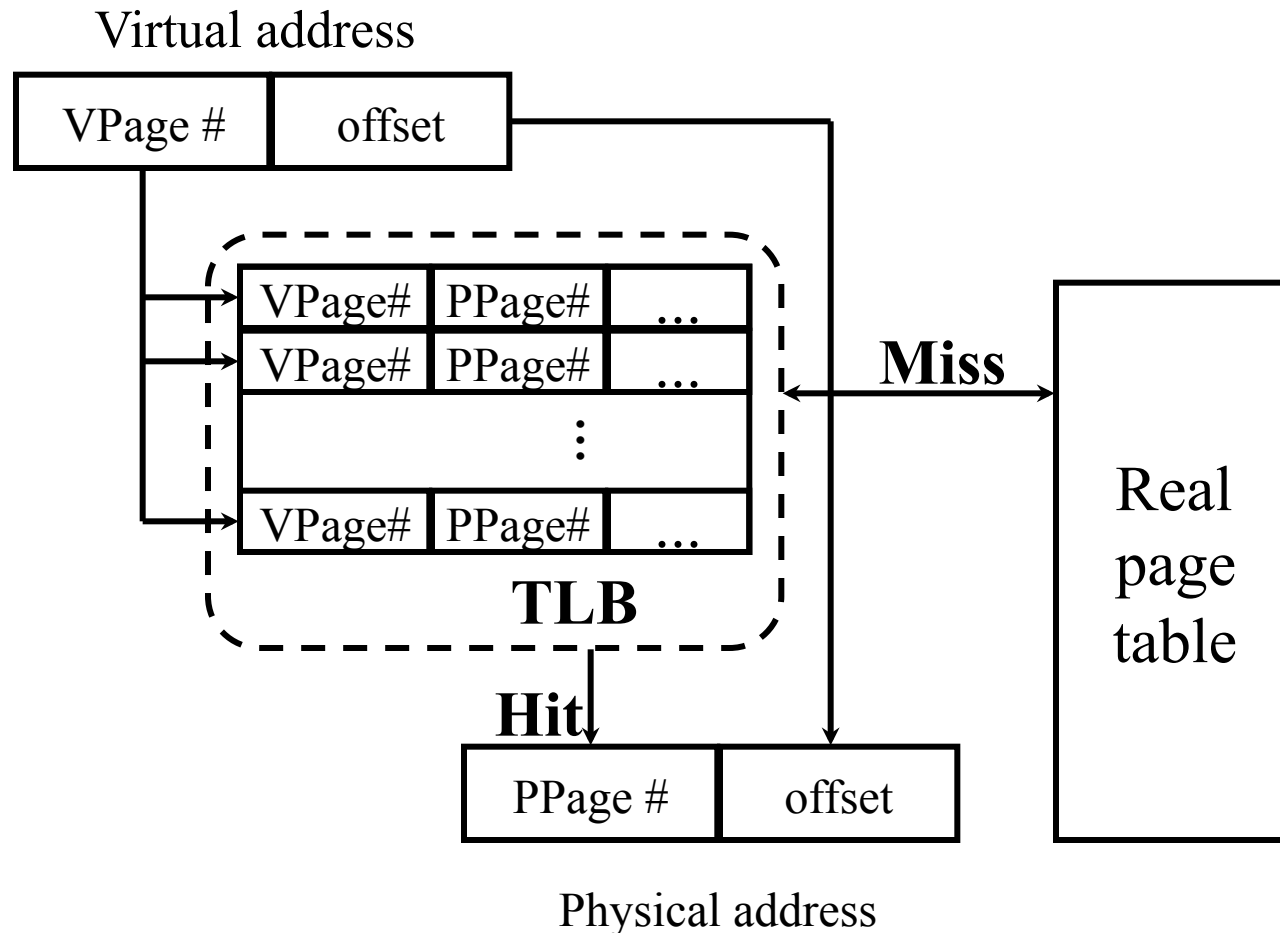


Performance problem with paging



- How many extra memory references to access page tables?
 - One-level page table?
 - Two-level page table?
- Solution: *reference locality!*
 - In a short period of time, a process is likely accessing only a few pages
 - Instead of storing only page table starting address in hardware (MMU), store part of the page table that is “hot”

Translation Look-aside Buffer (TLB)




TLB often fully set-associative → least conflict misses
Expensive → typically 64 – 1024 entries

Bits in a TLB Entry

VPage#	PPage#	...
VPage#	PPage#	...
⋮		
VPage#	PPage#	...

TLB



- Common (necessary) bits
 - Virtual page number: match with the virtual address
 - PTE
- Optional (useful) bits
 - ASIDs -- Address-space identifiers (process tags)

Miss handling:

Hardware-controlled TLB



- On a TLB hit, MMU checks the valid bit
 - If valid, perform address translation
 - If invalid (e.g. page not in memory), MMU generates a page fault
 - OS performs fault handling
 - Restart the faulting instruction
- On a TLB miss
 - MMU parses page table and loads PTE into TLB
 - Needs to replace if TLB is full
 - PT layout is fixed
 - Same as hit ...

Miss handling:

Software-controlled TLB



- On a TLB hit, MMU checks the valid bit
 - If valid, perform address translation
 - If invalid (e.g. page not in memory), MMU generates a page fault
 - OS performs page fault handling
 - Restart the faulting instruction
- On a TLB miss, HW raises exception, **traps to the OS**
 - **OS** parses page table and loads PTE into TLB
 - Needs to replace if TLB is full
 - PT layout is **flexible**
 - Same as in a hit...

Hardware vs. software controlled

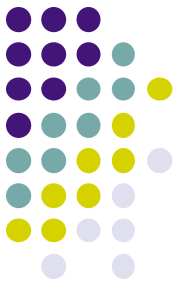


- Hardware approach

- Efficient – TLB misses handled by hardware
- OS intervention is required only in case of page fault
- Page structure prescribed by MMU hardware -- rigid

- Software approach

- Less efficient -- TLB misses are handled by software
- MMU hardware very simple, permitting larger, faster TLB
- OS designer has complete flexibility in choice of MM data structure
 - e.g. 2-level page table, inverted page table



Deep thinking

- Without TLB, how MMU finds PTE is fixed
- With TLB, it can be flexible, e.g. software-controlled is possible
- What enables this?
- TLB is an extra level of indirection!

More TLB Issues



- Which TLB entry should be replaced?
 - Random
 - LRU
- What happens when changing a page table entry (e.g. because of swapping, change read/write permission)?
 - Change the entry in memory
 - flush (eg. invalidate) the TLB entry
 - INGLPG on x86

What happens to TLB in a process context switch?



- During a process context switch, cached translations can not be used by the next process
 - Invalidate all entries during a context switch
 - Lots of TLB misses afterwards
 - Tag each entry with an ASID
 - Add a HW register that contains the process id of the current executing process
 - TLB hits if an entry's process id matches that reg




Cache vs. TLB

- Similarities:
 - Both cache a part of the physical memory
- Differences:
 - Associatively
 - TLB is usually fully associative
 - Cache can be direct mapped
 - Consistency
 - TLB does not deal with consistency with memory
 - TLB can be controlled by software

More on consistency Issues




- Snoopy cache protocols can maintain consistency with DRAM, even in the presence of DMA
- No hardware maintains consistency between DRAM and TLBs:
 - OS needs to flush related TLBs whenever changing a page table entry in memory
- On multiprocessors, when you modify a page table entry, you need to do “*TLB shoot-down*” to flush all related TLB entries on all processors



Memory Hierarchy Revisited

What does this
imply about L1
addresses?

Where do we hope
requests get satisfied?

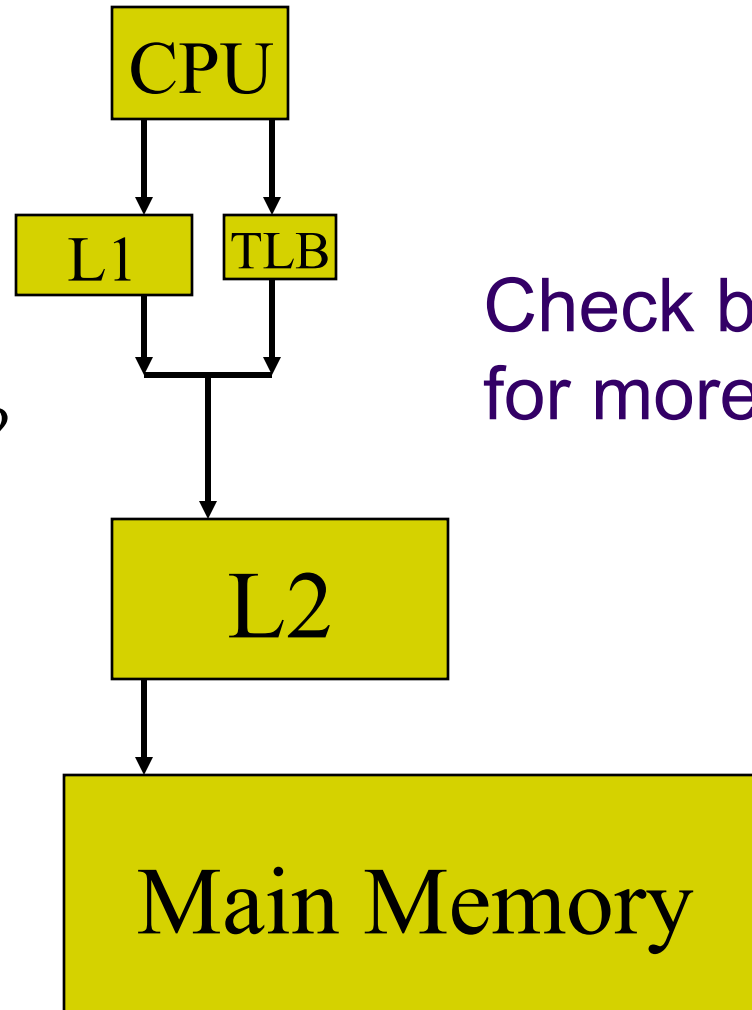




Memory Hierarchy Re-Revisited


What does this
imply about L1
addresses?

Any speed benefits?
Any drawbacks?



Check backup slides
for more details!

What about the kernel itself?

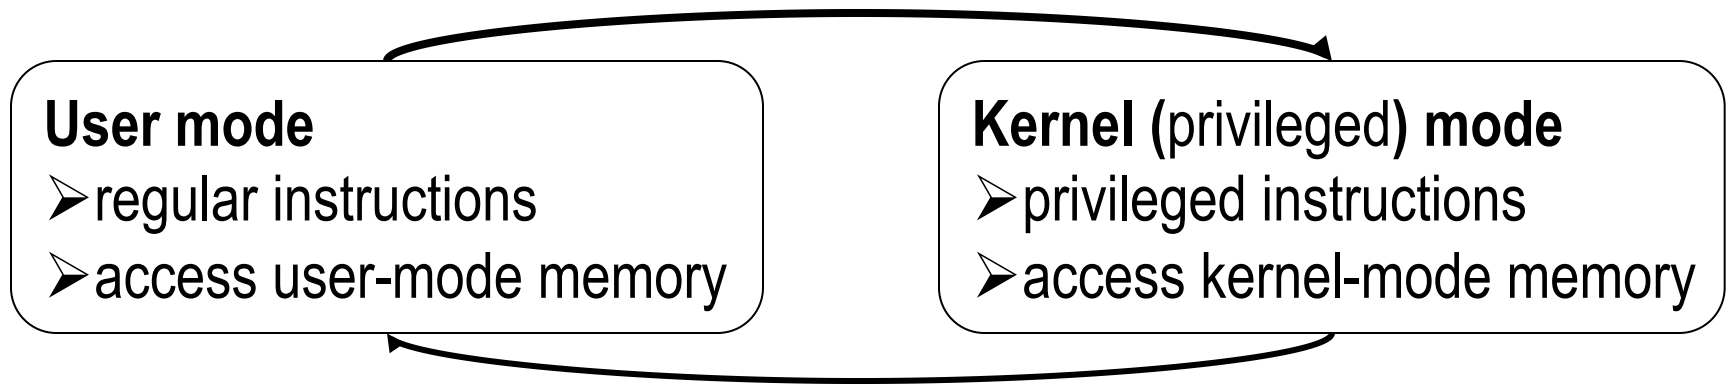


- What happens when OS is running?
 - OS runs with relocation turned off (a bit in processor status word (PSW) controls relocation)
- How to prevent users from controlling base & bound, relocation?
- Does kernel need multiple address spaces?
- How does OS regain control?
Need to atomically
 - Branch into/out of OS
 - Turn relocation on/off

Support in modern processors: User \Leftrightarrow Kernel



An interrupt or exception (INT)



A special instruction (IRET)

Privileged instructions



- Special Instructions
 - system call (invoked by user program)
 - memory mapping, TLB, etc.
 - device registers
 - I/O operations

Hardware support for Modes



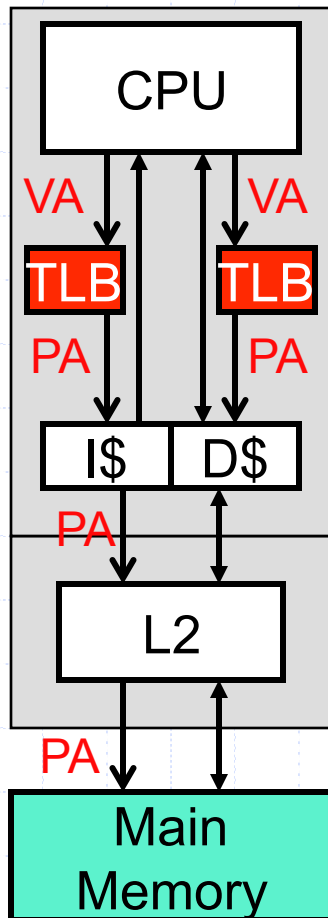
- *Mode bit* added to hardware to indicate the current mode:
 - kernel (0)
 - user (1)



Backup

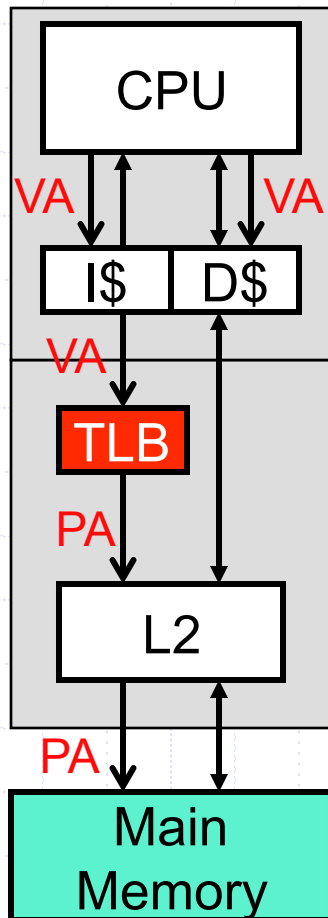
- Backup slides will not be in exams, but can help you understand topics that will be in exams

Physical (Address) Caches



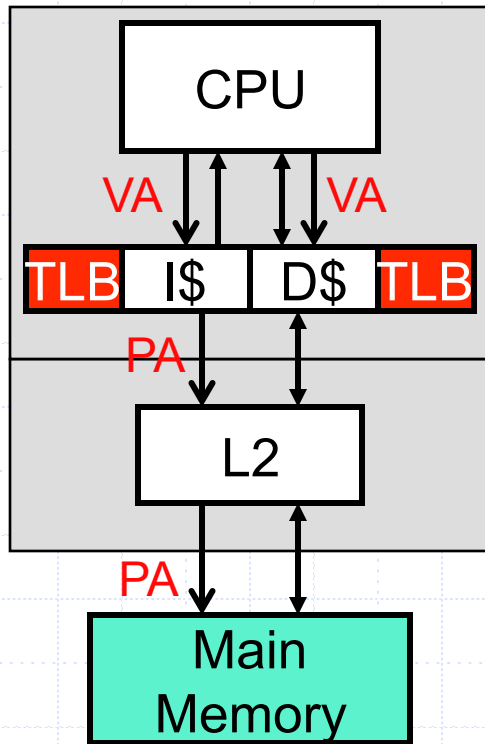
- Memory hierarchy so far: **physical caches**
 - Indexed and tagged by PAs
 - **Physically Indexed (PI)**
 - **Physically Tagged (PT)**
 - Translate to PA to VA at the outset
- + Cached inter-process communication works
 - Single copy indexed by PA
- Slow: adds at least one cycle to t_{hit}

Virtual Address Caches (VI/VT)



- Alternative: **virtual caches**
 - Indexed and tagged by VAs (VI and VT)
 - Translate to PAs only to access L2
 - + Fast: avoids translation latency in common case
 - Problem: VAs from **different processes** are distinct physical locations (with different values) (call **homonyms**)
- What to do on process switches?
 - Flush caches? Slow
 - Add process IDs to cache tags
- Does inter-process communication work?
 - **Synonyms**: multiple VAs map to same PA
 - Can't allow same PA in the cache twice
 - Also a problem for DMA I/O
 - Can be handled, but very complicated

Parallel TLB/Cache Access (VI/PT)

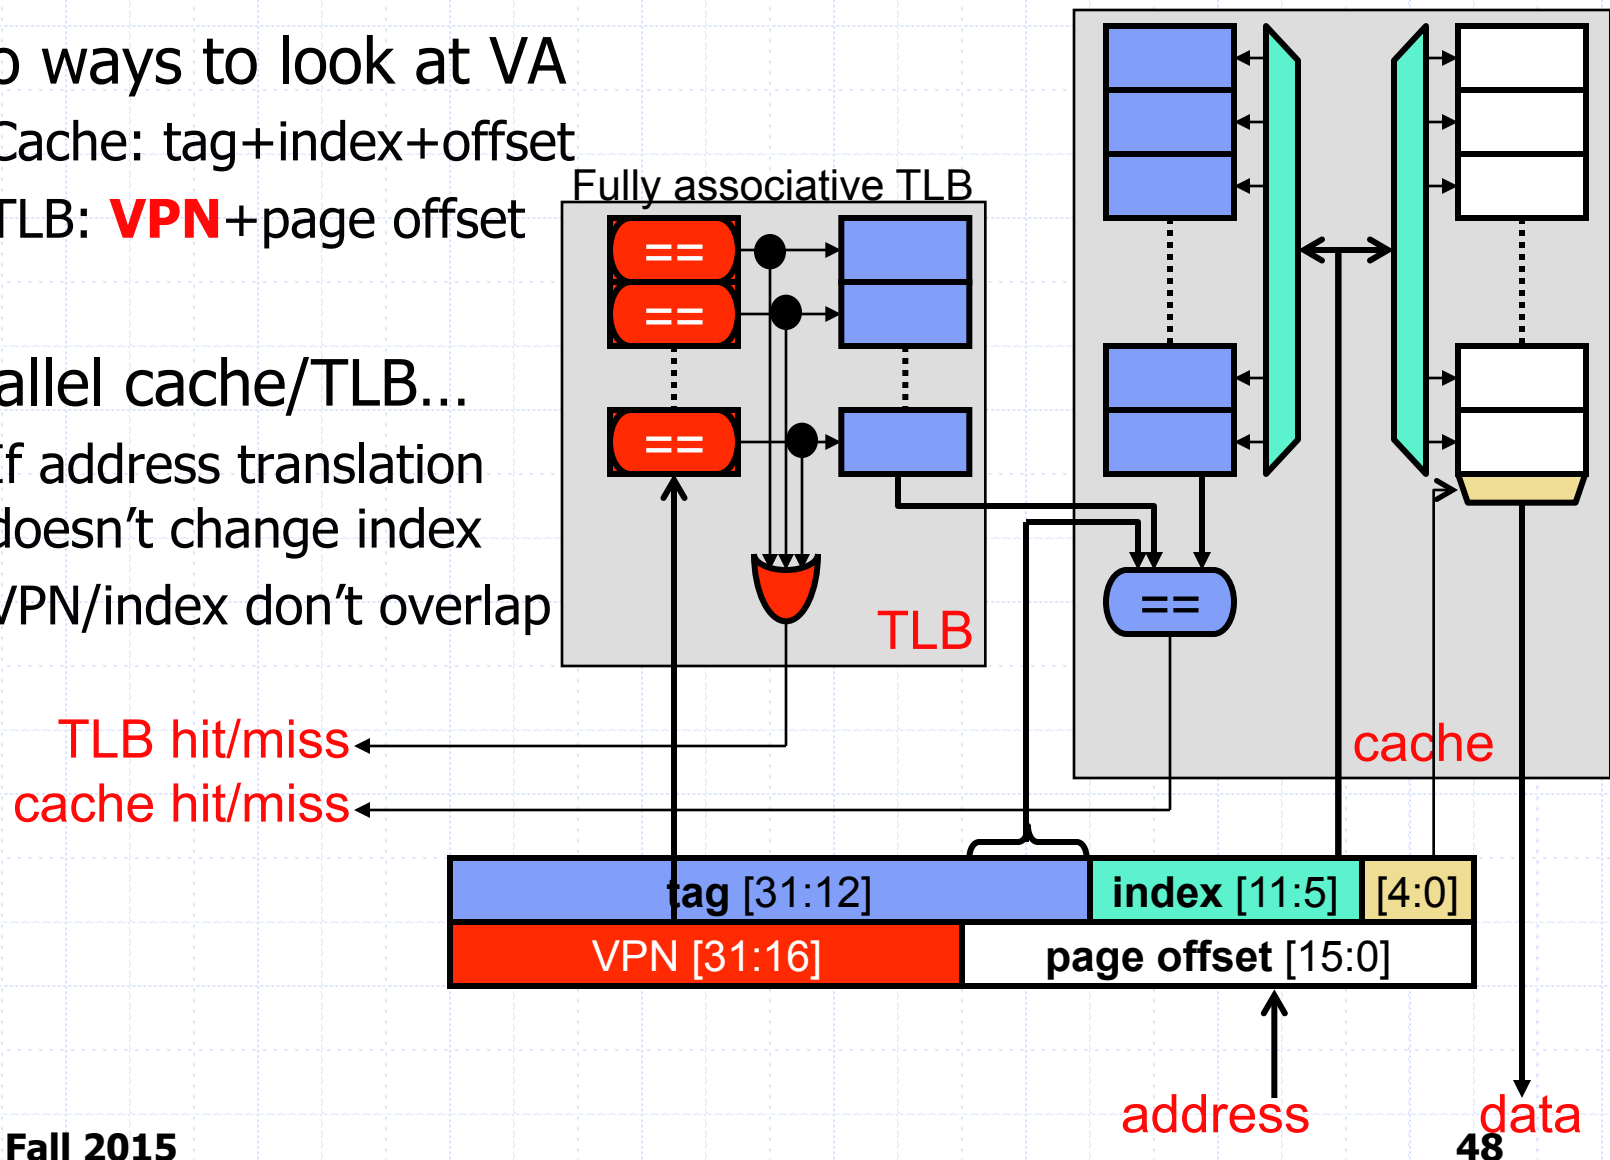


Compromise: **access TLB in parallel**


- *In small caches, index of VA and PA the same*
 - $VI == PI$
- Use the VA to index the cache
- Tagged by PAs
- Cache access and address translation in parallel
- + No context-switching/aliasing problems
- + Fast: no additional t_{hit} cycles
- Common organization in processors today

Parallel Cache/TLB Access

- Two ways to look at VA
 - Cache: tag+index+offset
 - TLB: **VPN**+page offset
- Parallel cache/TLB...
 - If address translation doesn't change index
 - VPN/index don't overlap



Cache Size And Page Size



- Relationship between page size and L1 cache size
 - Forced by non-overlap between VPN and IDX portions of VA
 - Which is required for TLB access
 - **Rule: (cache size) / (associativity) ≤ page size**
 - Result: associativity increases allowable cache sizes
 - Systems are moving towards bigger (64KB) pages
 - To use parallel translation with bigger caches
 - To amortize disk latency
 - Example: Pentium 4, 4KB pages, 8KB, 2-way SA L1 data cache