


Page Replacement, Midterm Review 2

ECE469, March 24

Yiying Zhang




Practice Exam



break




Linus Torvalds



- Linux
- git


Andrew Tanenbaum



- Minix
- OS textbooks

Reading

- Chapter 9




[lec16] The Policy Issue: Page selection and replacement



- Once the hardware has provided basic capabilities for VM, the OS must make two kinds of decisions
 - Page selection: *when* to bring pages into memory
 - Page replacement: *which* page(s) should be thrown out
 - Try to kick out the least useful page

[lec 17] Page replacement problem



Definition:

- Expect to run with all physical pages in use
- Every “page-in” requires an eviction to swap space
- How to select the victim page?
- Performance goal:
 - Give a sequence of memory accesses, minimize the # of page faults
 - given a sequence of virtual page references, minimize the # of page faults
- Intuition: kick out the page that is least useful
- Challenge: how do you determine “least useful”?
 - Even if you know future?



Least Recently Used (LRU)

- Algorithm
 - Replace page that hasn't been used for the longest time
- Advantage: with locality, LRU approximates Optimal
- Question
 - What hardware mechanisms are required to implement *exact* LRU?



Implementing LRU: hardware

- A counter for each page
- Every time page is referenced, save system clock into the counter of the page
- Page replacement: scan through pages to find the one with the oldest clock
- Problem: have to search all pages/counters!

Least recently used

[illegible]

--	--

0	1	2	3	\dots	254	255
---	---	---	---	---------	-----	-----

11

Approximate LRU



Interval 1	Interval 2	Interval 3	Interval 4	Interval 5
00000000	00000000	10000000	01000000	10100000
00000000	10000000	01000000	10100000	01010000
10000000	11000000	11100000	01110000	00111000
00000000	00000000	00000000	10000000	01000000

- Algorithm
 - At regular interval, OS shifts reference bits (in PTE) into counters (and clear reference bits)
 - Replacement: Pick the page with the “smallest counter”
- How many bits are enough?
 - In practice 8 bits are quite good
- Pros: Require one reference bit, small counter/page
- Cons: Require looking at many counters (or sorting)



Implementing LRU: software

- A doubly linked list of pages
- Every time page is referenced, move it to the front of the list
- Page replacement: remove the page from back of list
 - Avoid scanning of all pages
- Problem: too expensive
 - Requires 6 pointer updates for each page reference info
 - High contention on multiprocessor



Not Recently Used (NRU)

- Algorithm
 - Randomly pick a page from the following (in this order)
 - Not referenced and not modified
 - Not referenced and modified
 - Referenced and not modified
 - Referenced and modified
- Pros
 - Easy to implement
- Cons
 - No frequency → Not very good performance
 - Takes time to classify



Thinking

- Performance design philosophy:
- Make common case fast!
- Amadhl's Law

$$\text{Overall Speedup} = \frac{1}{(1 - F) + \frac{F}{S}}$$

F = The fraction enhanced


S = The speedup of the enhanced fraction

Factors that affect paging performance



- # of memory misses
 - So far, have talked about algorithms to reduce misses (increase hit rate)
- Cost of a memory miss (replacement)

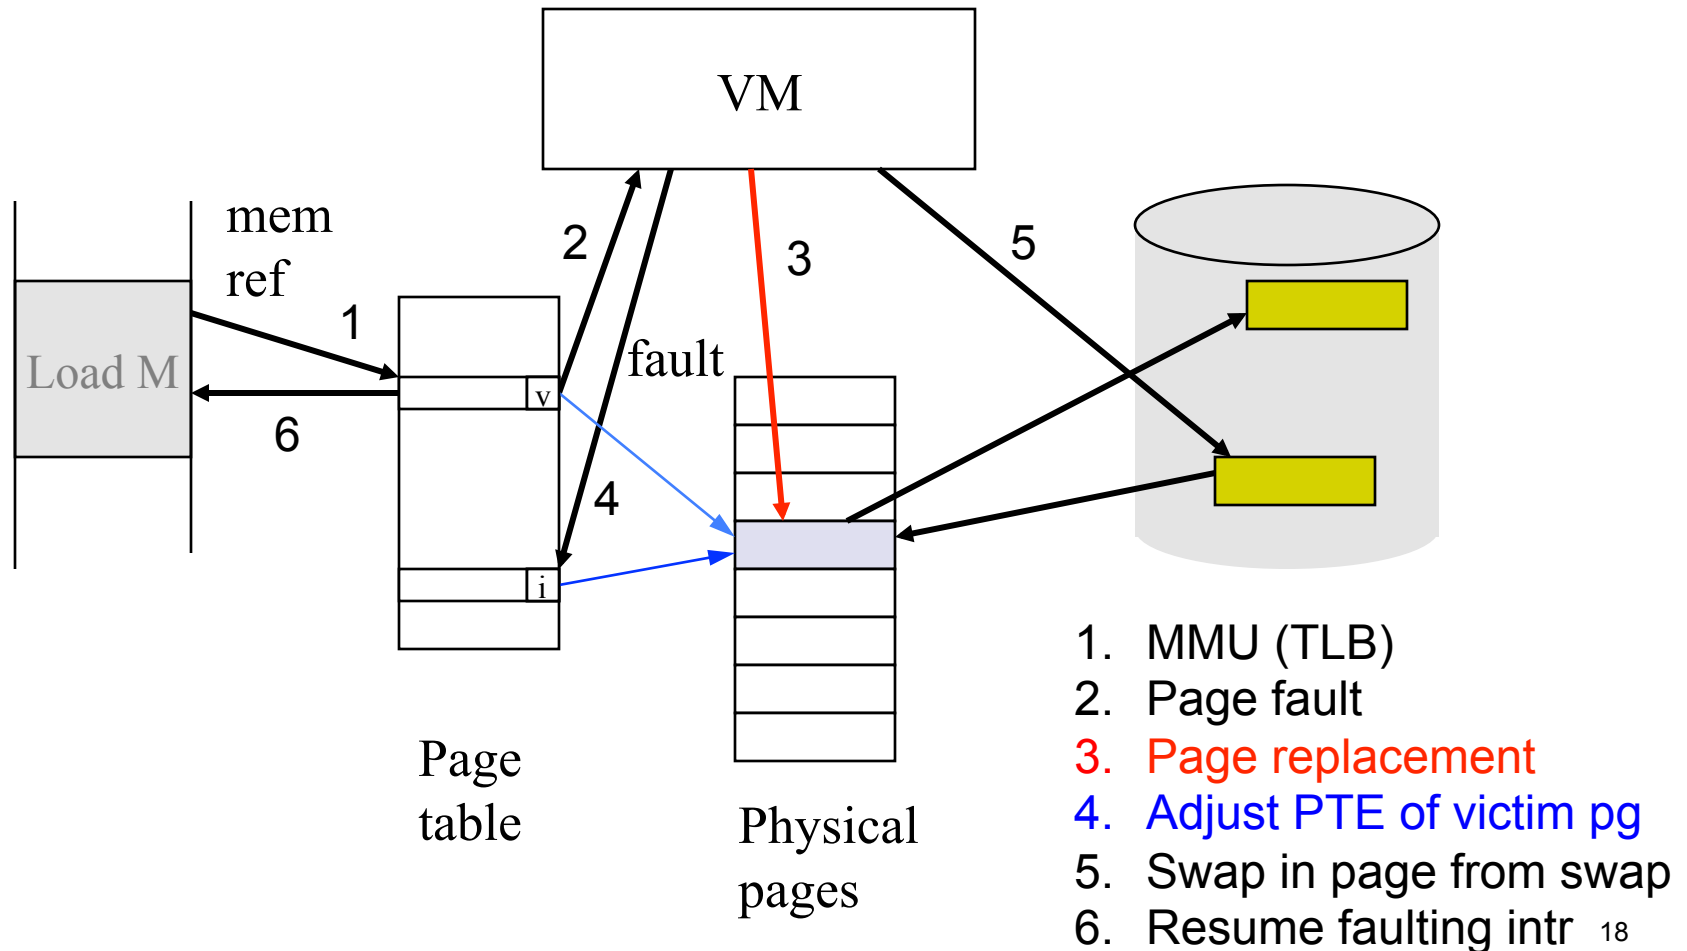
[lec17] Enhanced FIFO with 2nd-Chance Algorithm (used in Macintosh VM)



- Same as the basic FIFO with 2nd chance, except that it considers both (reference bit, modified bit)
 - (0,0): neither recently used nor modified (good)
 - (0,1): not recently used but dirty (not as good)
 - (1,0): recently used but clean (not good)
 - (1,1): recently used and dirty (bad)
 - When giving second chance, only clear reference bit
- Pros
 - Avoid write back
- Cons
 - More complicated, worse case scans multiple rounds



[lec16] Page Fault Handling in demand paging



Page out on critical path?



- If no free page pool, page in has to wait till page out is finished
 - Page fault handling time = proc. overhead + 2 * I/Os
- There is a chance of paged out page being referenced soon

Page buffering techniques



1. System maintains a pool of free pages
 - When a page fault occurs, victim page chosen as before
 - But desired page paged into a free page right away before victim page paged out
 - Dirty victim page written out when disk is idle
2. Also remember what was in free pages - maybe reused before reallocated