


Course Review

ECE469

Apr 25

Yiying Zhang




ABET




- ABET requirements and remedy homework online
 - Requirements under each objective have an **OR** relationship
- ABET outcome is orthogonal to your final grades
 - Passing ABET doesn't necessarily mean you pass the course (but in most cases it does)
 - Failing ABET however will most likely result in failing the course
- I will posted tentative ABET outcomes on blackboard based on your progress so far.
- If your outcome is fail (0 for fail, 1 for pass), **come talk to me** to see what you need to achieve in lab 5 **or** final **or** ABET remediation homework to pass ABET requirements.
- My goal is to pass as many as possible in ABET


Plan



- Final exam: HAMP 1144, May 3, 7-9pm
- Office hours: today until noon, Thur until noon
- Course evaluation close on May 1st
- 2nd half semester review
- Course summary (annotated slides from lec1)


Course review – 2nd half semester





Page Replacement Policies


- Optimal
- Random
- FIFO
 - Belady's anomaly
- Approximate LRU, NRU
- FIFO with 2nd chance
- Clock: a simple FIFO with 2nd chance



More Virtual Memory


- Thrashing
- Working set model, size, replacement algorithm
- Shared memory
- COW

Virtual Memory Questions to Think About




- What is the use of optimal algo?
- If future is unknown, what make us think there is a chance for doing a good job?
- Without addi. hardware support, the best we can do?
- What is the minimal hardware support under which we can do a decent job?
- Why is it difficult to implement exact LRU? Exact anything
- For a fixed replacement algo, more page frames → less page faults?
- How can we move page-out out of critical path?

More Virtual Memory Questions




- Per-process vs. global page replacement
- Thrashing
- What causes thrashing?
- What to do about thrashing?
- What is working set?
- What's the benefit of Copy-on-Write?




Practice Question (Quiz 2)

- Consider the following virtual page reference string: 1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5.
- How many page faults will there be under the LRU replacement algorithm on Nick's PC which has 4 physical pages? How many page faults will there be on Riley's machine which has 3 physical pages?
- Nick's: 8
- Riley's: 10


Storage and File System






Storage Device

- Disk Internals
 - Seek/rotation, random/sequential accesses
- SSD Internals
 - Flash read/write/erase, the granularity of them
 - Erase-before-write, flash wear
 - FTLs
 - Garbage collection and wear leveling
- Both Disks and SSDs can fail!




RAID

- Two motivations
 - Performance, reliability
 - Two main ideas
 - Striping, mirroring (parity)
 - RAID levels: 1-6
- no RAID level 2 3 in exam



Practice question


- Assume that
 - you have a mixed configuration comprising disks organized as RAID Level 1 and as RAID Level 5 disks;
 - the system has flexibility in deciding which disk organization to use for storing a particular file.
 - you have a mixed workload of *frequently-read* and *frequently-written* files
- Which files should be stored in the RAID Level 1 disks and which in the RAID Level 5 disks in order to optimize performance?




File System Overview

- File system abstraction
 - File, directory, FS APIs
- Directories
 - Different directory organizations
 - Directory internals
 - Path walk
 - Hard link, soft link
- Metadata
 - Inode

File Allocation



- Two tasks:
 - How to allocate blocks for a file?
 - How to design inode to keep track of blocks?
- Allocation methods:
 - Contiguous
 - Extent-based
 - Linked
 - File-allocation Tables
 - Indexed
 - Multi-level Indexed
- Free space management
 - linked list
 - bitmap




UFS

- UFS
 - multi-level indexed files
 - Inodes all stored on outermost track
 - Free-block linked list
 - Two problems of UFS
 - data blocks are allocated randomly in aging file systems
 - inodes are allocated far from blocks

Indirect blocks addressing ranges




- Assume blocksize = 1K
 - a block contains $1024 / 4 = 256$ block addresses
- direct block address: 10K
 - indirect block addresses: 256K
 - double indirect block addresses: $256 * 256K = 64M$
 - tripe indirect block addresses: $256 * 64M = 16G$



FFS and LFS




- FFS
 - Cylinder group
 - Inodes and data within same dir
 - Free block managed with bitmap
- File system buffer cache
- LFS
 - All writes buffered into chunks and go to an append-only log
 - Locating data is more difficult
 - Needs background cleaner




Journaling File System

- Data reliability
 - Three threats: three directions to prevent data loss
- The problem with write back cache under system crash
 - One file system operation consists of multiple sub-operations, they should be atomic
- Journaling
 - Write to a redo log first, in a transactional way
 - Journal checkpointing, crash recovery
 - Ext3: three journaling modes




Distributed File System

- DFS and client/server model
- NFS
 - Transparency through indirection of VFS
 - Two key ideas for fast crash recovery
 - Stateless server
 - Idempotent operations
 - Client cache and cache consistency



Course Summary

(annotated slides from Lecture 1)




[lec1] What is an OS?

“Code” that *sits between*:

- programs & hardware
- different programs
- different users

But what does it do/achieve?




[lec1] What is an OS?

- Resource manager
- Extended (abstract) machine
- A giant interrupt handler!

Makes computers efficient and simple to use

[lec1] Separating Policy from Mechanism



Policy – decisions on how to use tool

Examples:


- CPU scheduling policies
- Page replacement policies
- Buffer cache replacement policies
- Disk allocation policies

Mechanism – tool to achieve some effect

Examples:

- Priority scheduling vs. lottery scheduling
- FIFO w/ 2nd chance vs. Clock: a simple FIFO w/ 2nd chance

Separation leads to flexibility



[Ice1] Is there a perfect OS?

Portability

Security


Fairness

Robustness

Efficiency

Interfaces

- Conflicting goals
 - Fairness vs efficiency
 - SJF vs. RR
 - FIFO vs. SCAN
 - Efficiency vs robustness
 - Buffer caching
- Don't know future
 - CPU scheduling
 - Page replacement
 - Disk scheduling




[lec1] There is no magic in OS design

This is Engineering

- Imperfection
 - Don't know future
- Tradeoffs
 - Segmentation vs. paging
 - Read/write API vs. mmap
- Constraints
 - hardware, cost, time
 - FIFO w/ 2nd chance
 - Enhanced version
 - Approx. LRU
- Optimizations
 - After functionality
 - 1-level paging -> 2-level
 - Buffer caching

Nothing's Permanent

- High rate of change
 - Killer-app: Databases/web servers
 - Arch: Multi-core
- Cost / benefit analyses
 - motivation for mmap
 - Semaphore impl on multiprocessor
- One good news:
 - Lots of inertia
 - Principle of locality
 - TLB
 - Demand paging
 - Buffer caching
 - Extra level of indirection
 - Dynamic memory relocation
 - 1-level paging -> 2-level paging
 - UNIX multi-level indexed files



[lec1] About this course...

Principles of OS design

- Some theory
 - SJF optimal
 - Working set modeling
- Some rational
 - Optimize the common case
 - Sequential file access
 - Locality -> caching
 - Why mmap()?
- Lots of practice
 - How much locality?
 - Dist. of file size (UFS inode)
 - I/O access pattern (dir/files)
 - Buffer cache size vs. VM size?

Goals

- Understand OS design decisions
- Basis for future learning

To achieve the goals:

- Learn concepts in class
- Get hands dirty in labs

Great ideas in Computer System Design (1)




- *“All computer science problems can be solved with an extra level of indirection”*

-- David Wheeler

1. Dynamic memory relocation
 - Base&bound, segmentation, paging
2. One-level paging → Two-level paging
3. UFS multi-level indexed files
4. NFS: transparency via VFS

Great ideas in Computer System Design (2)



- Principle of locality → Caching

1. TLB
 2. Demand paging (VM)
 3. Buffer cache in FS
 4. (On-disk cache)
 5. (Client caching in NFS)
-
5. (Hardware cache, L1, L2, etc.)